



Parallel Performance-Energy Predictive Modeling of Browsers: Case Study of Servo

Rohit Zambre*, Lars Bergstrom^, Laleh Aghababaie Beni*, Aparna Chandramowlishwaran*

*University of California, Irvine

^Mozilla Research



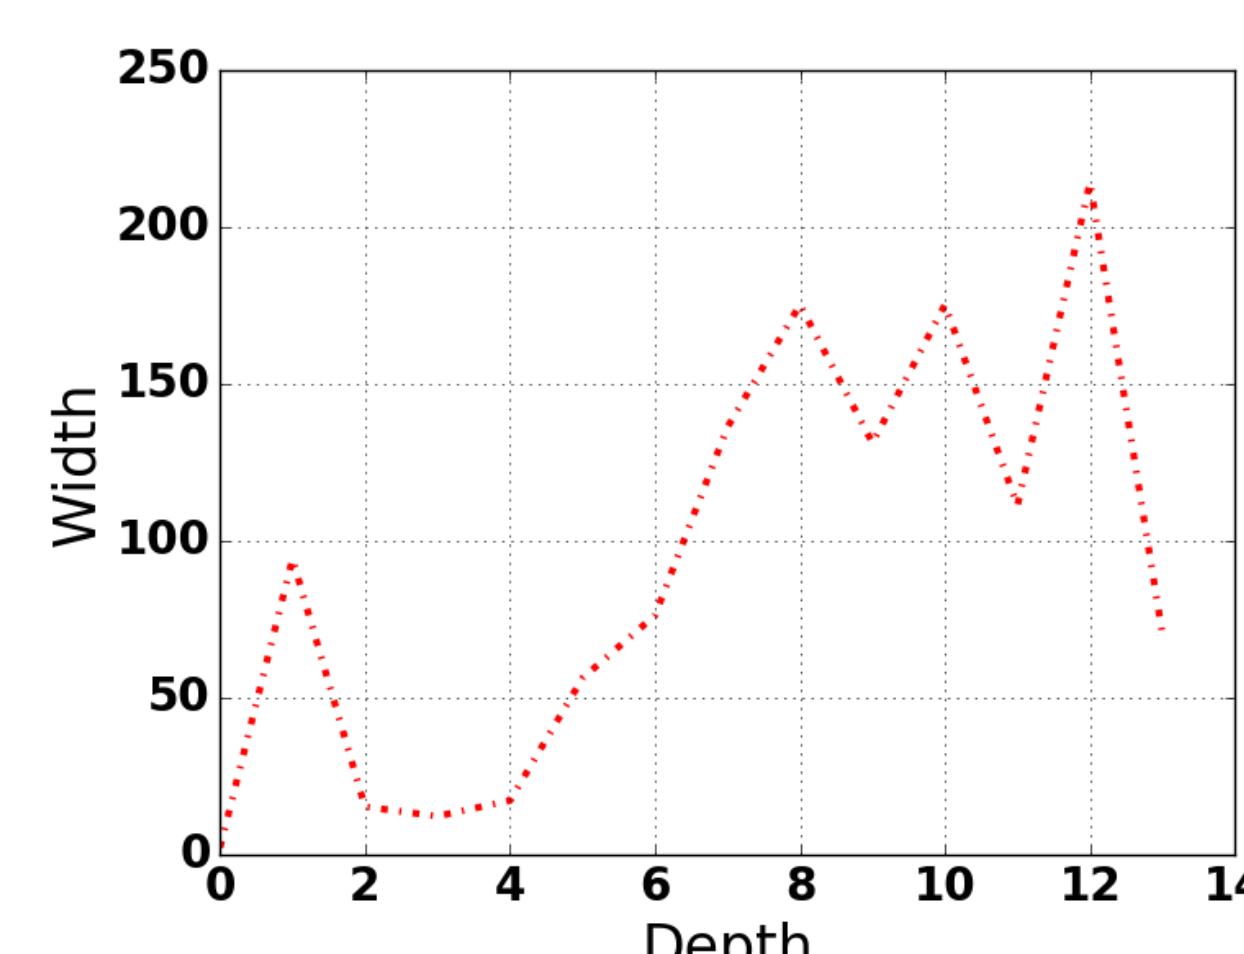
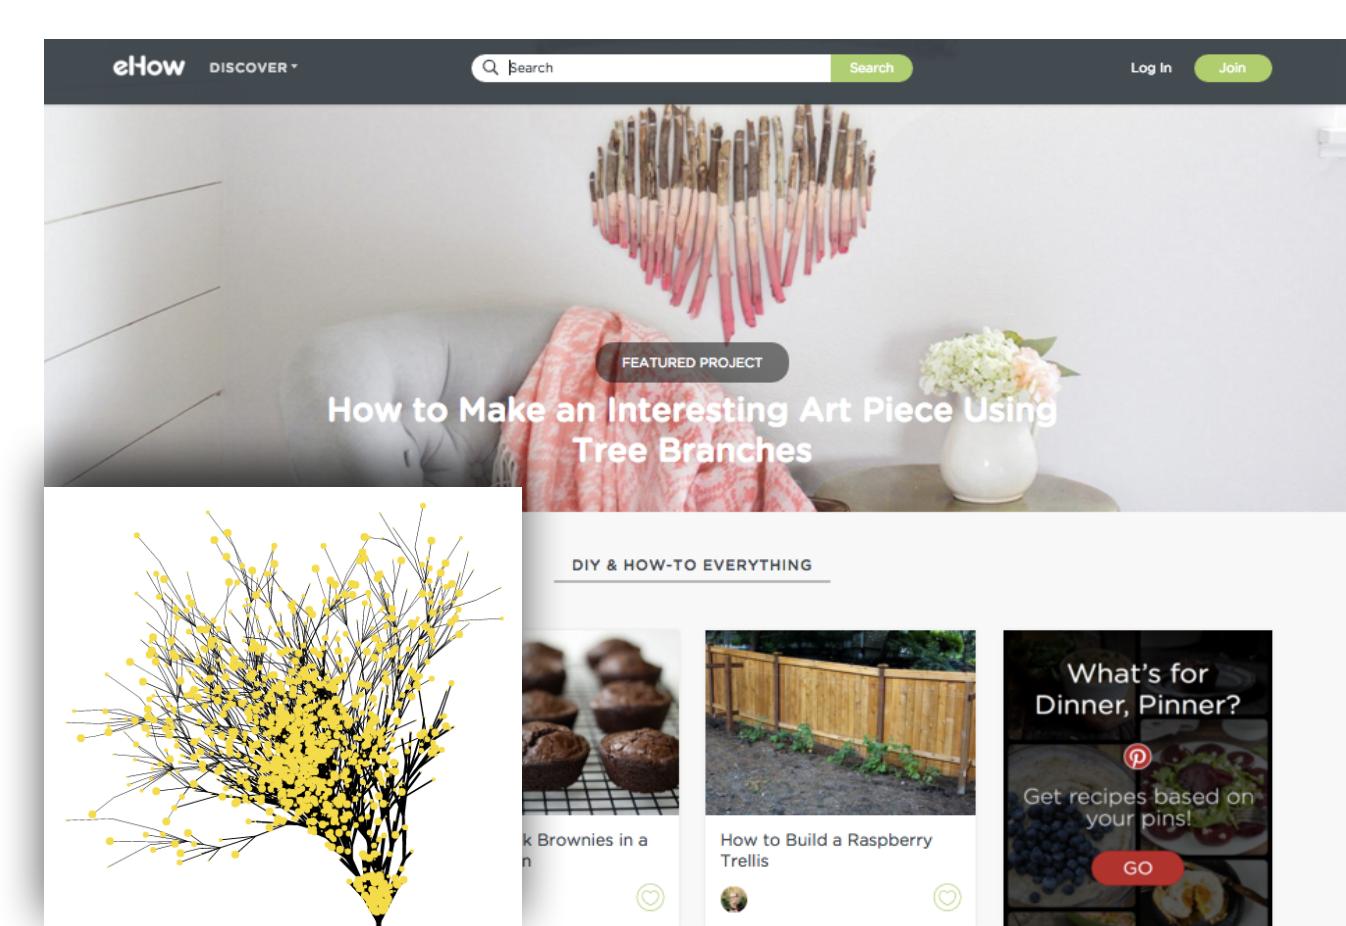
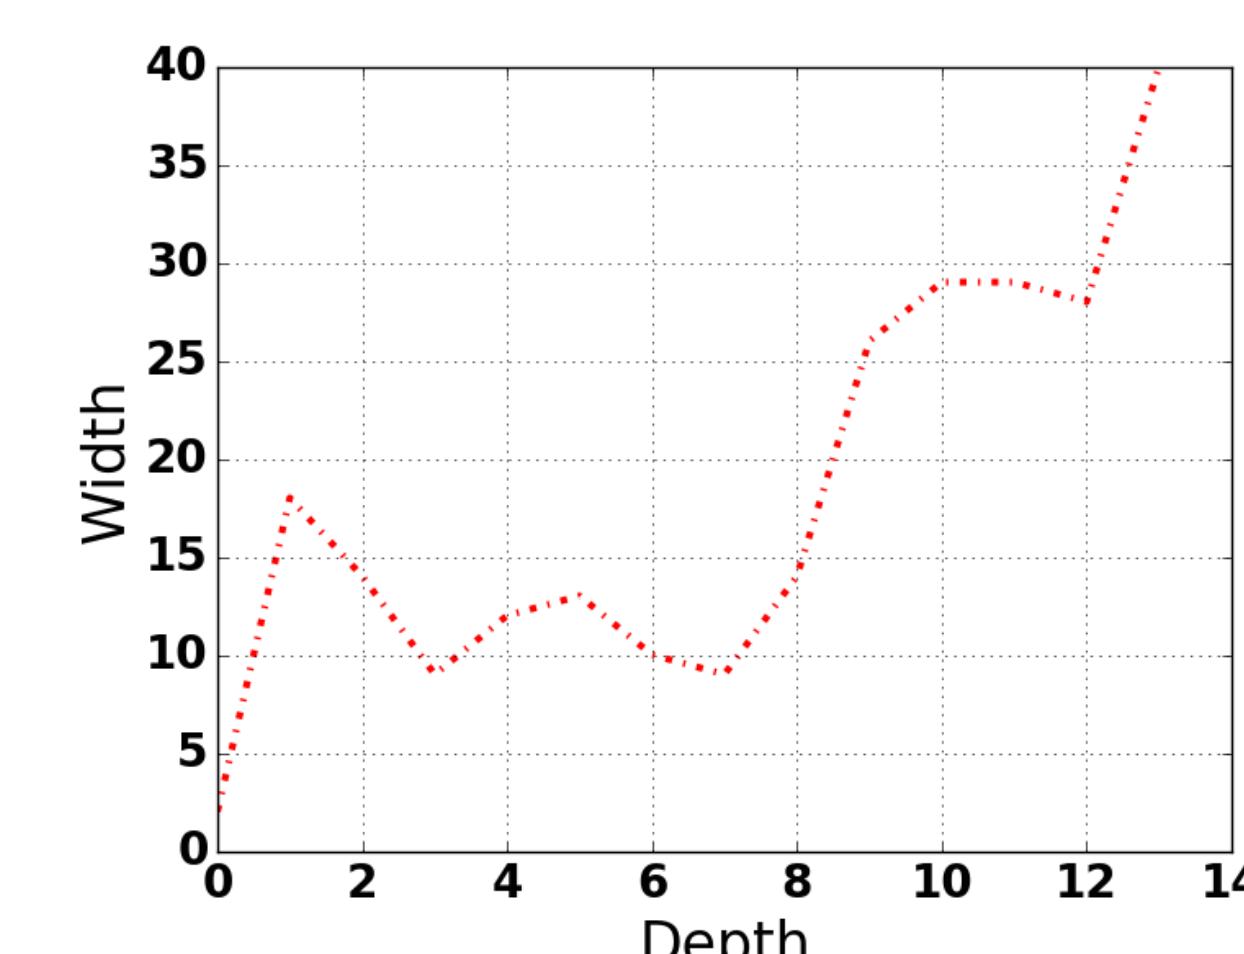
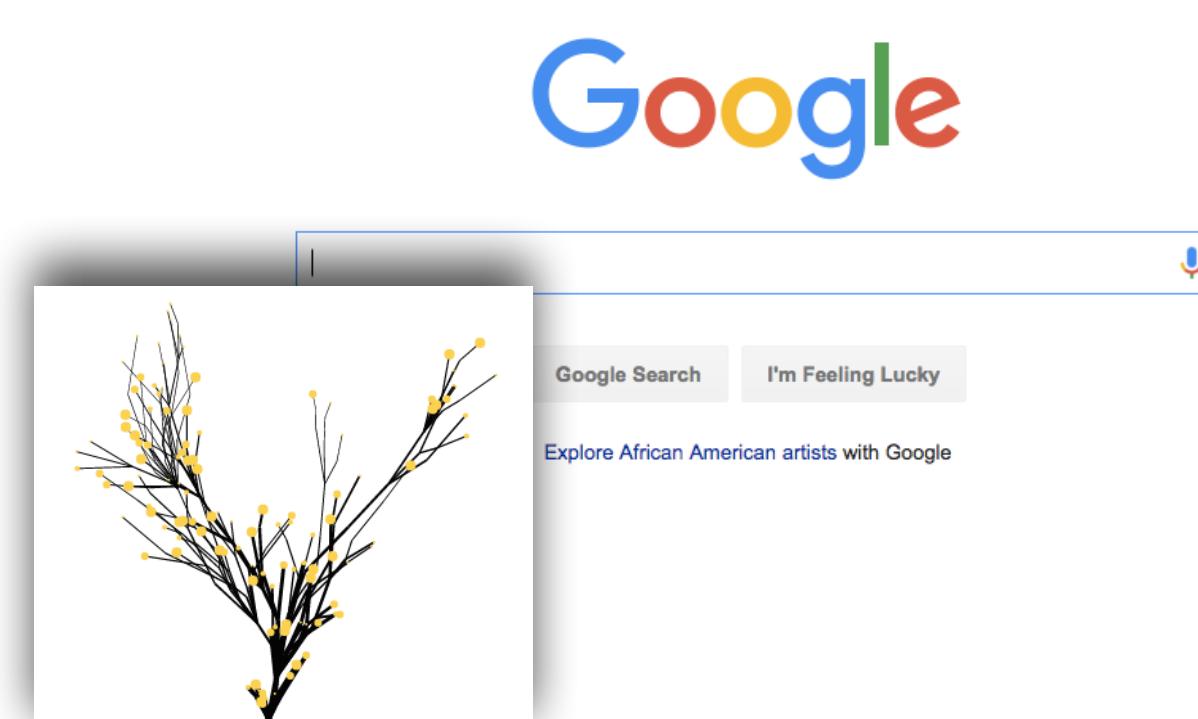
HPC Factory

The Problem & Solution

- What is wrong with today's web browsers?
 - ▷ None exploit multi-core parallelism benefits.
 - ▷ None were designed to handle today's highly complex web pages.
 - ▷ Slow page-load times are a growing concern, especially on mobile platforms.
- What is a solution?
 - ▷ A web browser engine that parallelizes its web rendering tasks.
- But what about parallel overhead?
 - ▷ Need to avoid parallelism when its overhead outweighs its benefits.
- What is a practical solution?
 - ▷ Smart parallelism!
 - ▷ Construct a model to predict degree of parallelism in a web page.
 - ▷ Use model to decide how many threads to spawn for rendering tasks.

Web Page Characterization

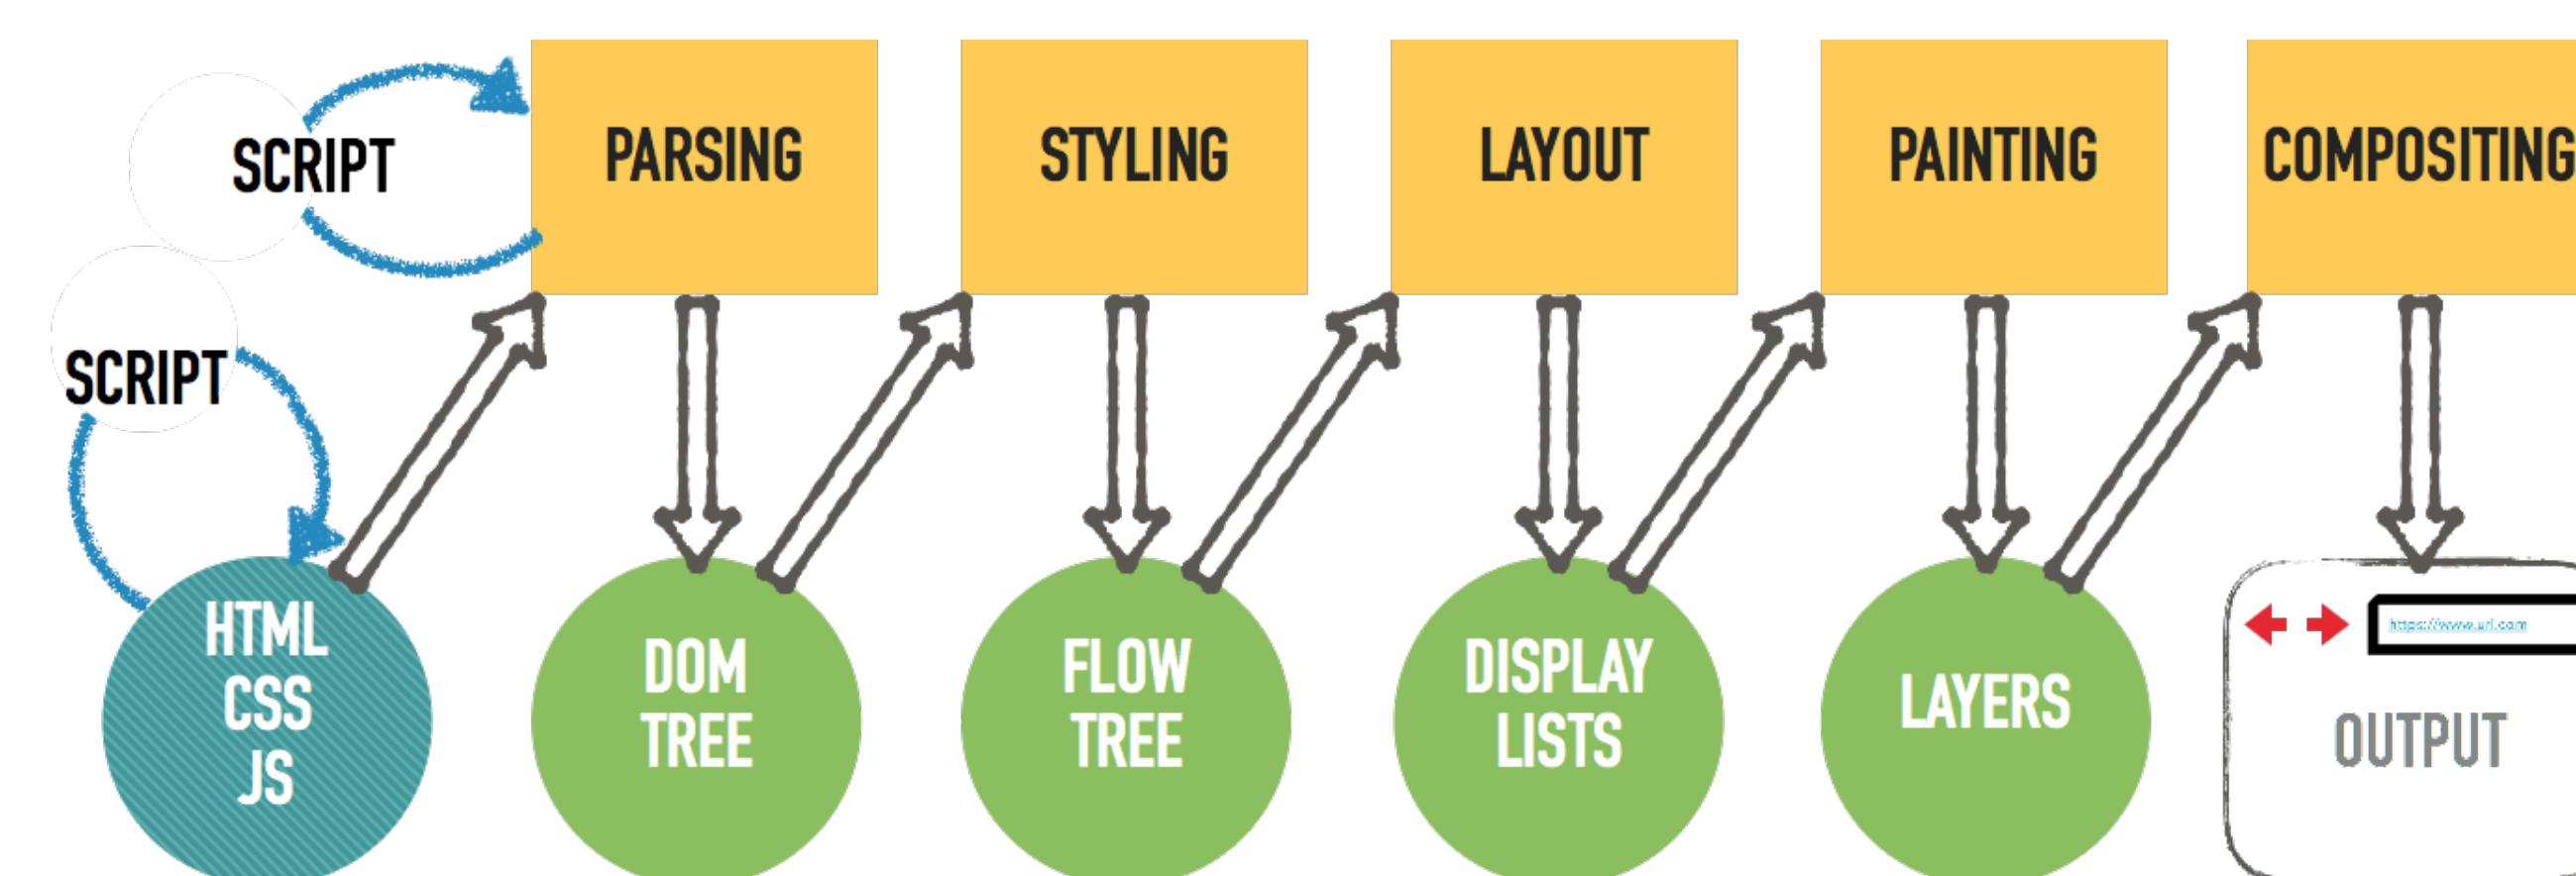
- The workload of a browser is dependent on the web page it is rendering.
- Document Object Model (DOM) tree of a web page describes its structure.



- DOM tree characteristics intuitively predictive of available parallelism:
 - ▷ **DOM-size:** total number of nodes in the DOM tree
 - ▷ **attribute-count:** total number of attributes in the HTML tags
 - ▷ **web-page-size:** size of the web page's HTML in bytes
 - ▷ **number-of-leaves:** number of leaves in the DOM tree
 - ▷ **avg-tree-width:** average number of nodes at each level of the tree
 - ▷ **max-tree-width:** largest number of nodes at a level of the tree
 - ▷ **avg-work-per-level:** average work per level of the tree
- Why characterize a web page?
 - ▷ These DOM tree characteristics are the input features of the predictive model.

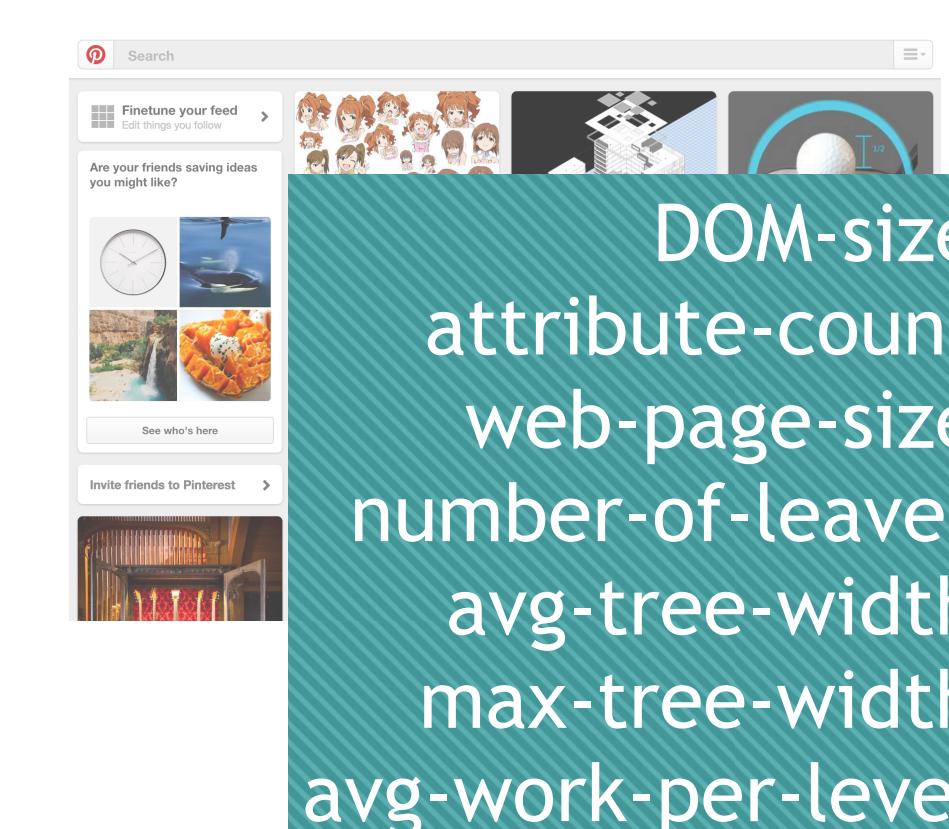
Servo Overview

- Mozilla Research's Servo
 - ▷ A new, parallel web rendering engine
 - ▷ Written in the Rust programming language.
 - ▷ Uses work-stealing scheduler.
 - ▷ Currently parallelism in *Styling* and *Layout* stages.

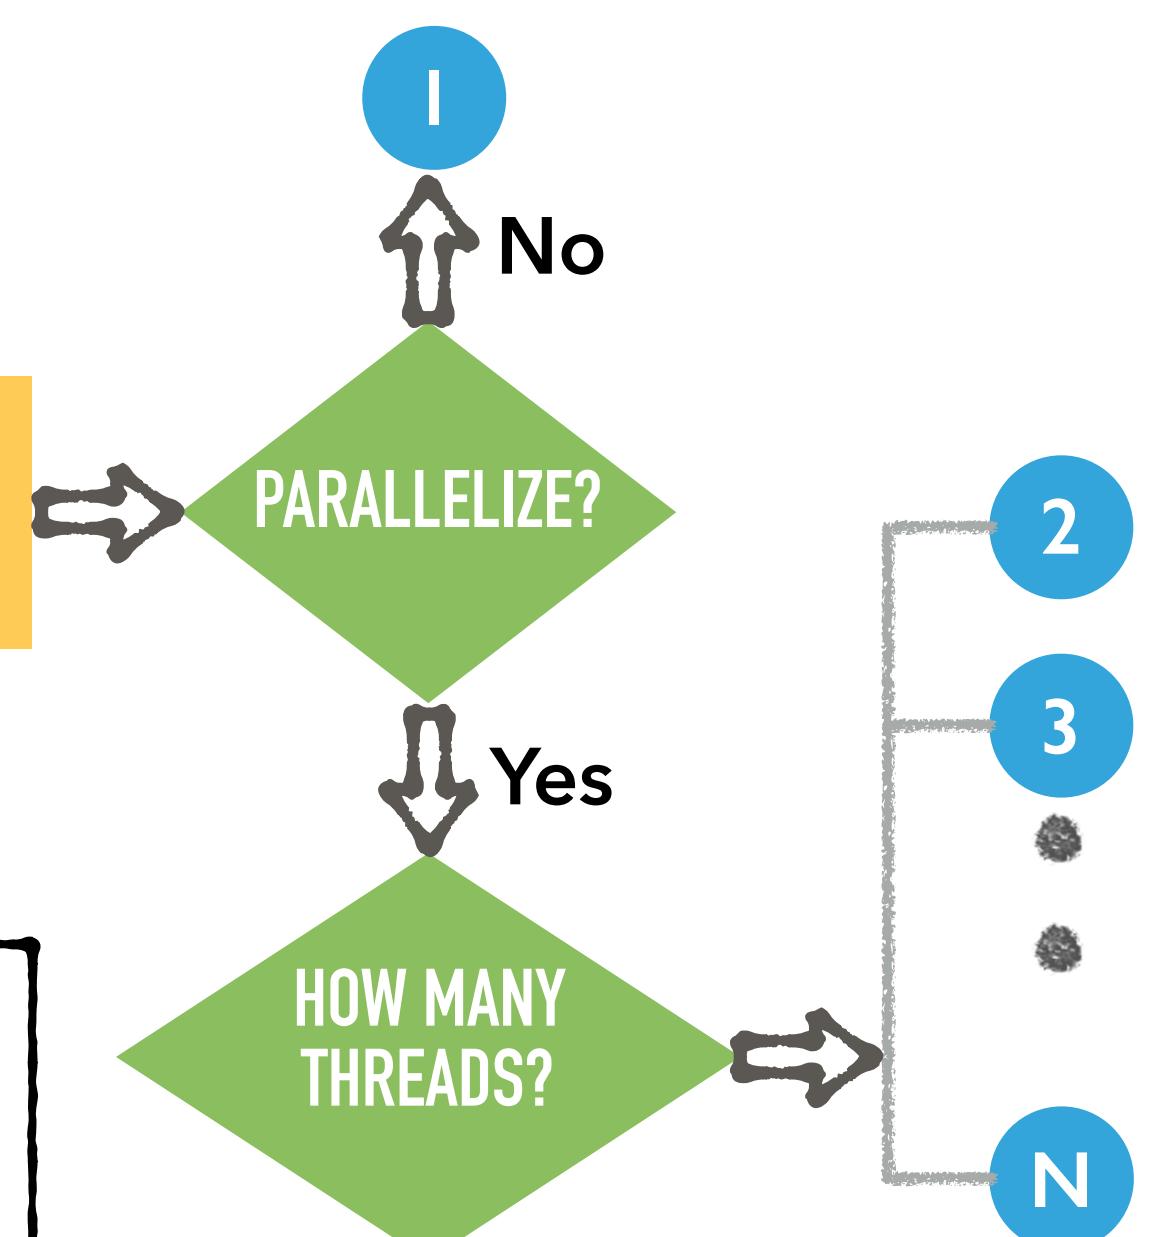


Servo in action!

Modeling & Results



- Off-the-shelf supervised learning.
- Implementation-blind feature space.
- Within complete browser execution.
- Confident accuracies & high savings!



Supervised Learning Algorithm	Evaluation Method	Average Accuracy	Maximum Accuracy
Multinomial Logistic Regression (MNR)	Cross-validation	72.22%	87.27%
Ensemble Learning	Cross-validation	71.27%	83.63%
Neural Network	Holdout-set	77.8%	-

- Why is accuracy not > 90%?

- ▷ Heavy class imbalance.
- ▷ Feasible training data set is small.
- ▷ Servo is a young software.

- How much are we saving?
 - ▷ Up to 91.39% on performance.
 - ▷ Up to 46.32% on energy.

Automated Labeling

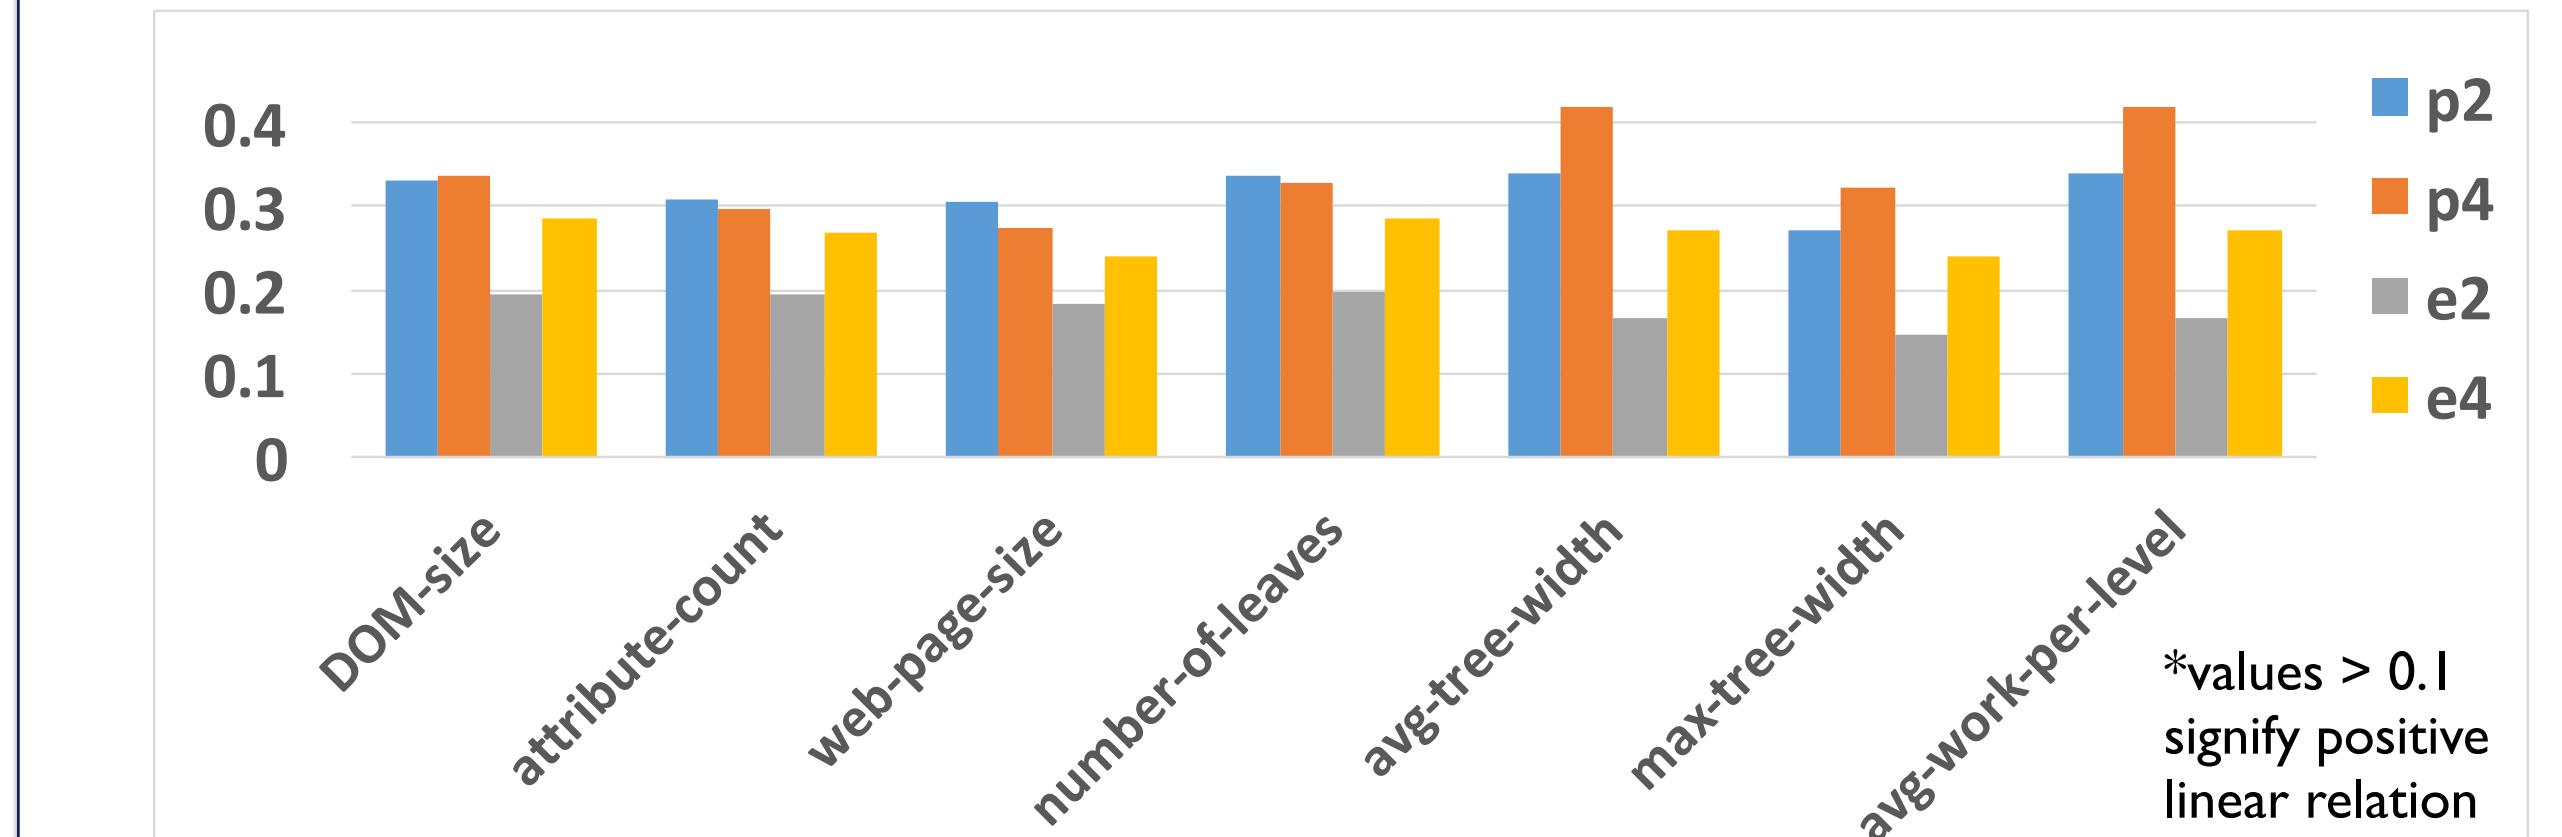
- Why automated labeling?
 - ▷ Eliminates need for domain expert to manually label training data.
- Three cost models
 - ▷ Performance Cost Model
 - ▷ Energy Cost Model
 - ▷ Performance & Energy Cost Model
- Labeling flow
 1. Collect performance & energy data for each web page with 1, 2, 4, 8, ... threads.
 2. For each web page, compute speedups and greenups for each thread.
 3. Using a tuned cost model, label the web page.
- How to label using the Performance/Energy Cost Model?
 - ▷ Using thread that achieves optimal performance/energy.
- How to label using the Performance & Energy Cost Model?
 - ▷ Using Performance-Energy Tuple (PET) buckets.
 - ▷ Each PET bucket has a greenup limit.
 - ▷ Choose thread from highest PET bucket that satisfies greenup limit.

Our Labeling Settings	
X_1, X_2	$X_1 = 1.1; X_2 = 1.3$
X_2, X_3	$X_2 = 1.3; X_3 = 12.87$
Y_1	0.9
Y_2	0.85

Our Experimental Data			
Samples	Threads	Performance & Energy Labels	
535 (Alexa Top 500 & 2012 Fortune 1000)	I, 2, 4		
Platform	Quad-core Intel Ivy Bridge MacBook Pro	I: 59.25% 2: 9.34% 4: 31.40%	

Conclusion & Vision

- Evident correlation* between DOM-tree features and parallel work.



- Guided, work-load aware scheduling
 - ▷ Use model to guide scheduling and maintain load-balancing.
- big.LITTLE for parallel web browsers
 - ▷ Exploit and model big.LITTLE parallelism for web rendering tasks.
- App-clutter-free smartphones
 - ▷ Demonstrate greener and higher performance than native apps.
- Revive the universal Web platform
 - ▷ Shift focus of app development to the Web, accessible by all device-types.