

ROHIT ZAMBRE

---

# SCALABLE COMMUNICATION ENDPOINTS FOR MPI+THREADS APPLICATIONS

## WHO WILL YOU BE LISTENING TO FOR THE NEXT 40 MINS?

- ▶ An intern




- ▶ Ph.D. student @ University of California, Irvine
  - ▶ HPC Forge research group
- ▶ Visiting student @ Argonne National Laboratory
  - ▶ Programming Models and Runtime Systems group
- ▶ This talk: research on multiple endpoints for MPI

**A STANDARD FOR EXPLICIT  
COMMUNICATION BETWEEN PROCESSES  
RUNNING IN PARALLEL.**

**Message Passing Interface (MPI)**

## MPI AND PARALLELISM

- ▶ Parallelism from multiple processes. MPI for coordinating.



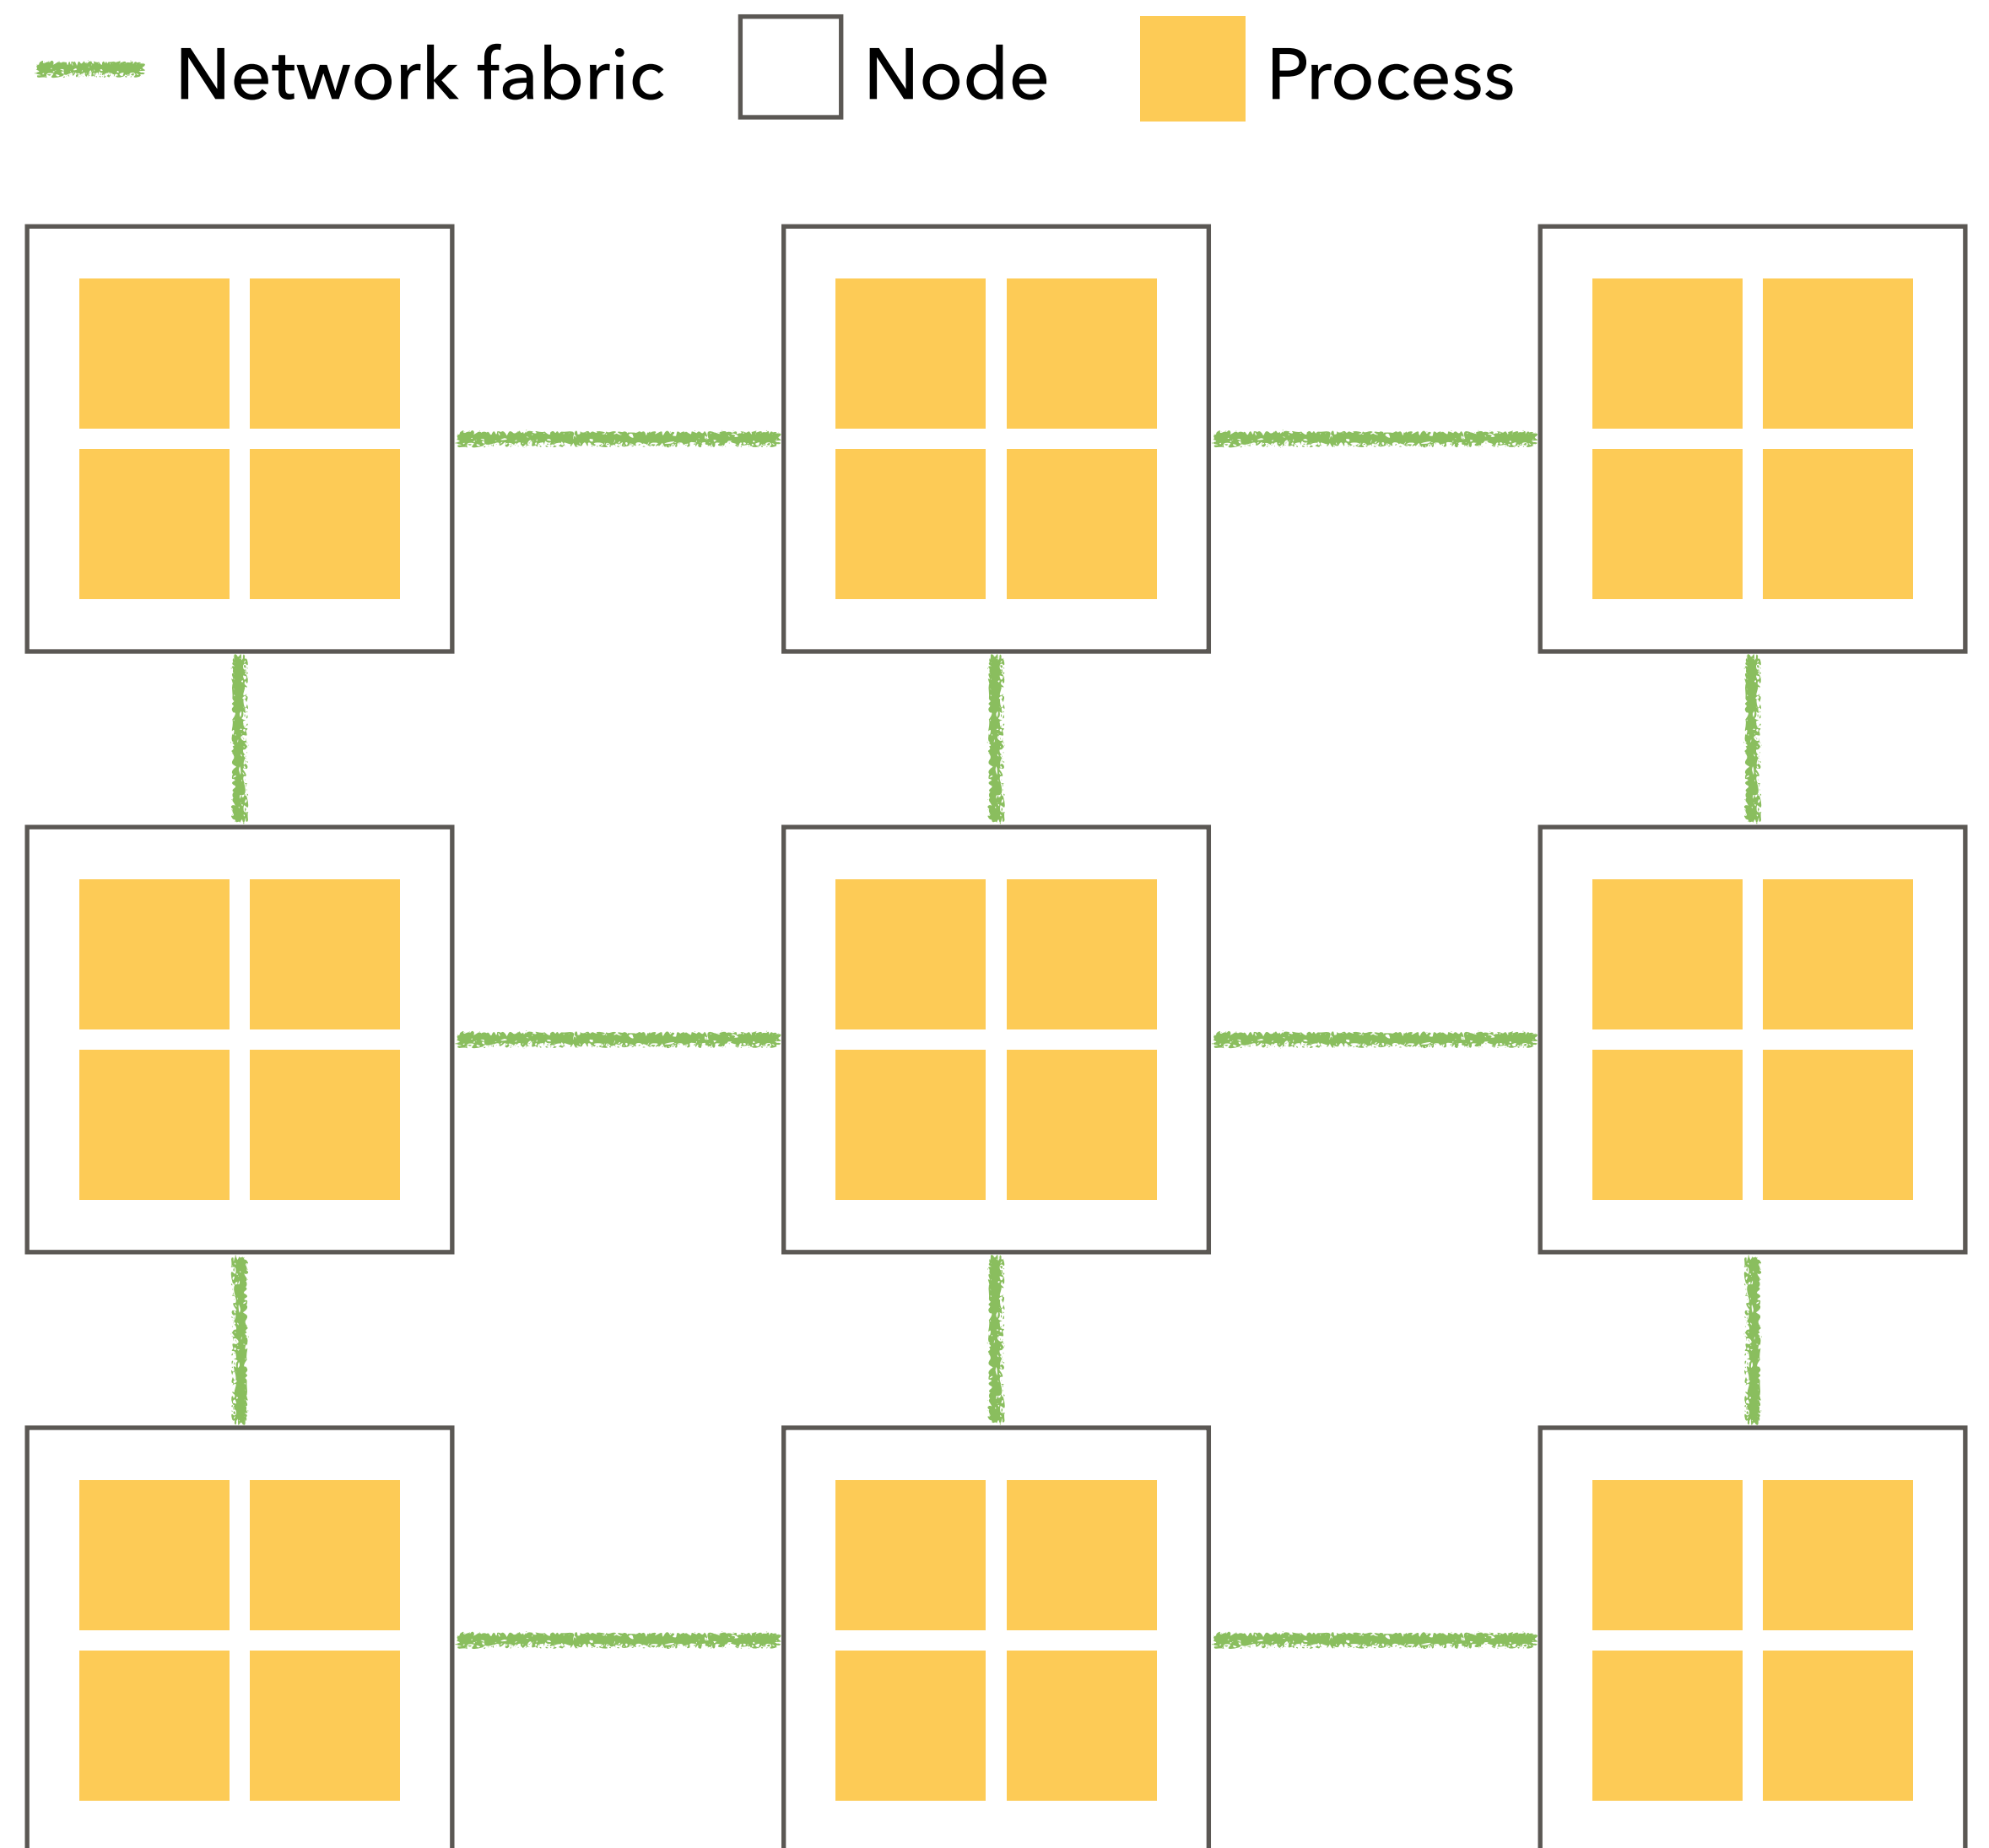
- ▶ MPI not an implementation. MPI forum (<http://www.mpi-forum.org>)
- ▶ MPICH, OpenMPI, MVAPICH, Intel MPI, Microsoft MPI, etc.
- ▶ Widely used: scientific simulations, modeling, USPS, etc.

# TRADITIONAL MPI: MPI EVERYWHERE

App ignores location of processes

Leads to static split of on-node  
resources

Cores scaling a lot faster than TLB,  
network resources, etc.




# HYBRID MPI: MPI+THREADS

e.g. MPI+OpenMP

All threads access all on-node  
resources

Handy compiler support for MPI user




# STATE OF THE ART

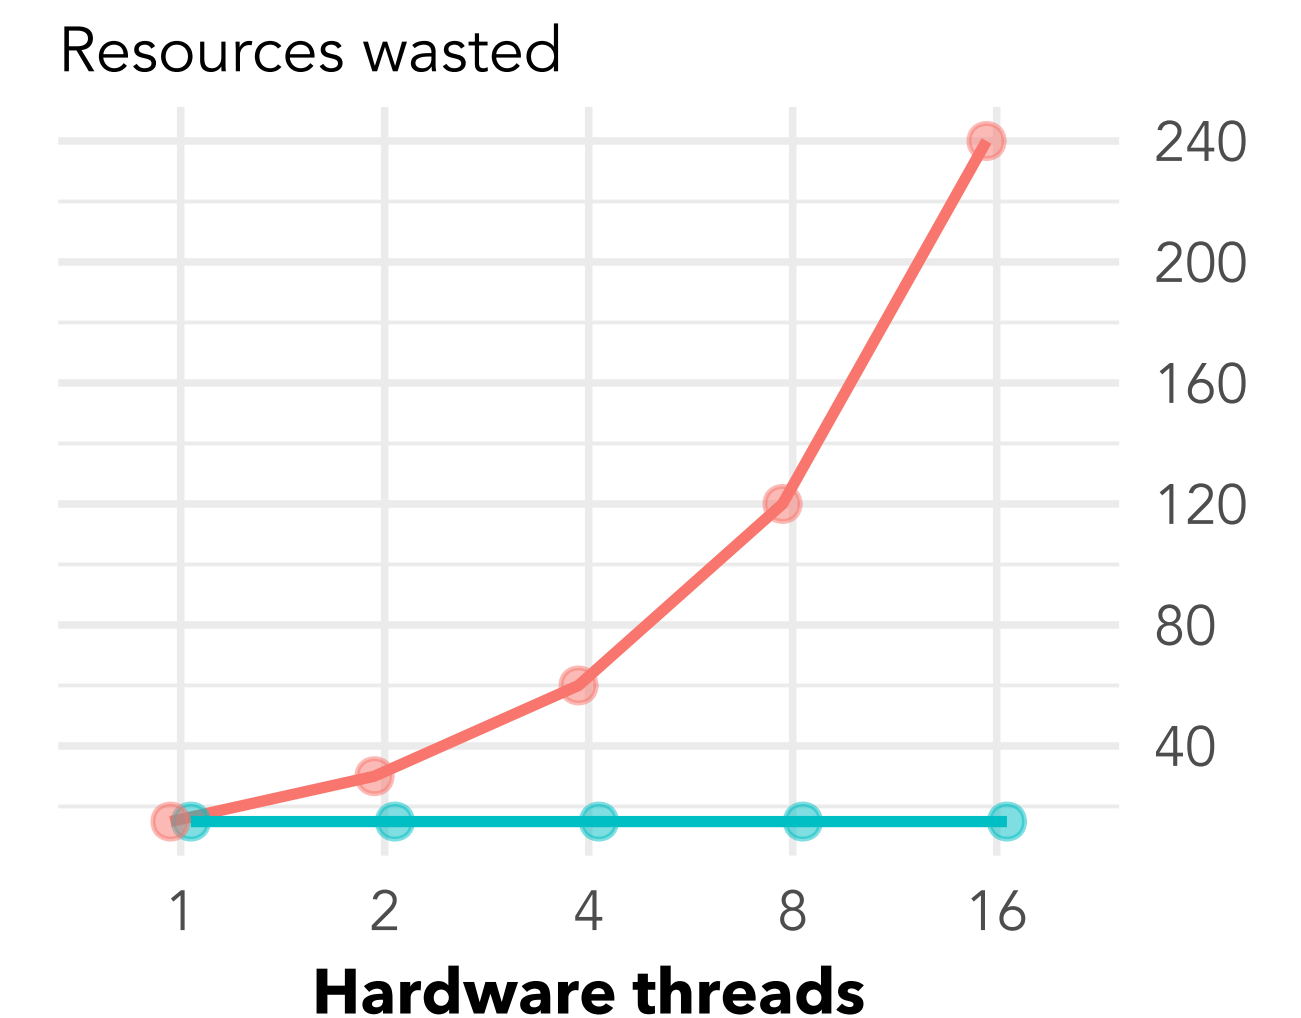
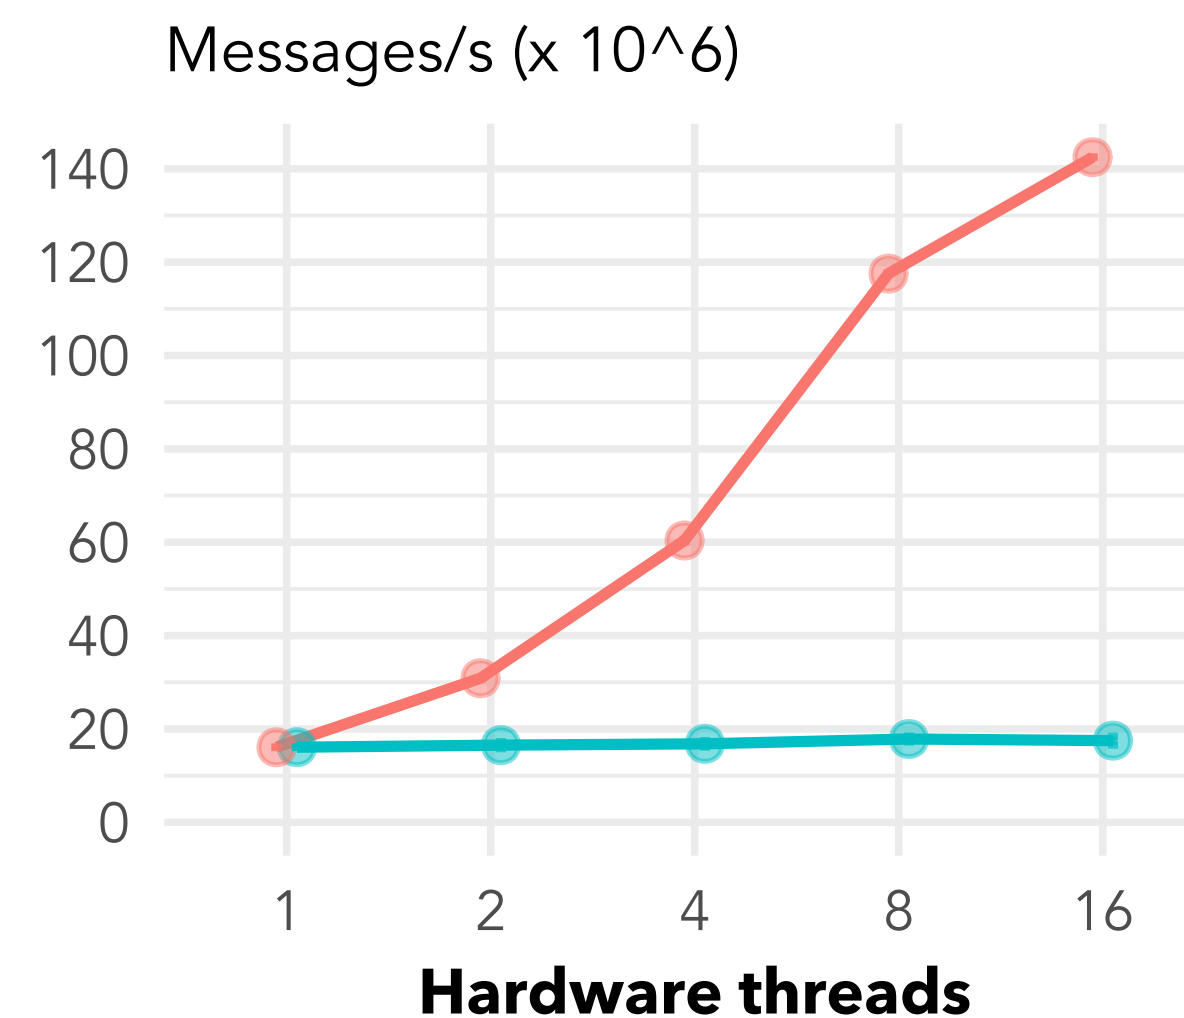

MPI everywhere achieves maximum possible throughput but wastes 93.75% of hardware resources

MPI+threads achieves maximum resource efficiency but performs up to 9x worse

MPI everywhere



MPI+threads



Endpoint type — MPI everywhere — MPI+threads


**WHAT LEVEL OF  
RESOURCE SHARING IS  
IDEAL?**



# WHAT LEVEL OF RESOURCE SHARING IS IDEAL?

**It depends**

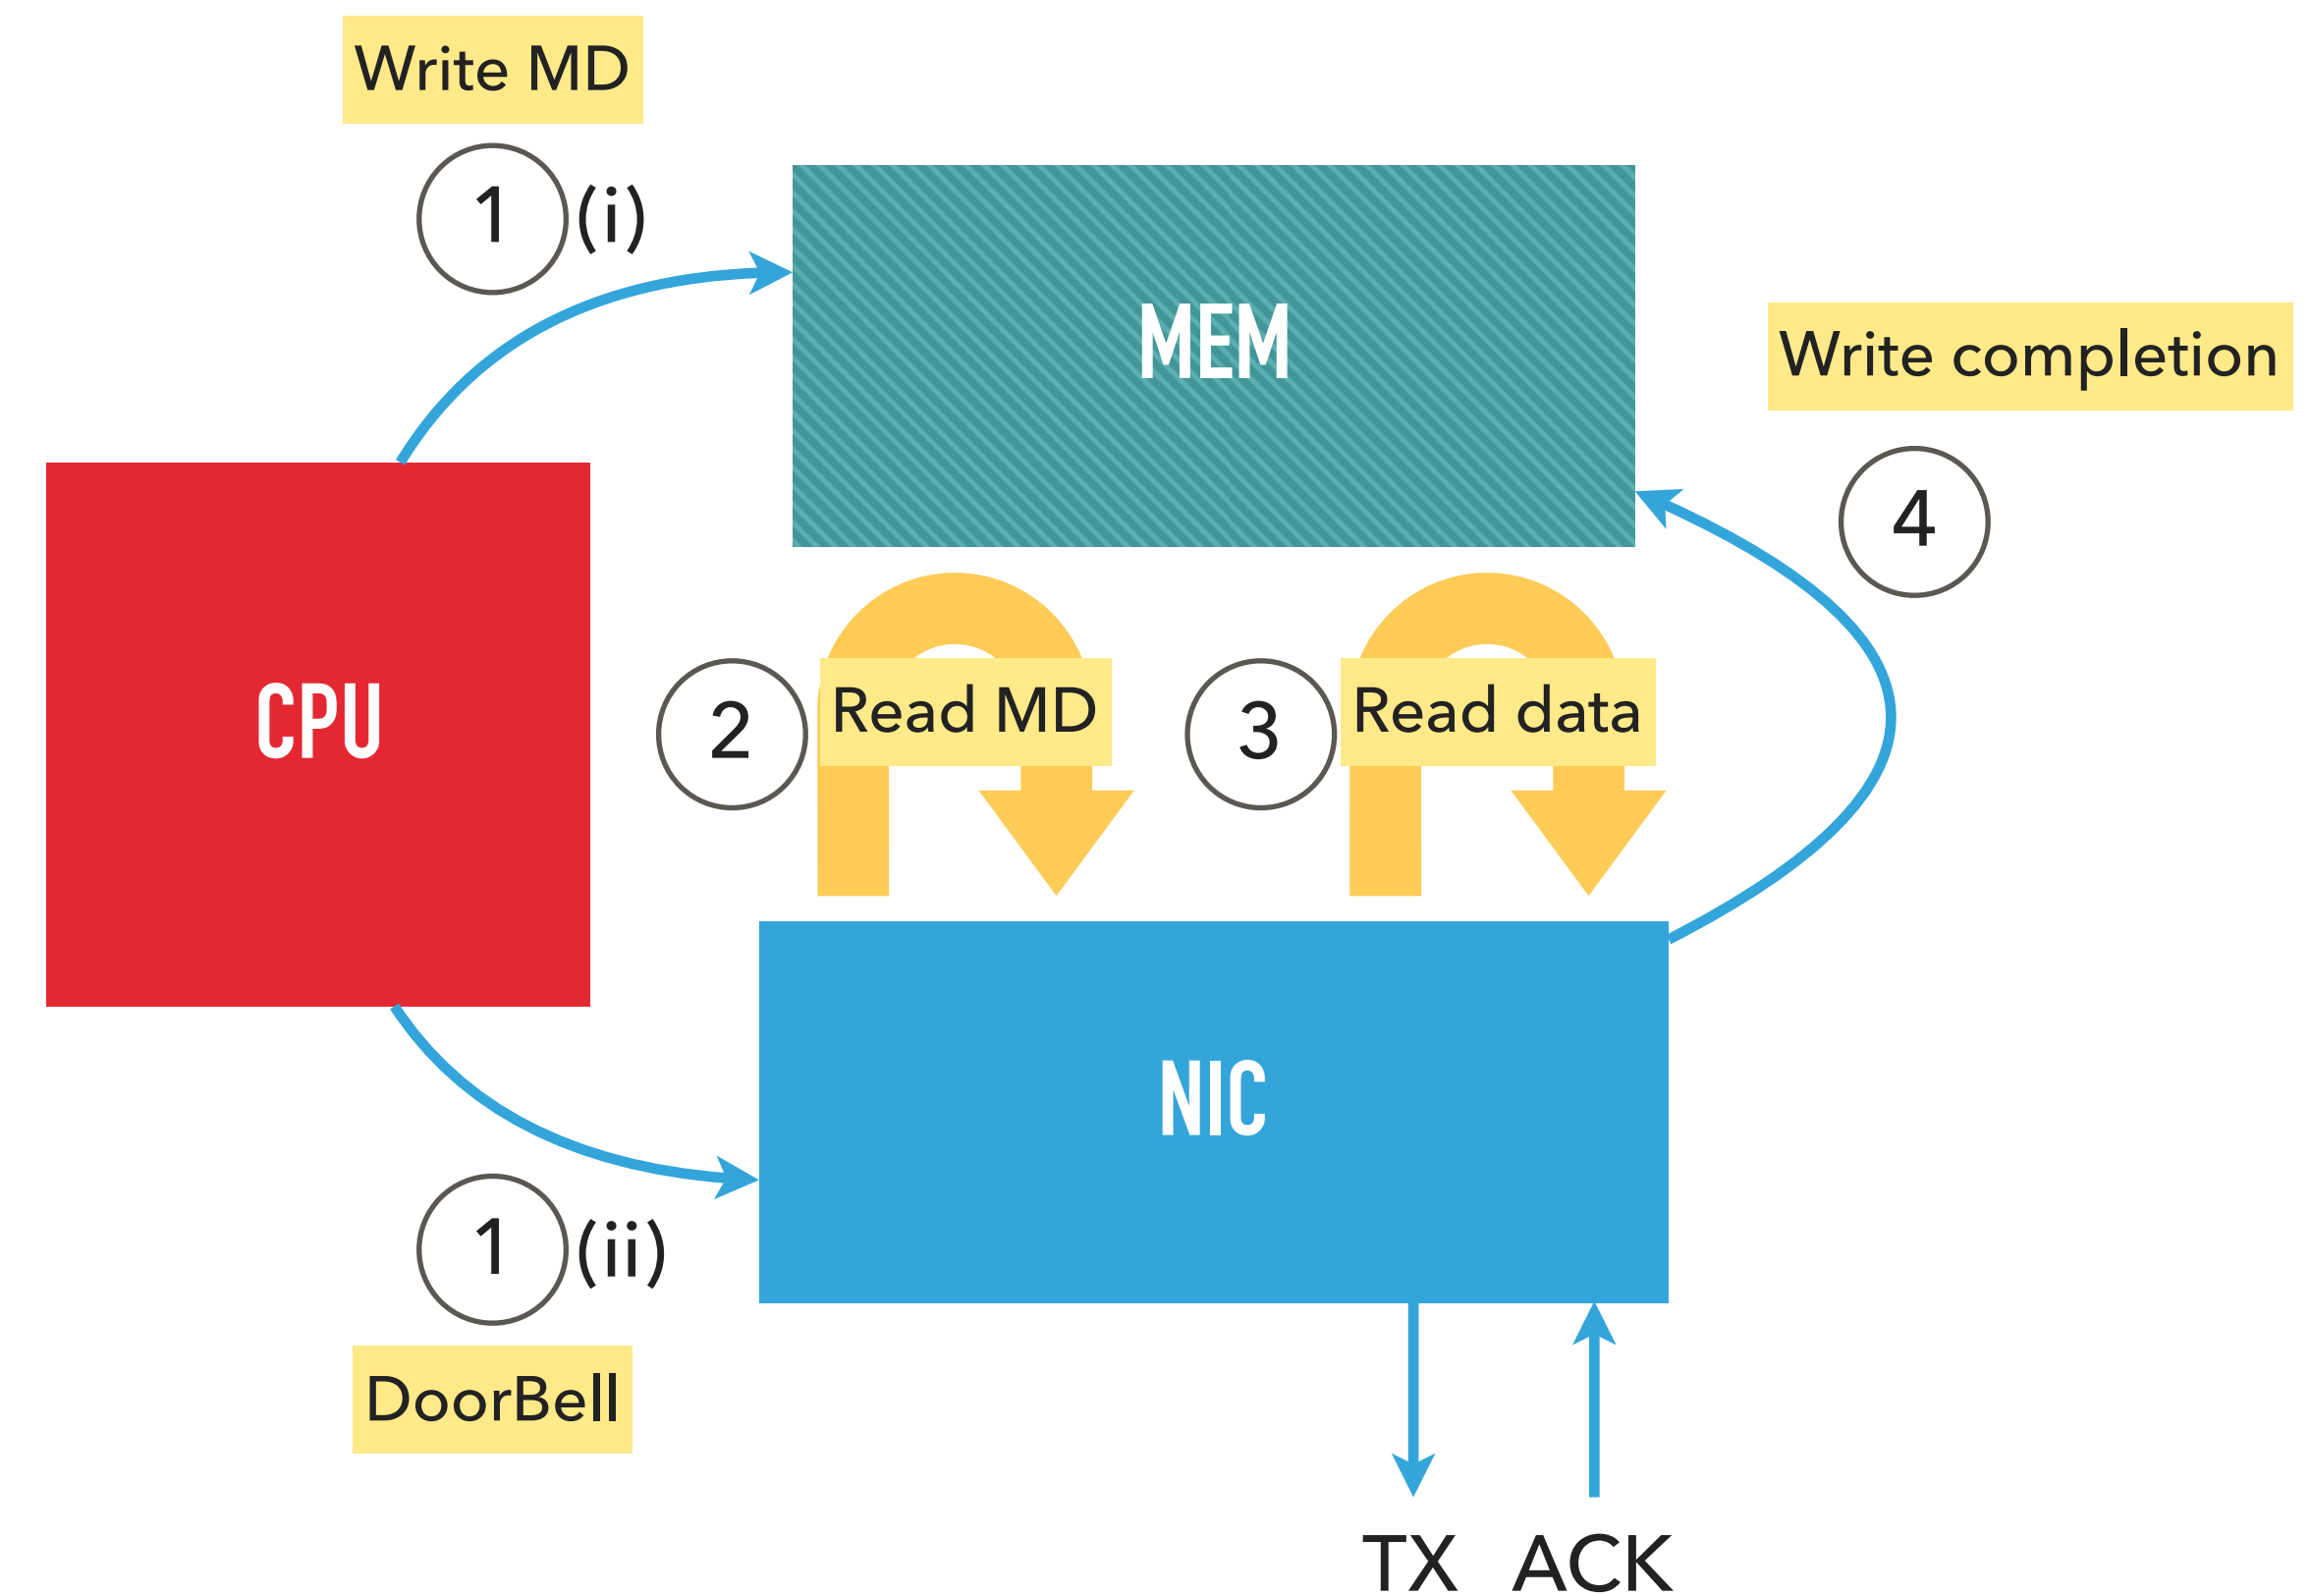
## INTER-NODE COMMUNICATION



- ▶ The Network Interface Card (NIC) is the node's communication portal.
- ▶ The software (CPU) coordinates with the hardware (NIC) to *initiate* a transfer and confirm its *completion*.

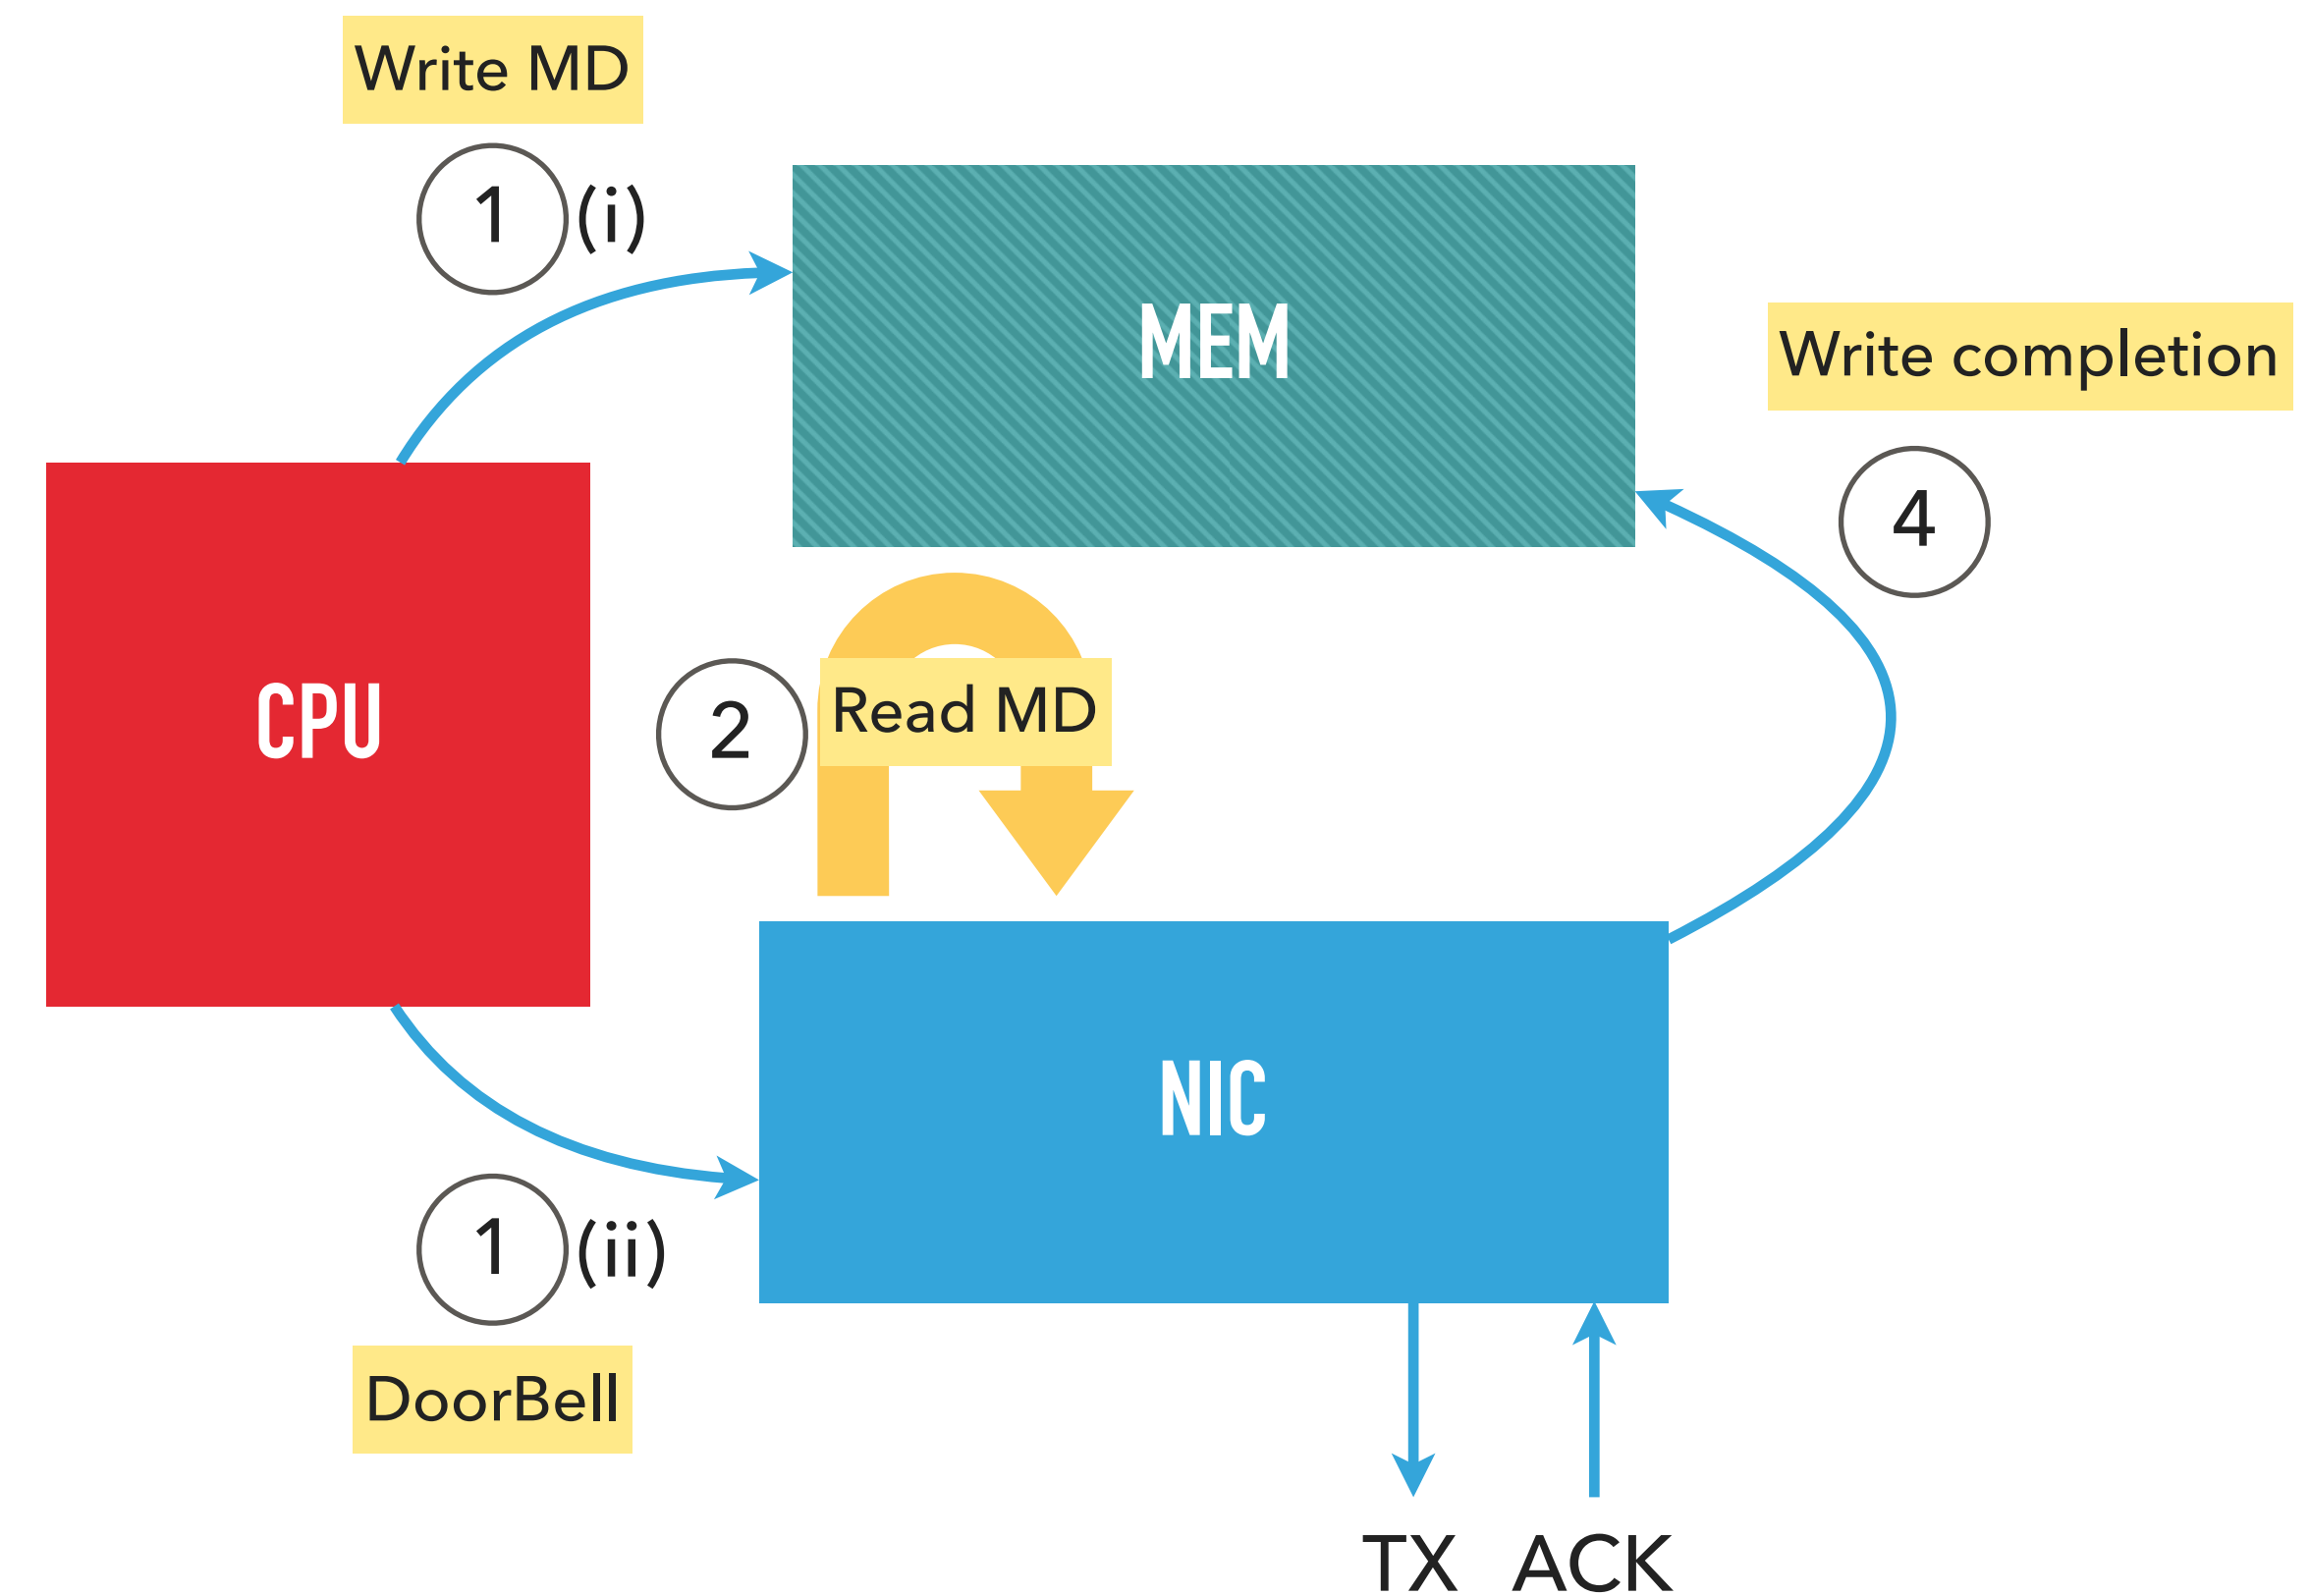
## CPU-NIC INTERACTION: OVERVIEW

1. (i) Write message descriptor to MEM  
(ii) DoorBell (atomic 8-byte MMIO-write to NIC)
2. DMA-read to get message descriptor
3. DMA-read to get payload (data)
4. DMA-write completion



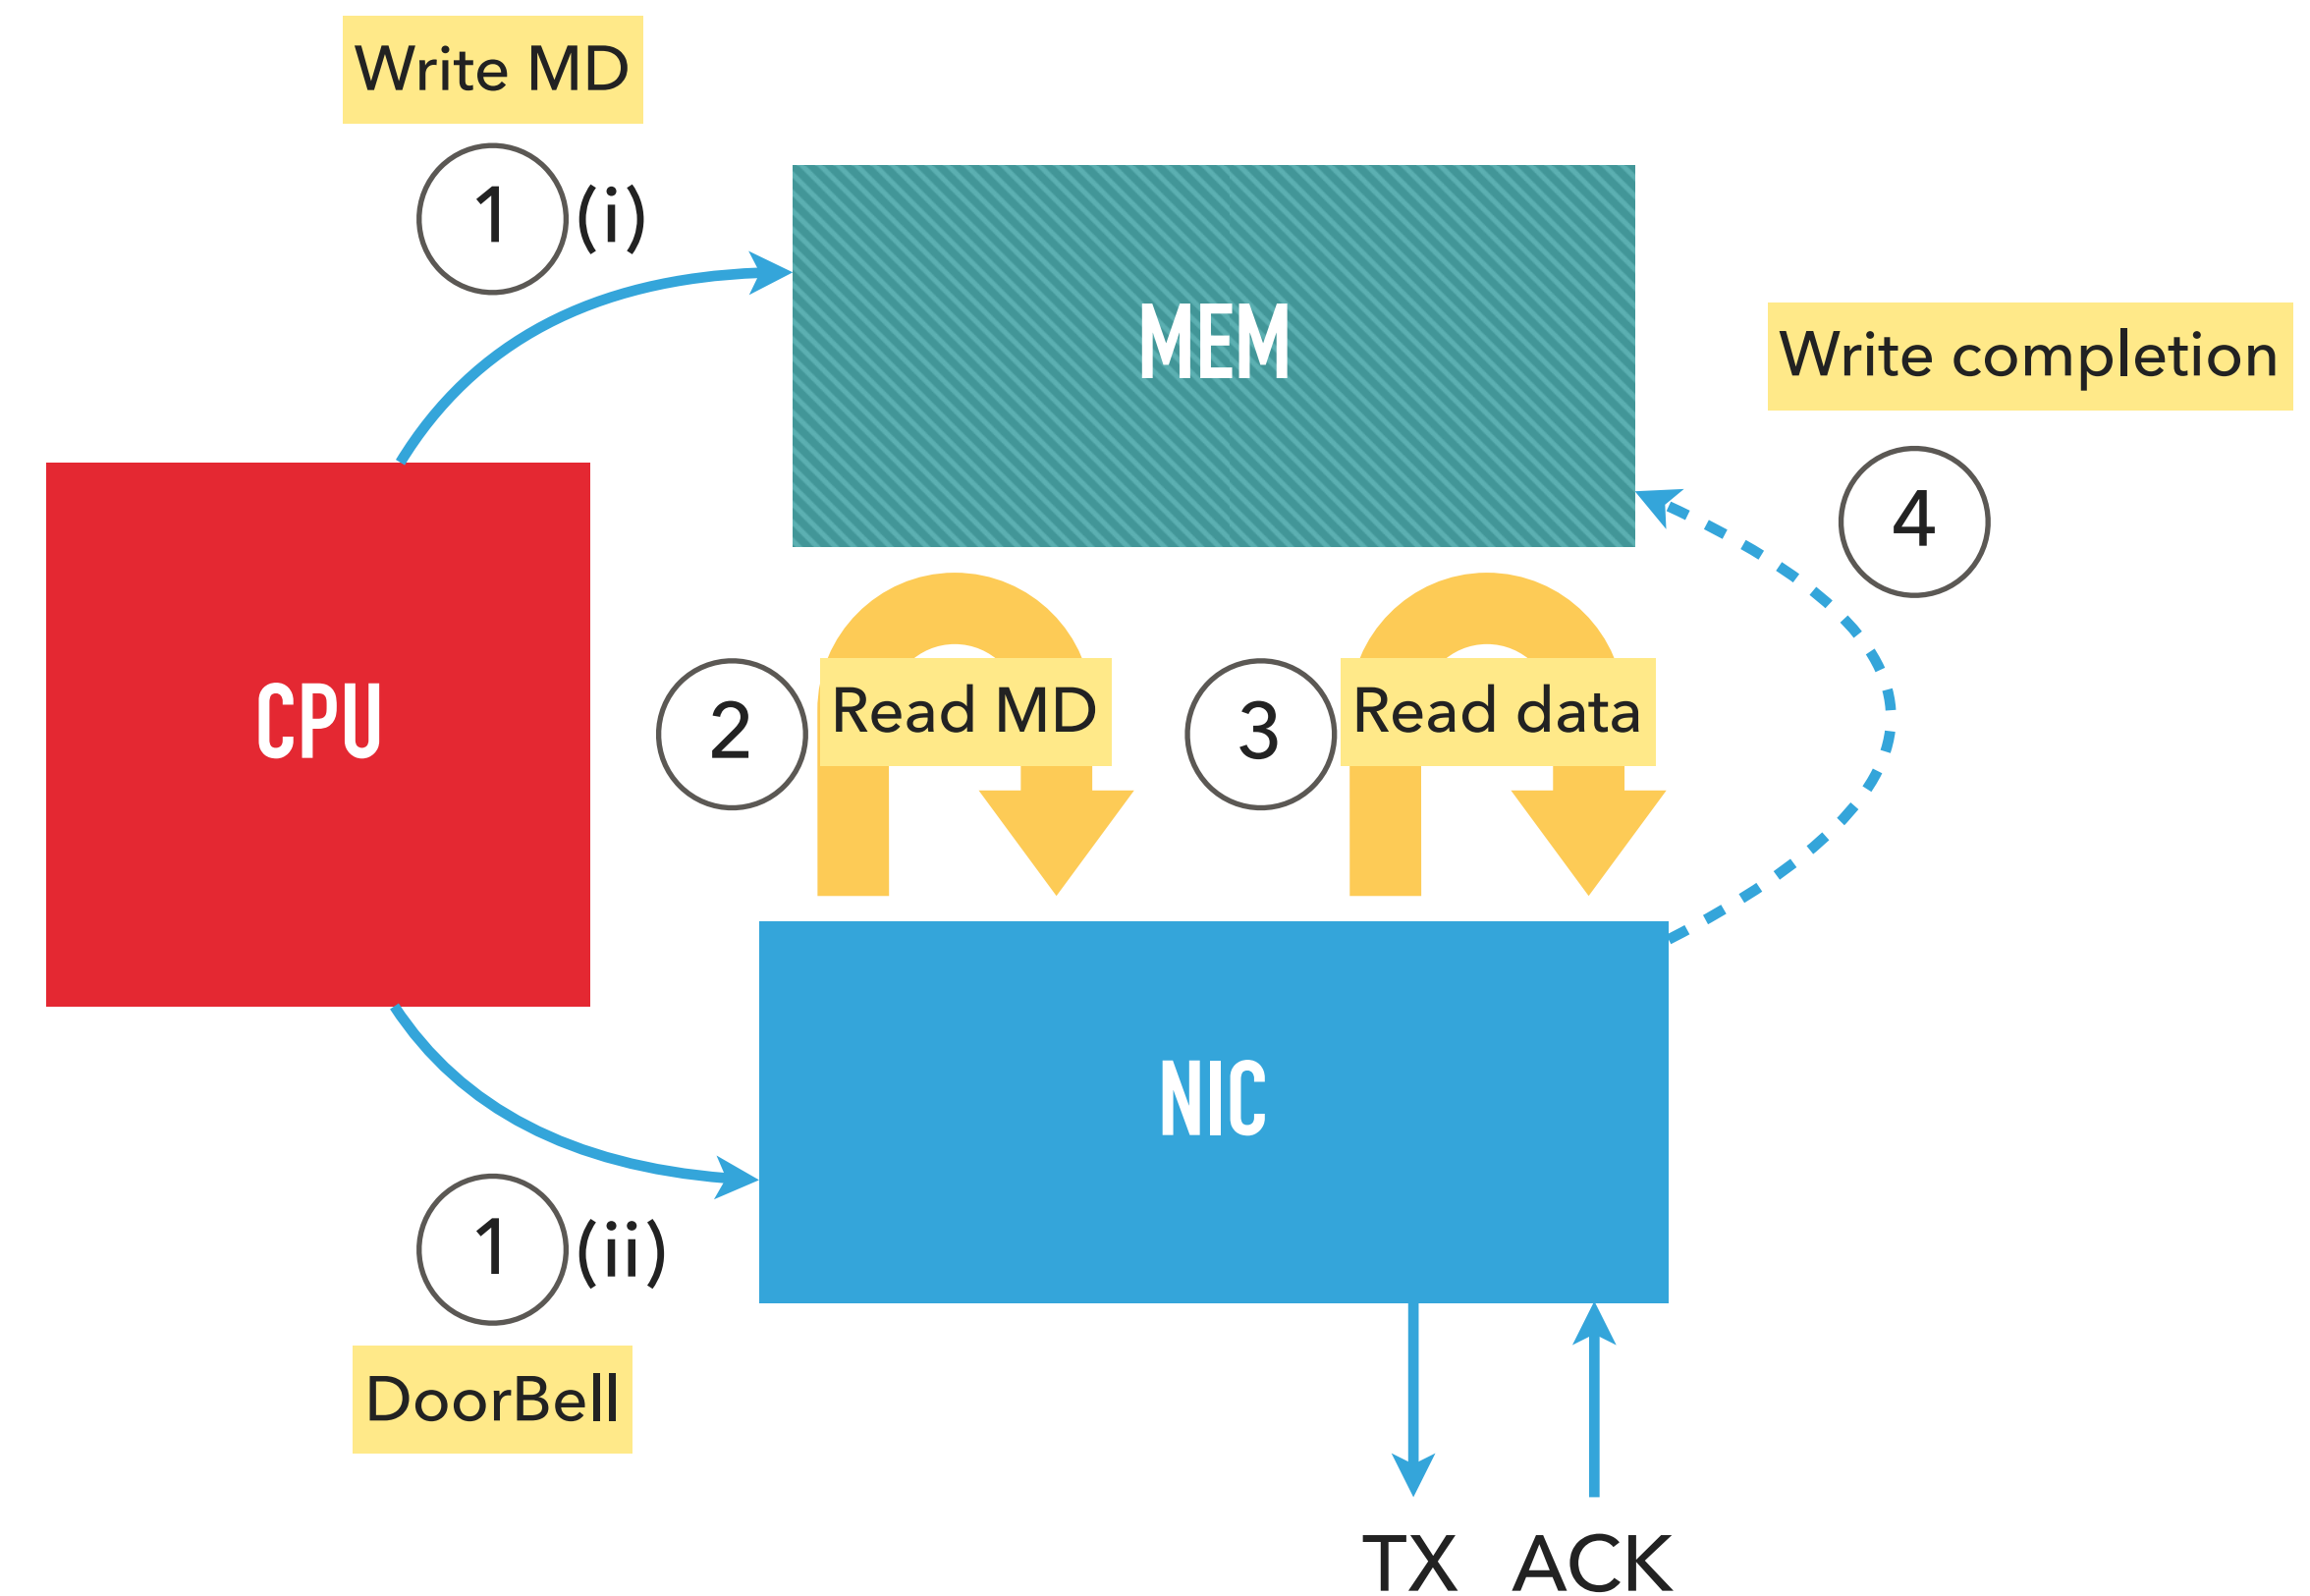
## CPU-NIC INTERACTION FEATURE: INLINING

- ▶ Payload part of message descriptor.
  - ▶ CPU reads the data instead of NIC
- ▶ Removes (3)



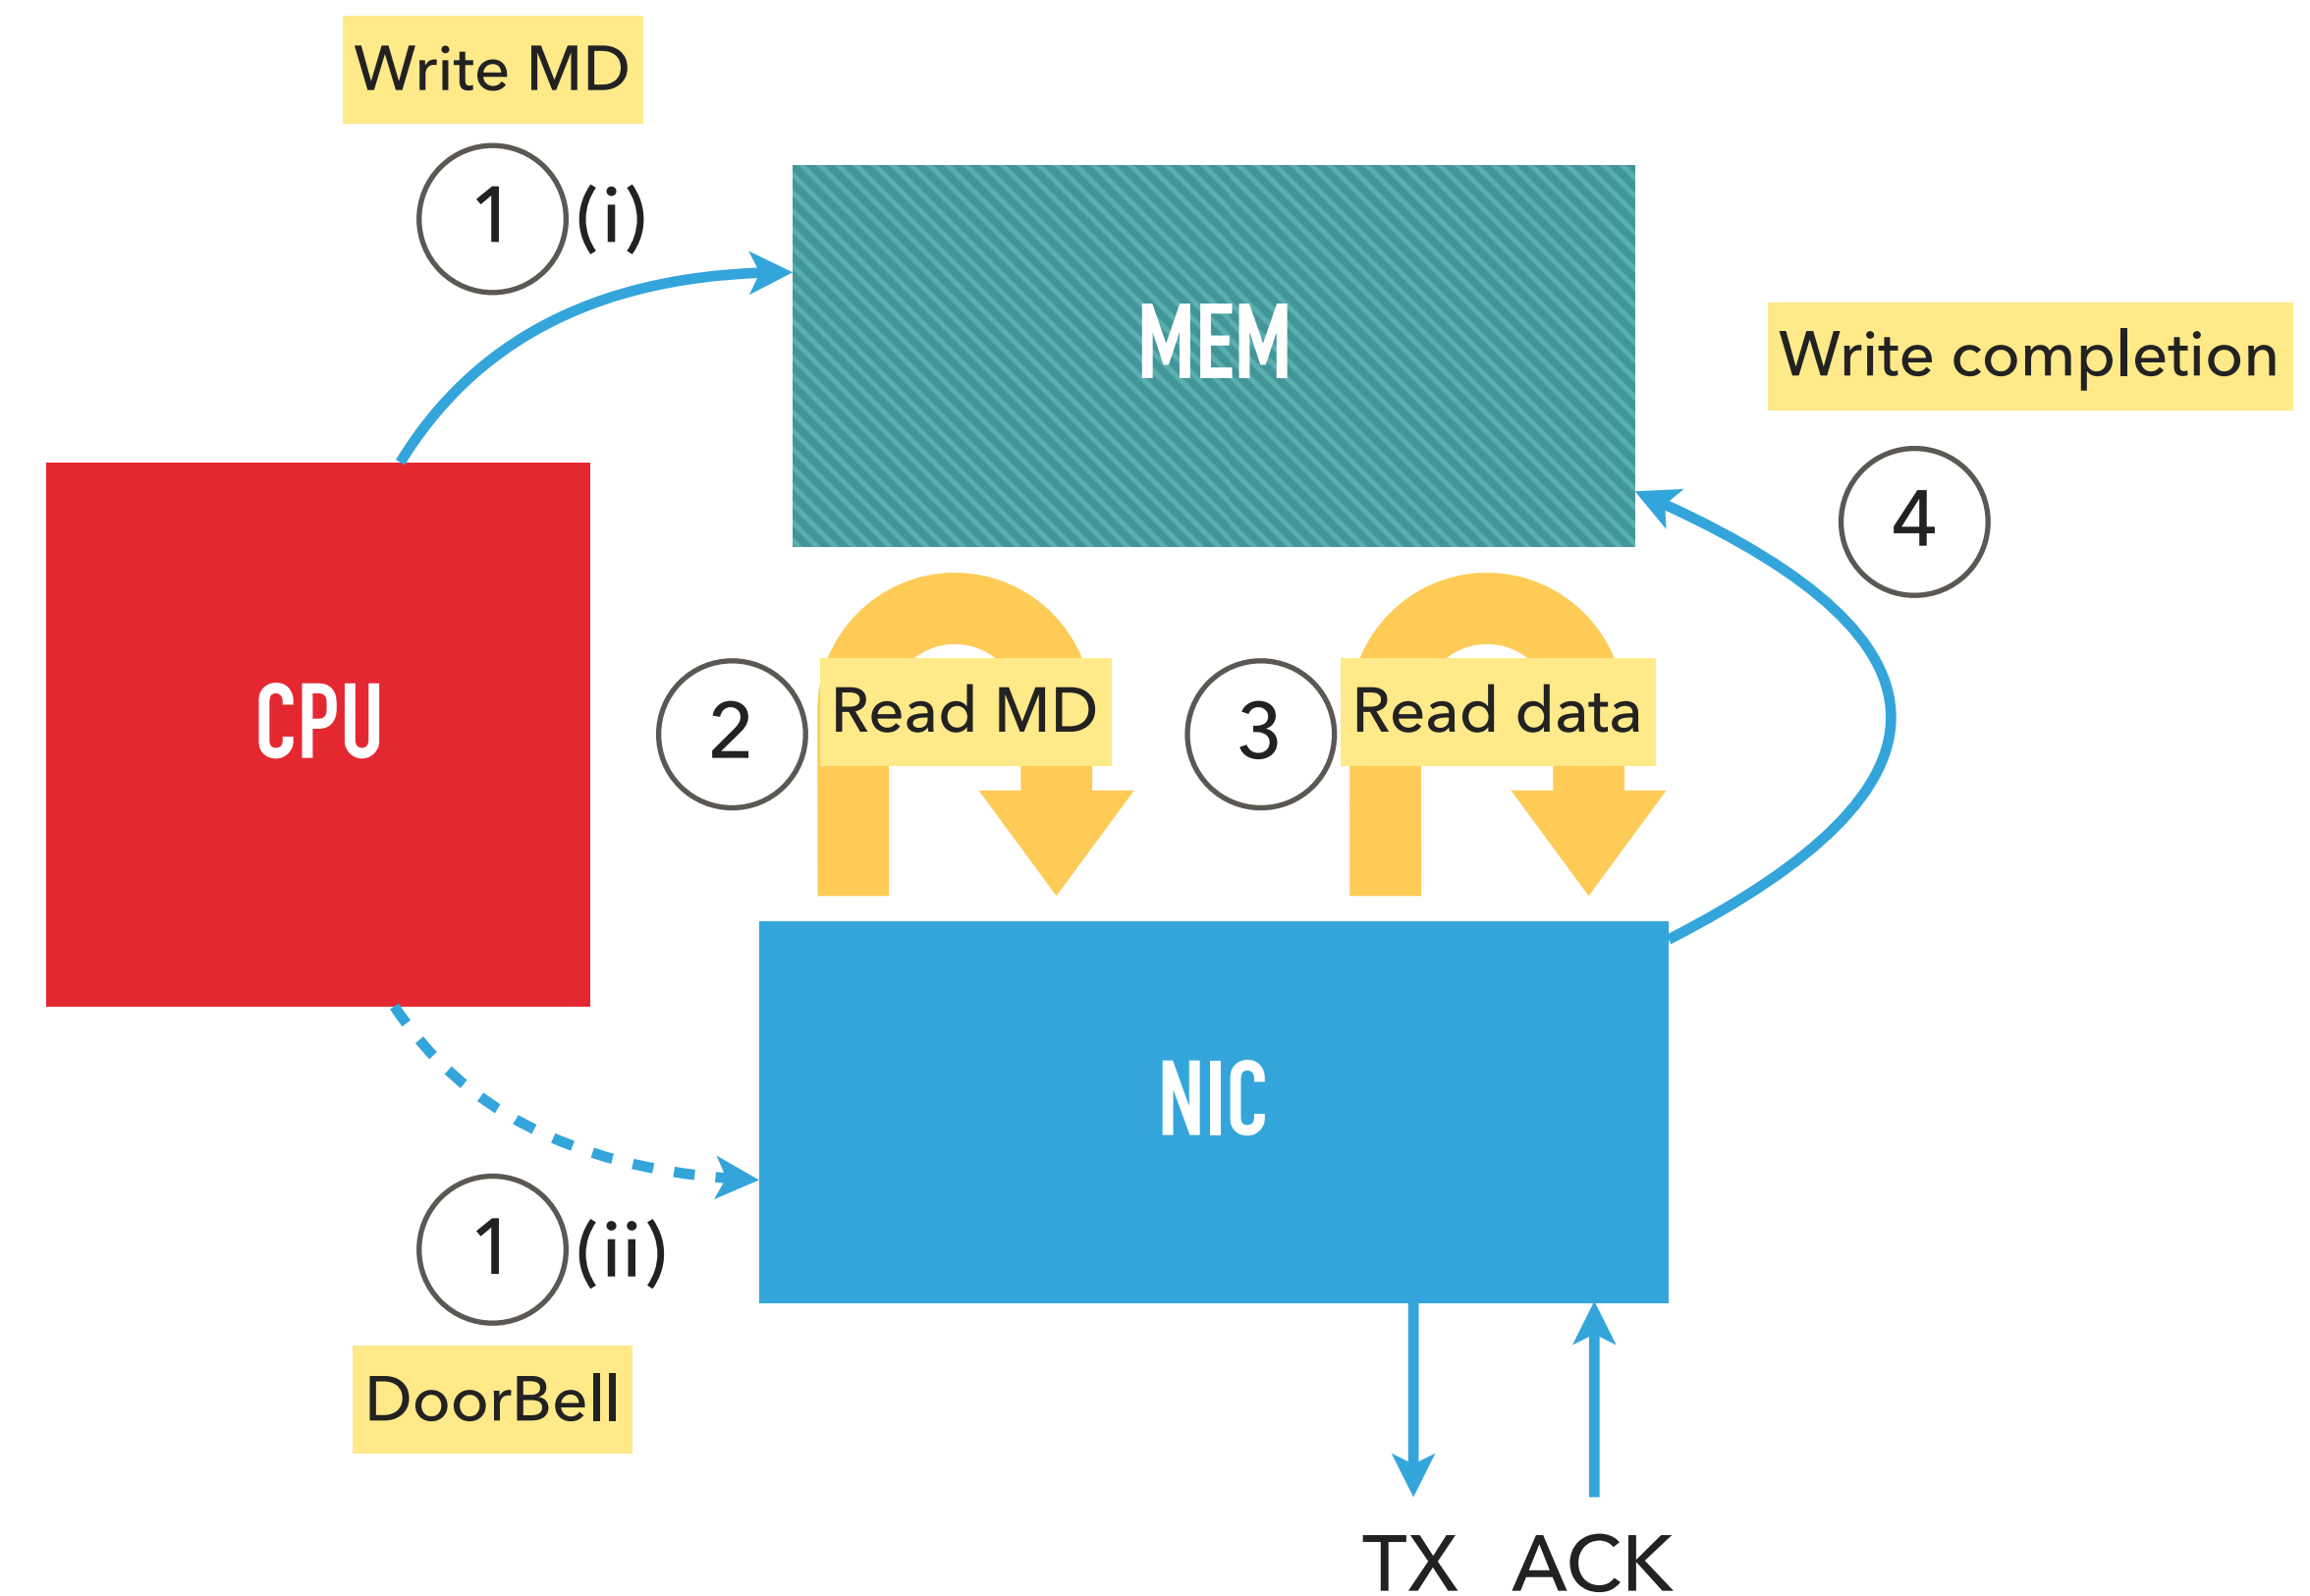
## CPU-NIC INTERACTION FEATURE: UNSIGNALED COMPLETIONS

- ▶ We can choose whether or not to generate a completion for each message. If not, then the NIC will not DMA-write a completion for a successfully executed message request.
- ▶ Reduces the number of (4)



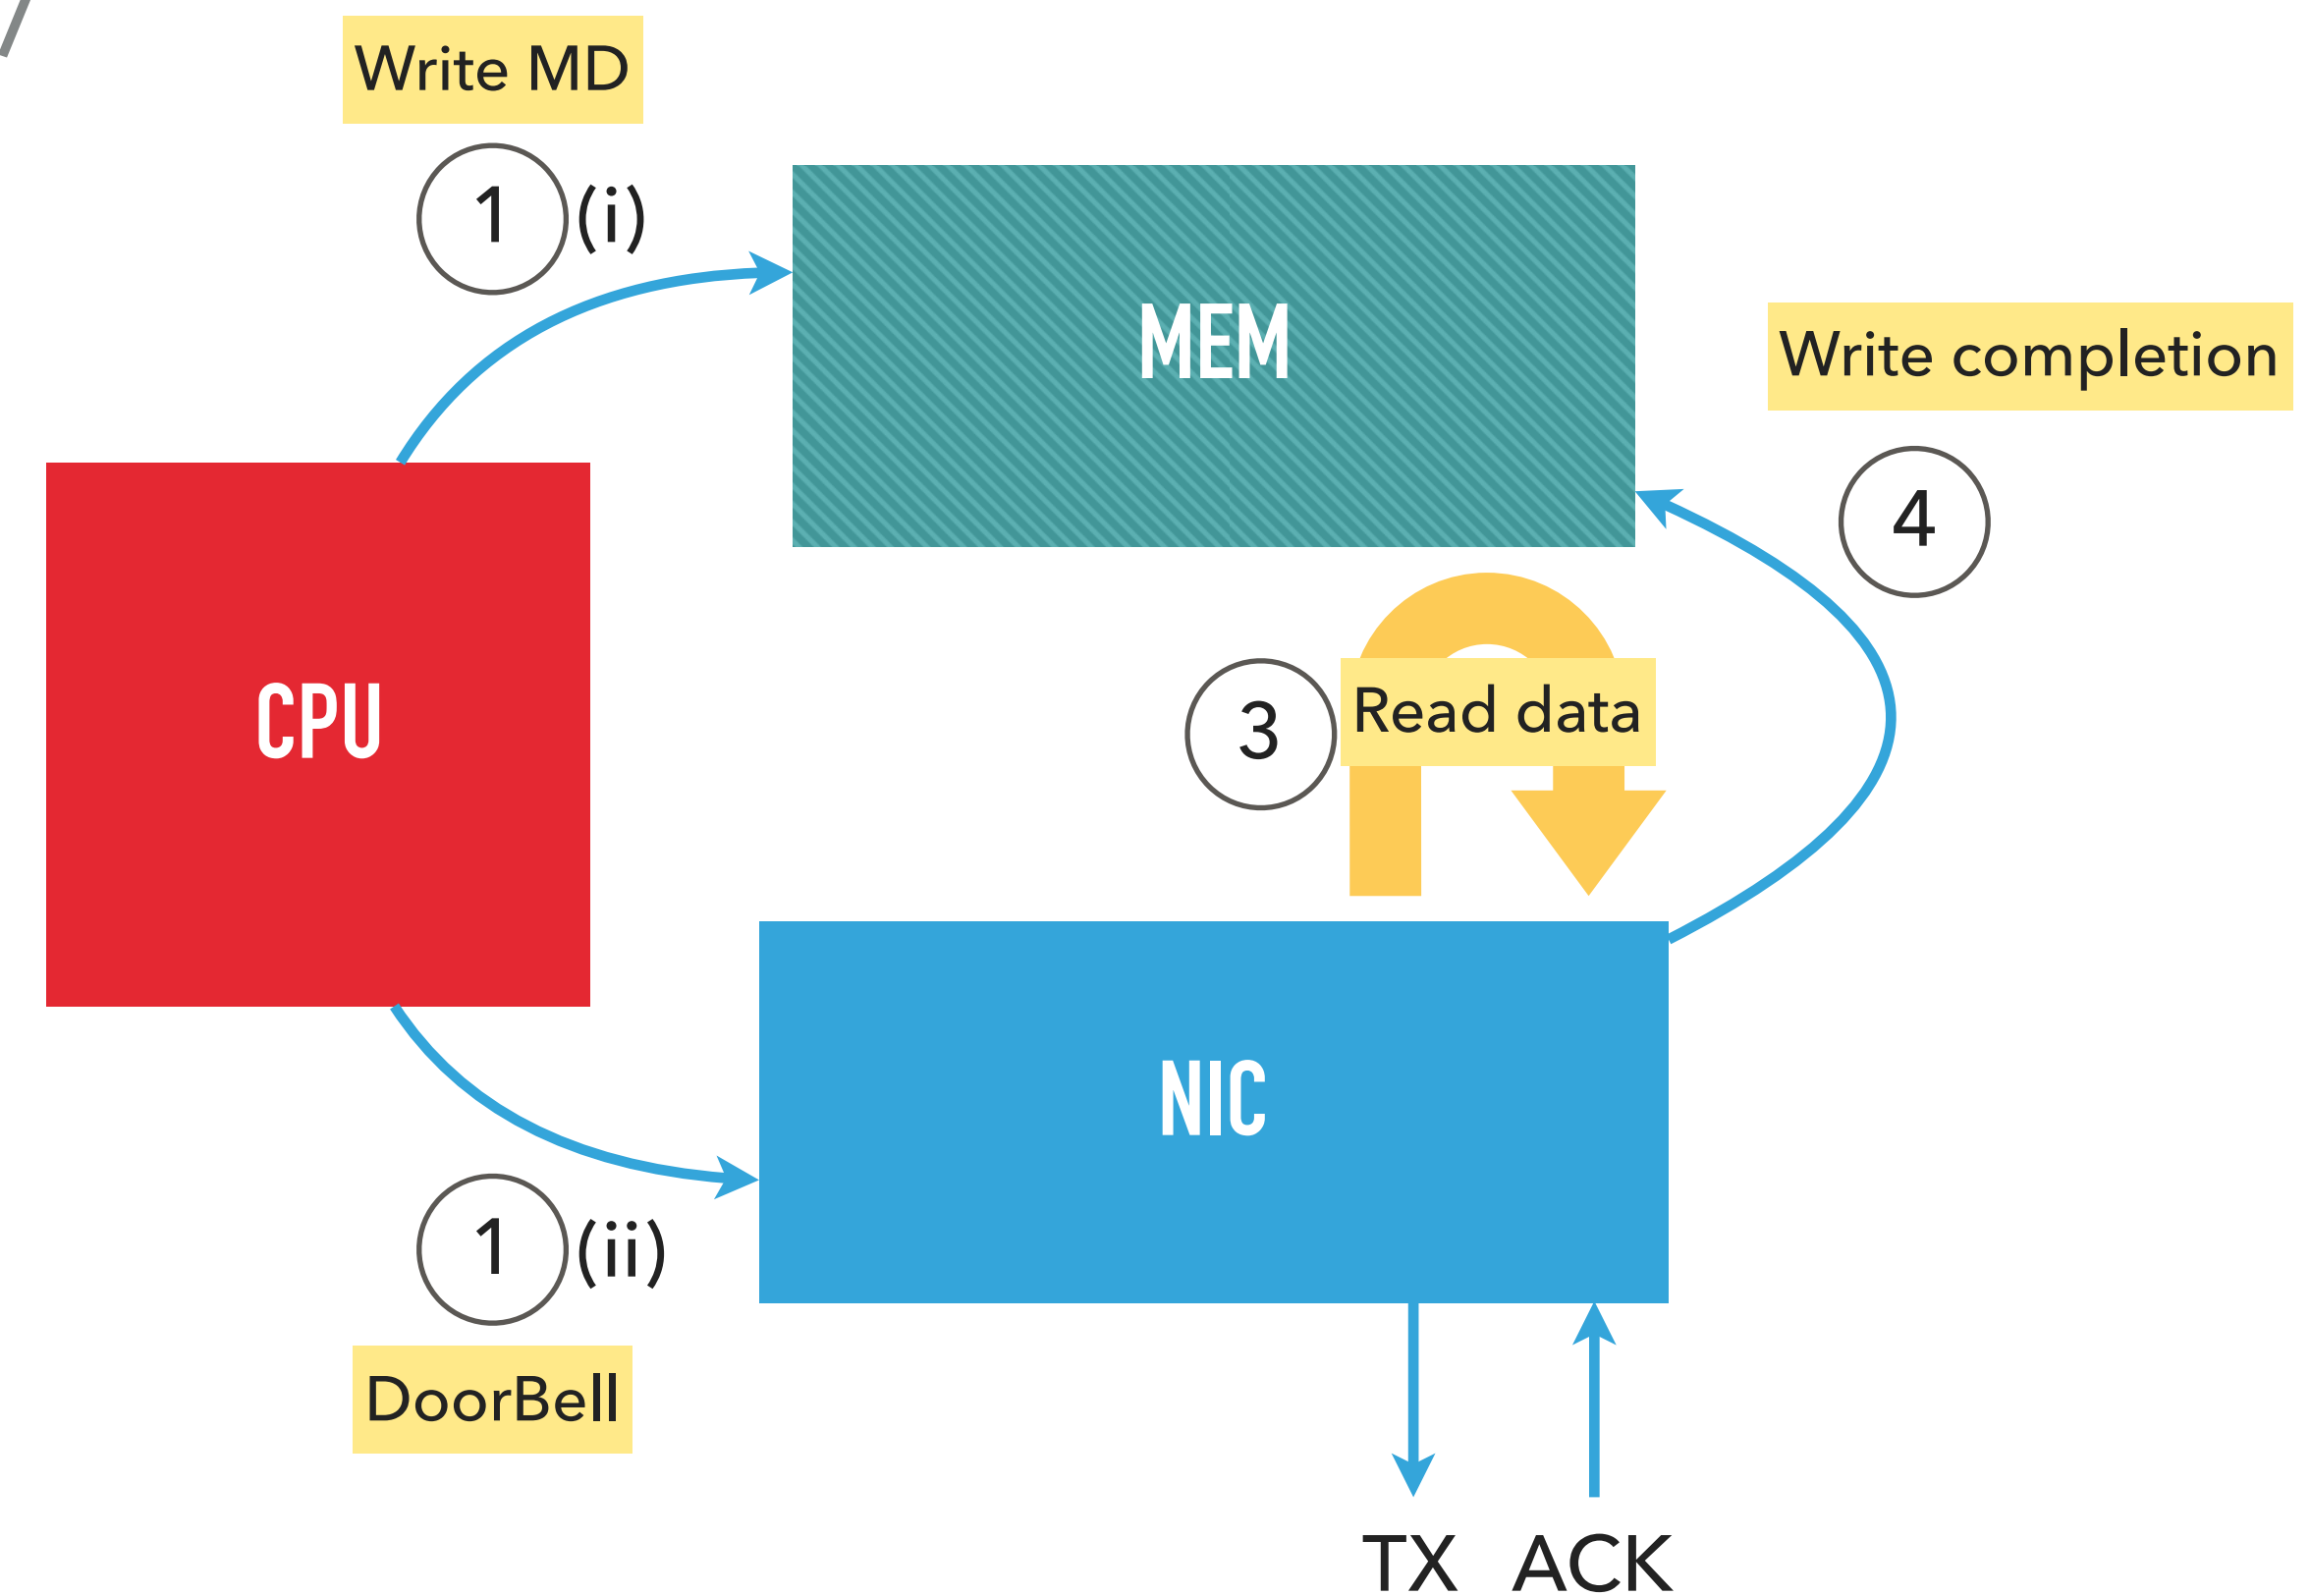
## CPU-NIC INTERACTION FEATURE: POSTLIST

- ▶ Instead of sending one message every DoorBell, we can send a linked list of message descriptors with 1 DoorBell
- ▶ Reduces number of (1)(ii)



## CPU-NIC INTERACTION FEATURE: BLUEFLAME

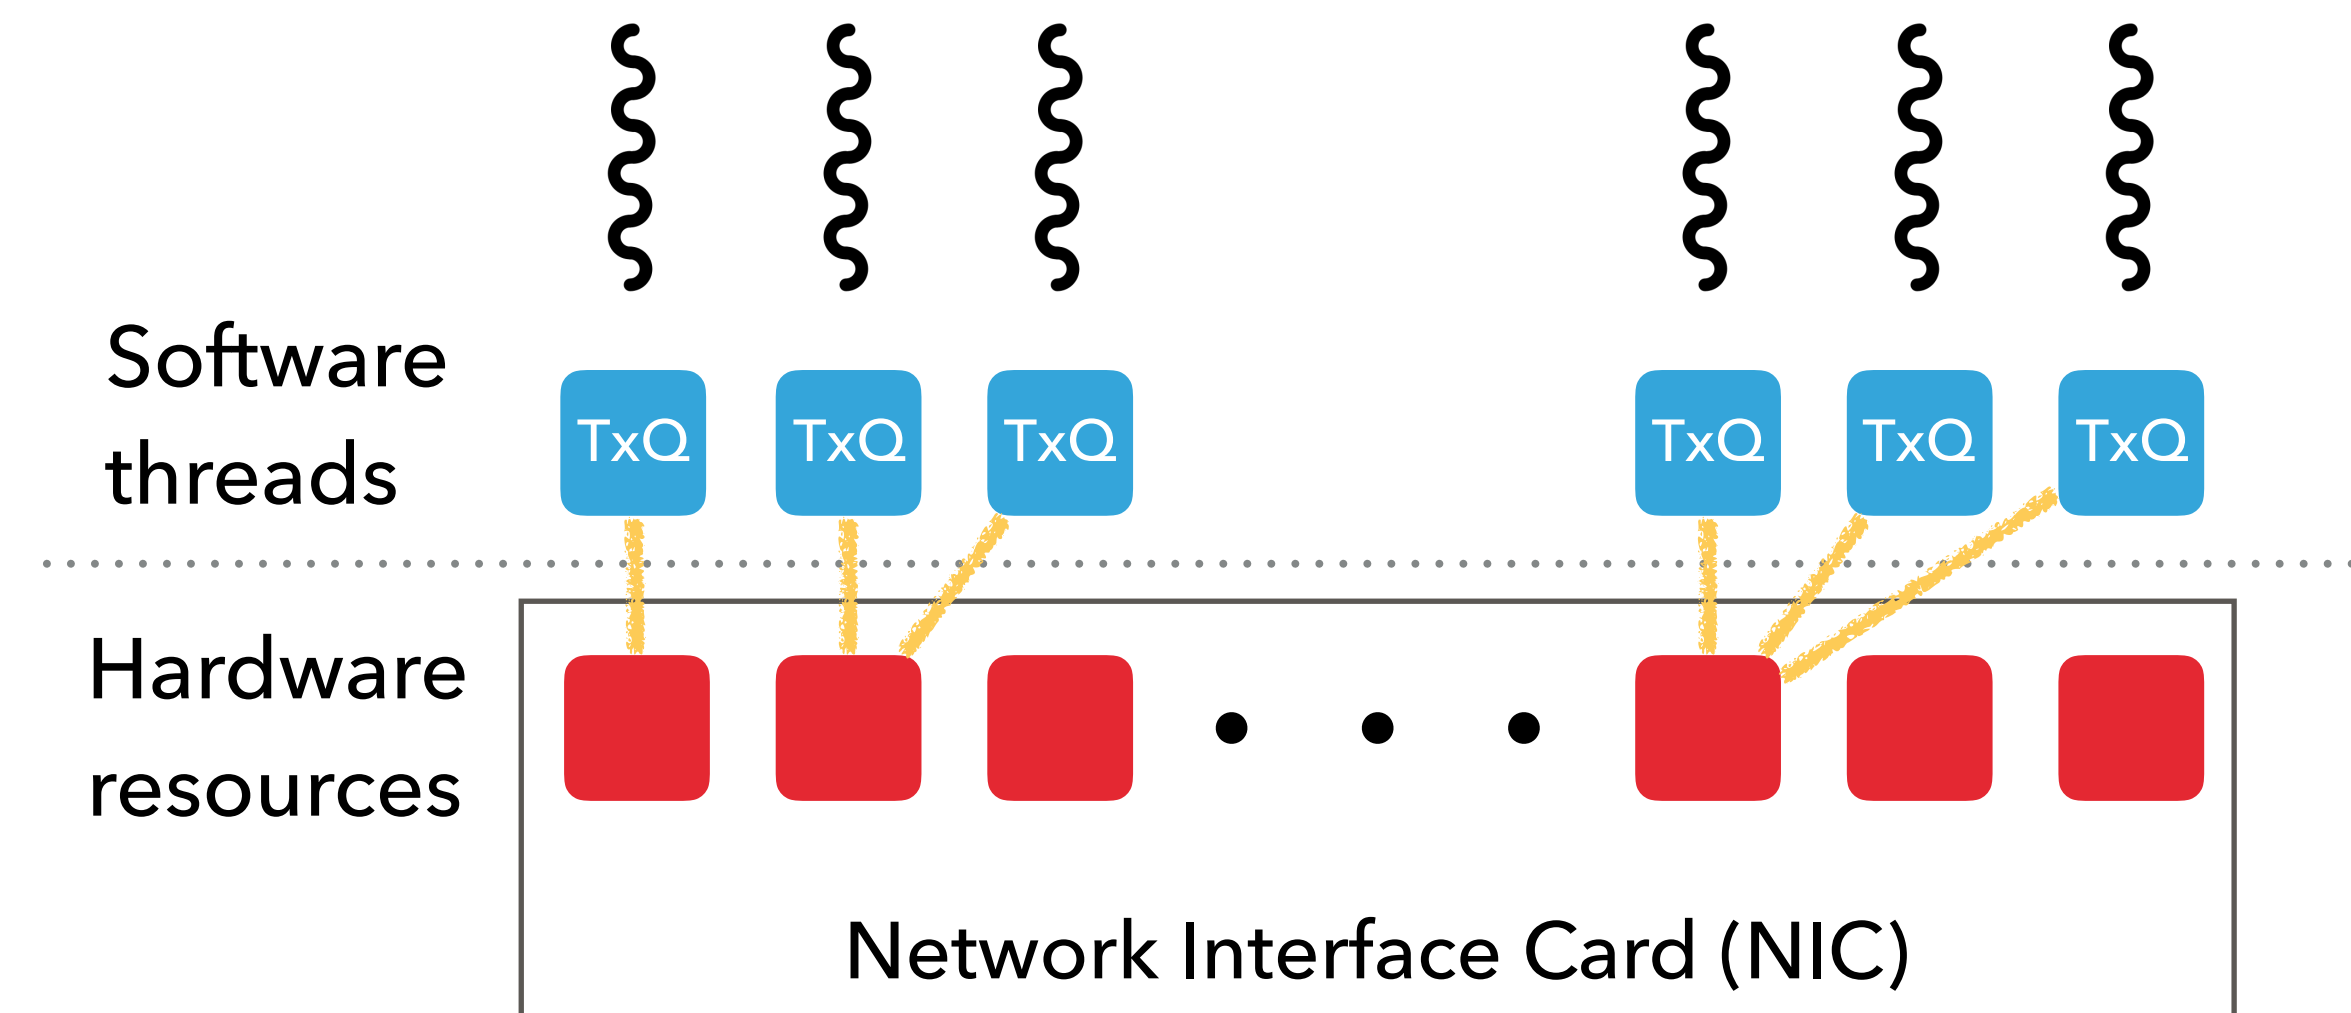
- ▶ Mellanox terminology for Programmed I/O (PIO)
- ▶ Send message descriptor as a part of DoorBell
- ▶ If used with inlining, will also cut off (3)
- ▶ Meaningless if used with Postlist





## INITIATION INTERFACE

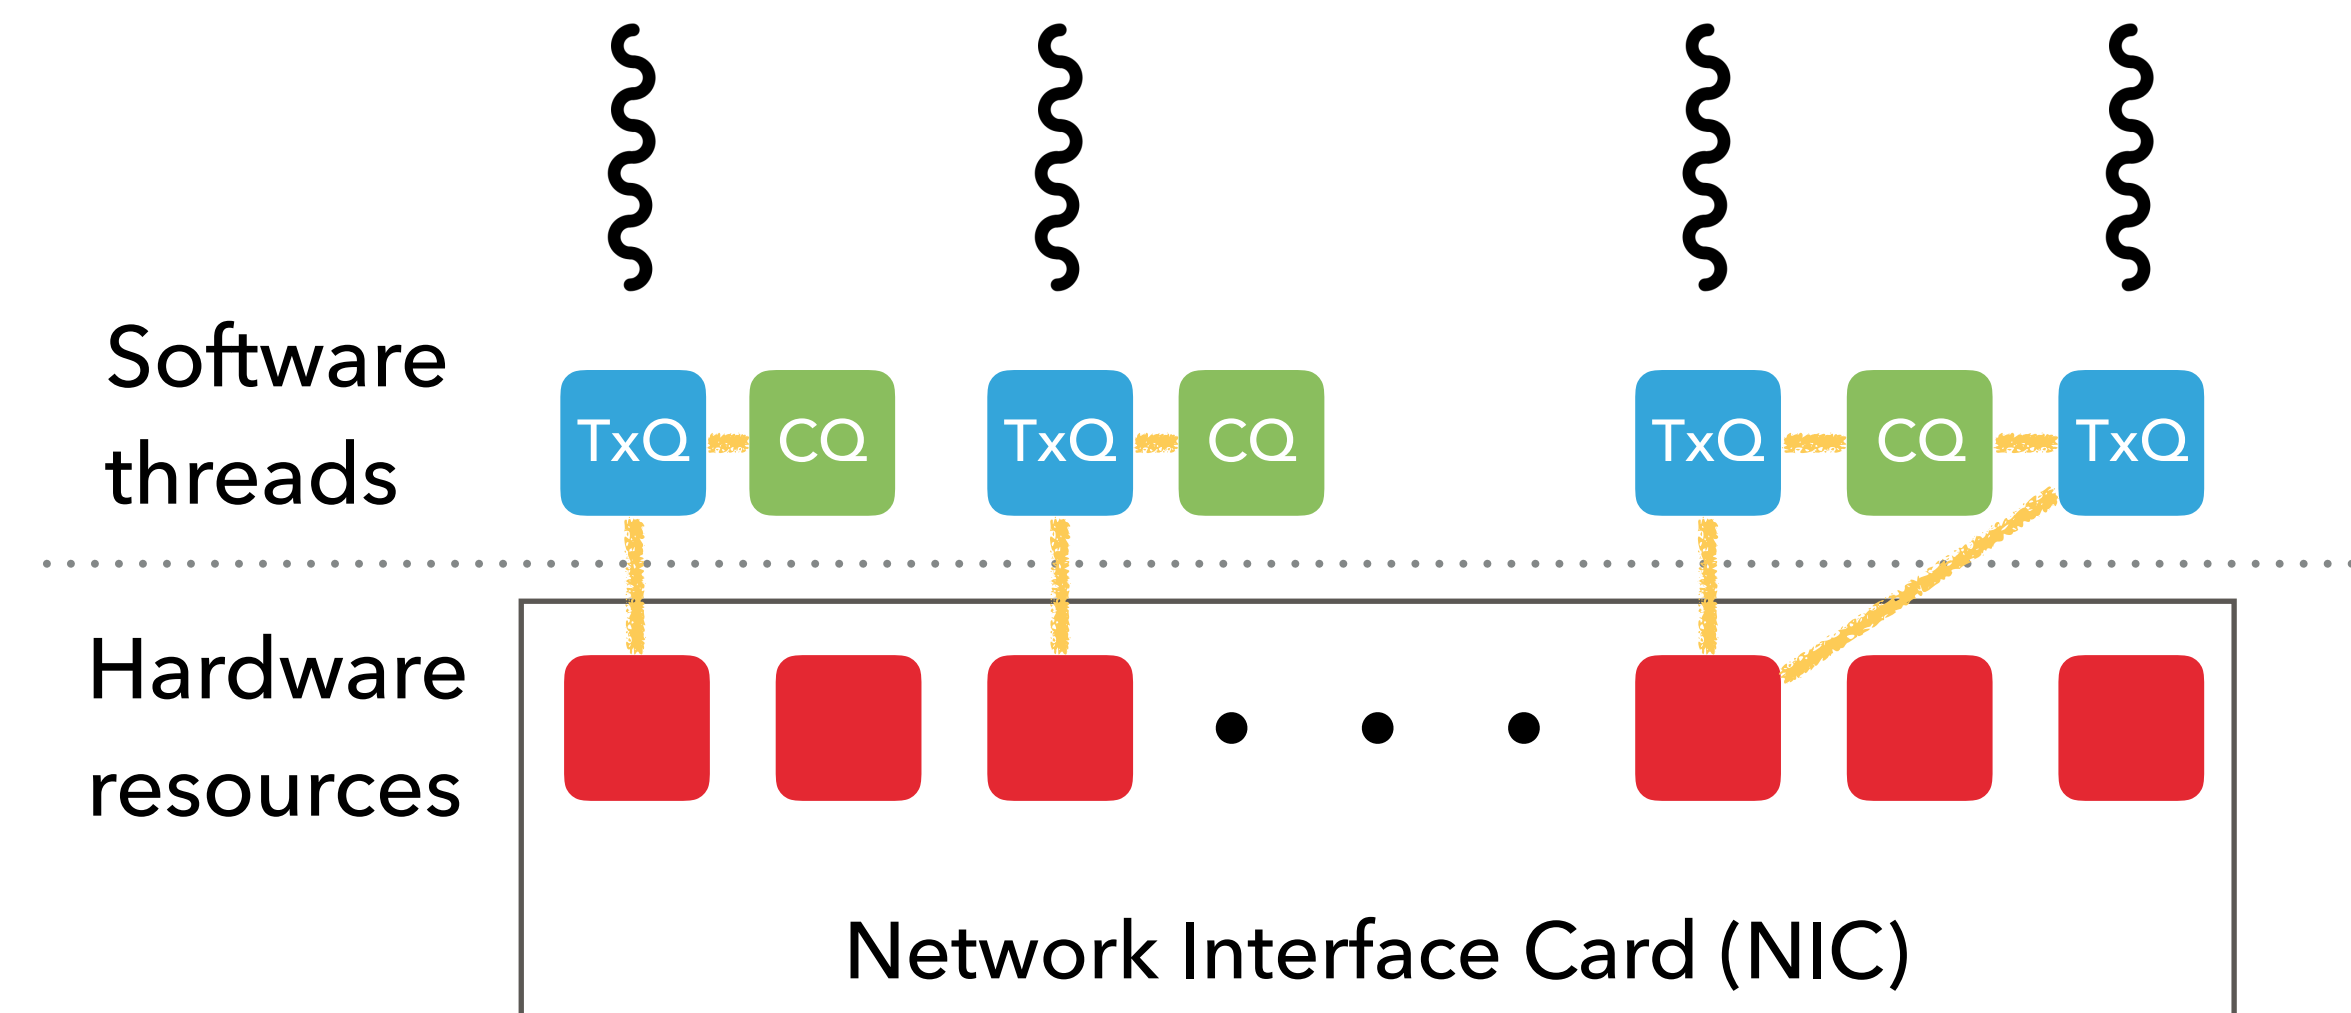
- ▶ Software transmit queue (TxQ) is a FIFO queue of messages.
- ▶ High-performance NICs feature multiple hardware resources (for scalability).
- ▶ The mapping between the TxQ and the hardware resource depends on the interconnect's driver. Hardware resource is `mmap()`ed.
- ▶ The thread posts messages to the endpoint, driver MMIO-writes to hardware resource to signal the NIC to send the message





## COMPLETION INTERFACE

- ▶ A FIFO completion queue (CQ) contains completions that indicate a successful, or unsuccessful completion of a message request.
- ▶ Mapping of CQs to the TxQ depends on software design, not the interconnect's driver.
- ▶ The NIC DMA-writes the completion, signaling the thread about the status of its message request.

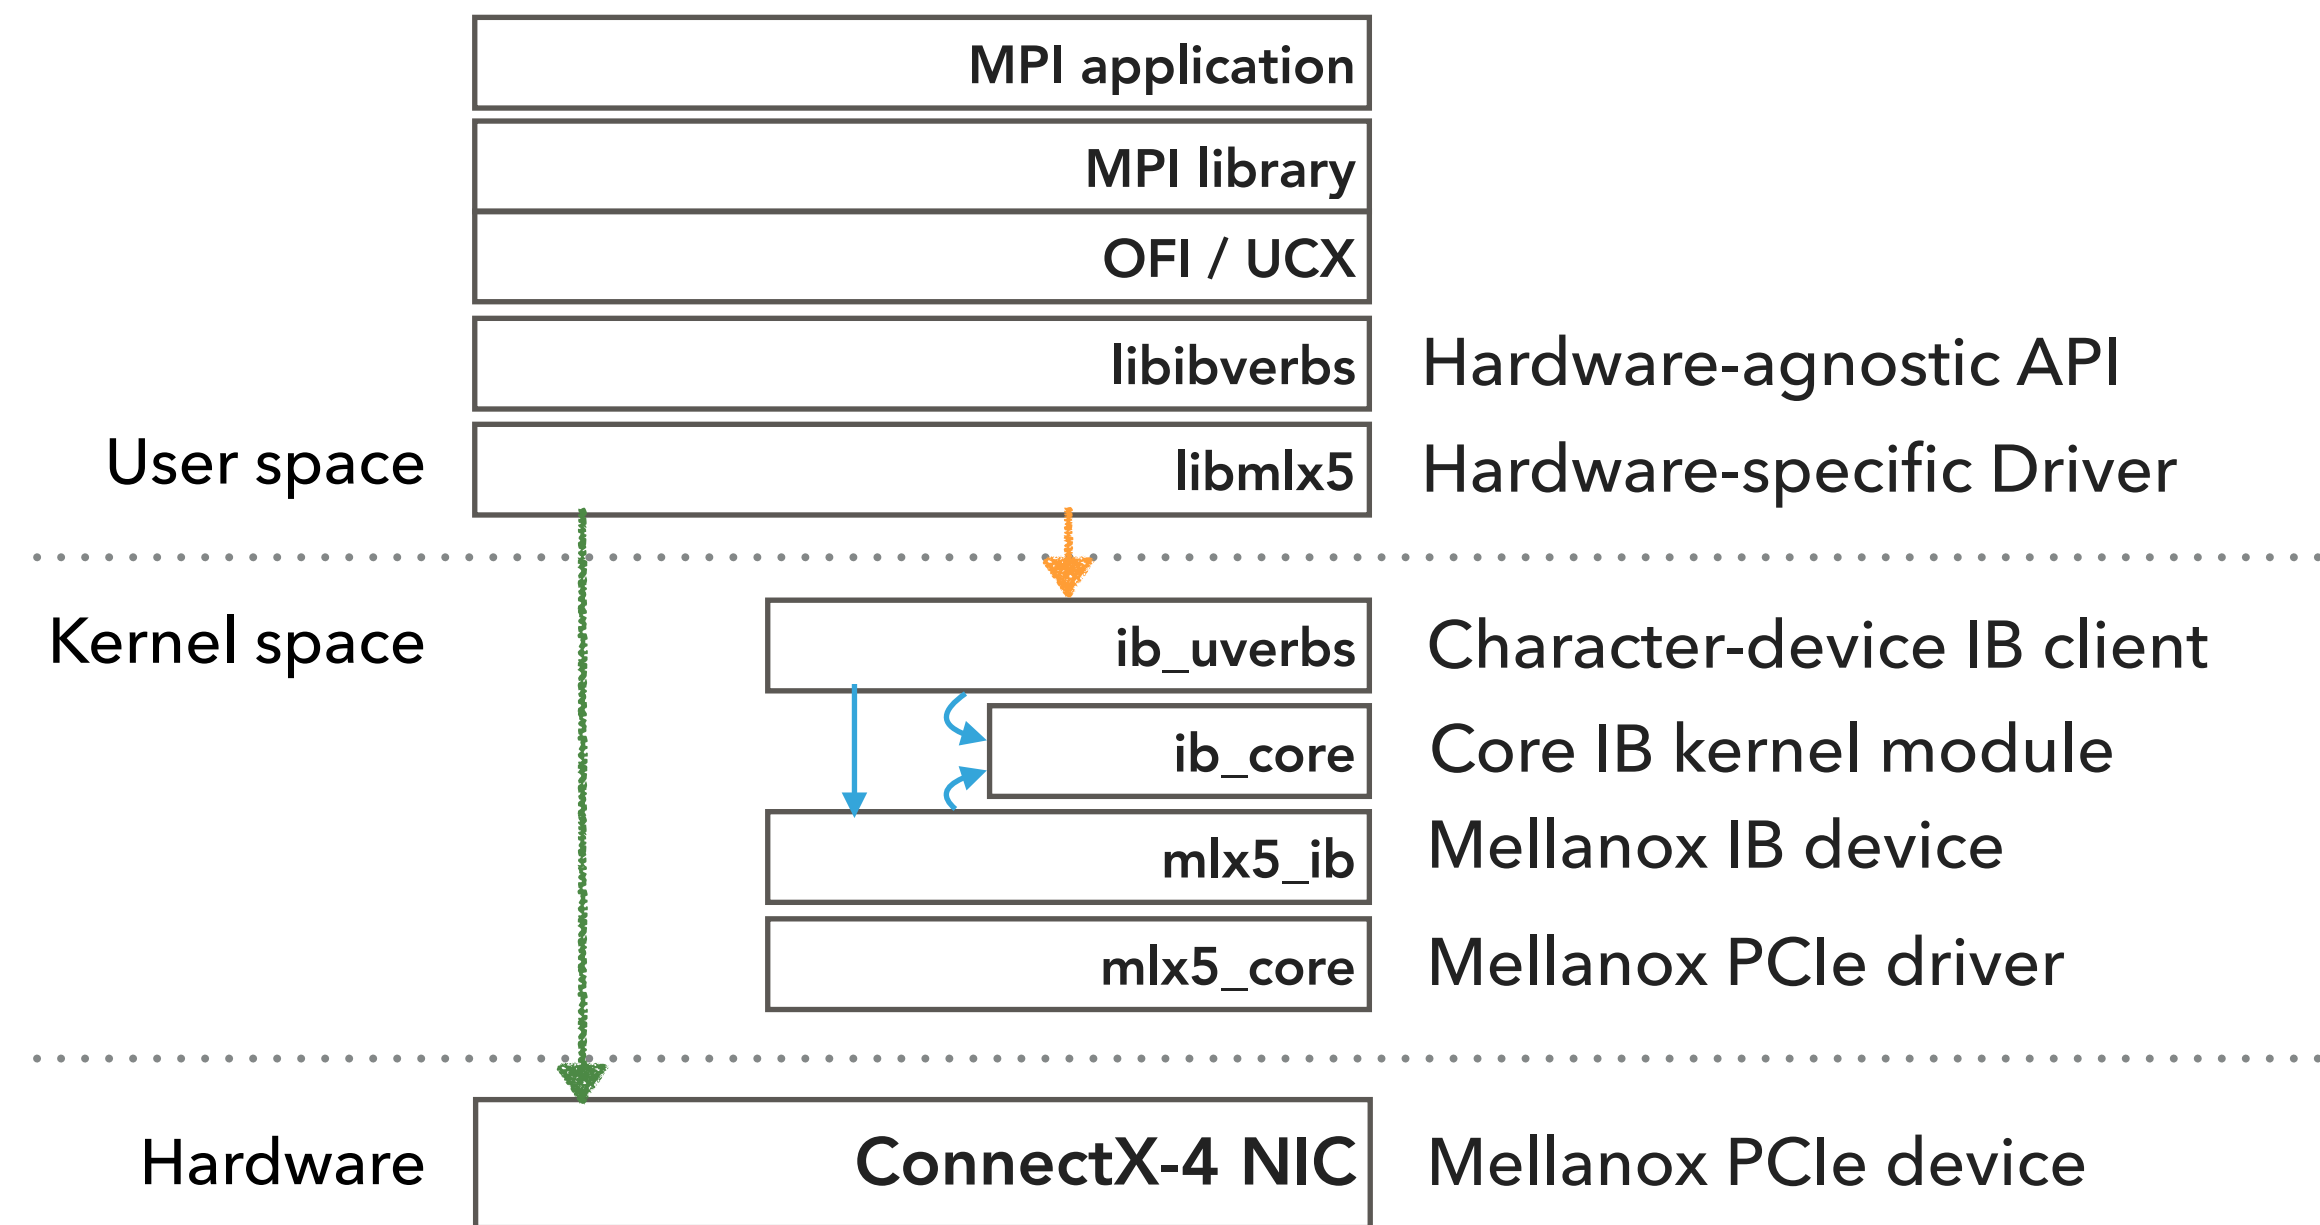


# COMMUNICATION RESOURCES: TXQ + CQ + HARDWARE RESOURCES



- ▶ Traditionally, the TxQs and CQs are exposed only to the kernel. For example, `send()` or `recv()` in TCP/IP's `sockets` interface.
- ▶ High-performance interconnects expose TxQs and CQs to the user-space to prevent system call overheads (kernel by-pass).
- ▶ Impact **memory usage** and **hardware resource** consumption

# HIGH-PERFORMANCE INTERCONNECT OF CHOICE: MELLANOX INFINIBAND




- ▶ **Slow-path:** user code interacts with kernel-code by `write()` on `/dev/infiniband/verbs0`. Slow-path contains resource setup operations.
- ▶ **Fast-path:** user code writes to NIC's hardware resources

# INFINIBAND COMMUNICATION RESOURCES

Transmit queue: Queue Pair (QP)

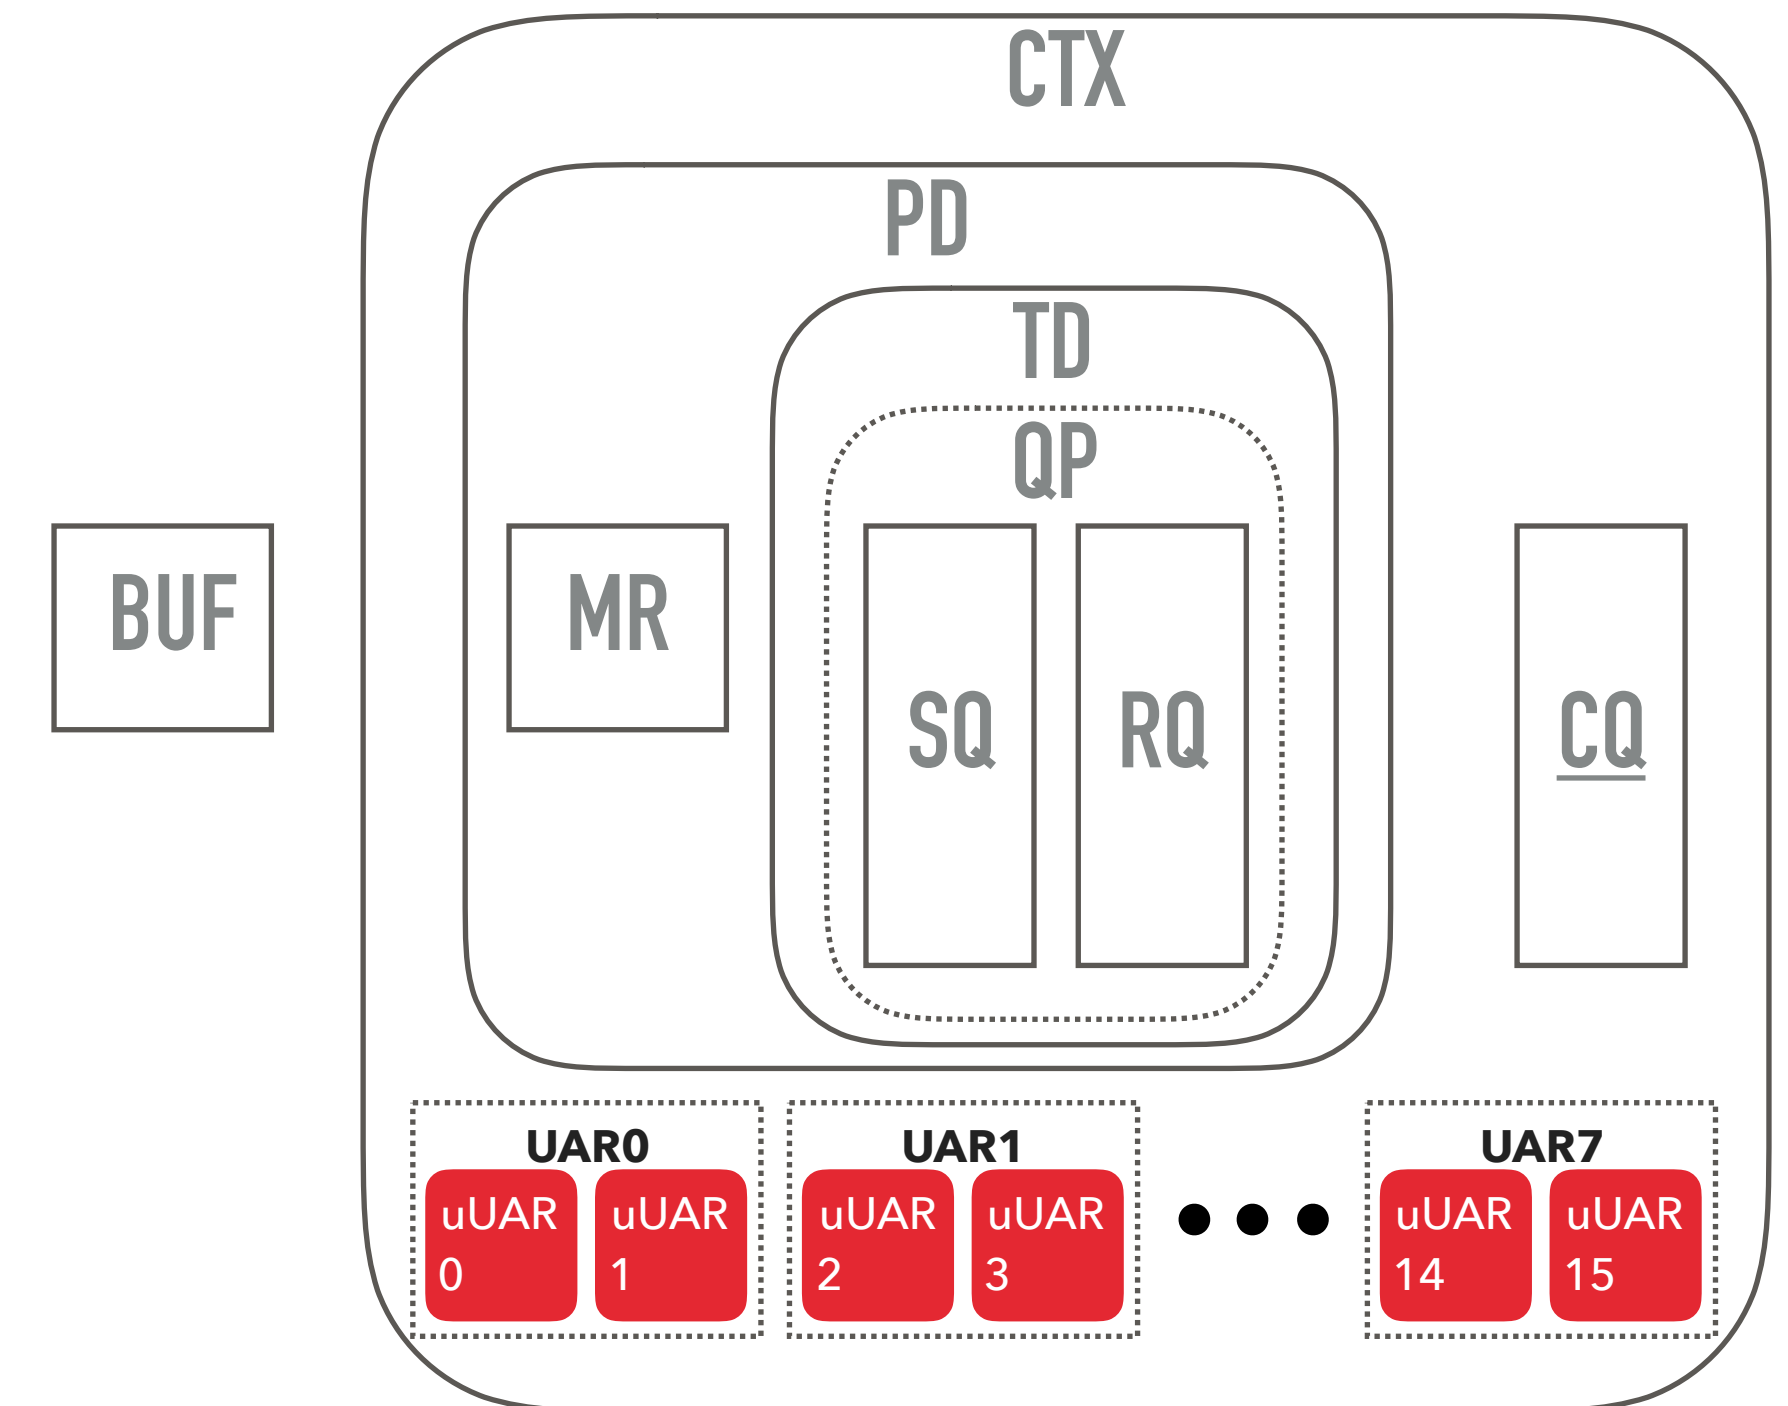
Completion queue: Completion  
Queue (CQ)

Hardware resources: User Access  
Region (UAR) + micro UARs (uUARs)



# INFINIBAND API

- ▶ **1 Device Context** (`ibv_open_device`): 8 UARs; 16 uUARs
- ▶ **1 Protection Domain** (`ibv_alloc_pd`)
- ▶ **1 Memory Region** (`ibv_create_mr`)
- ▶ **1 Completion Queue** (`ibv_create_cq`)
- ▶ **Memory buffer** (`malloc`)
- ▶ **1 Queue Pair** (`ibv_create_qp`)
- ▶ **1 Thread Domain** (`ibv_alloc_td`): 1 UAR; 2 uUARs



### NAIVE ENDPOINTS

- ▶ To achieve maximum possible throughput, create  $n$  contexts with TD-assigned-QPs for  $n$  threads
  - ▶ 18 uUARs per thread: 8 static UARs + 1 dynamic UAR
- ▶ This naive approach impacts
  - ▶ Memory usage
  - ▶ Hardware resource usage

# MEMORY USAGE

- ▶ Each QP and CQ occupy memory with their circular buffers (FIFO queue)
- ▶ CTX uses 74% of the memory required for 1 endpoint
- ▶ Creating 16 endpoints, for example, with the naive approach will use 5.15 MB.
- ▶ Not of immediate concern since the memory on nodes of clusters and supercomputers is in the order of GB.

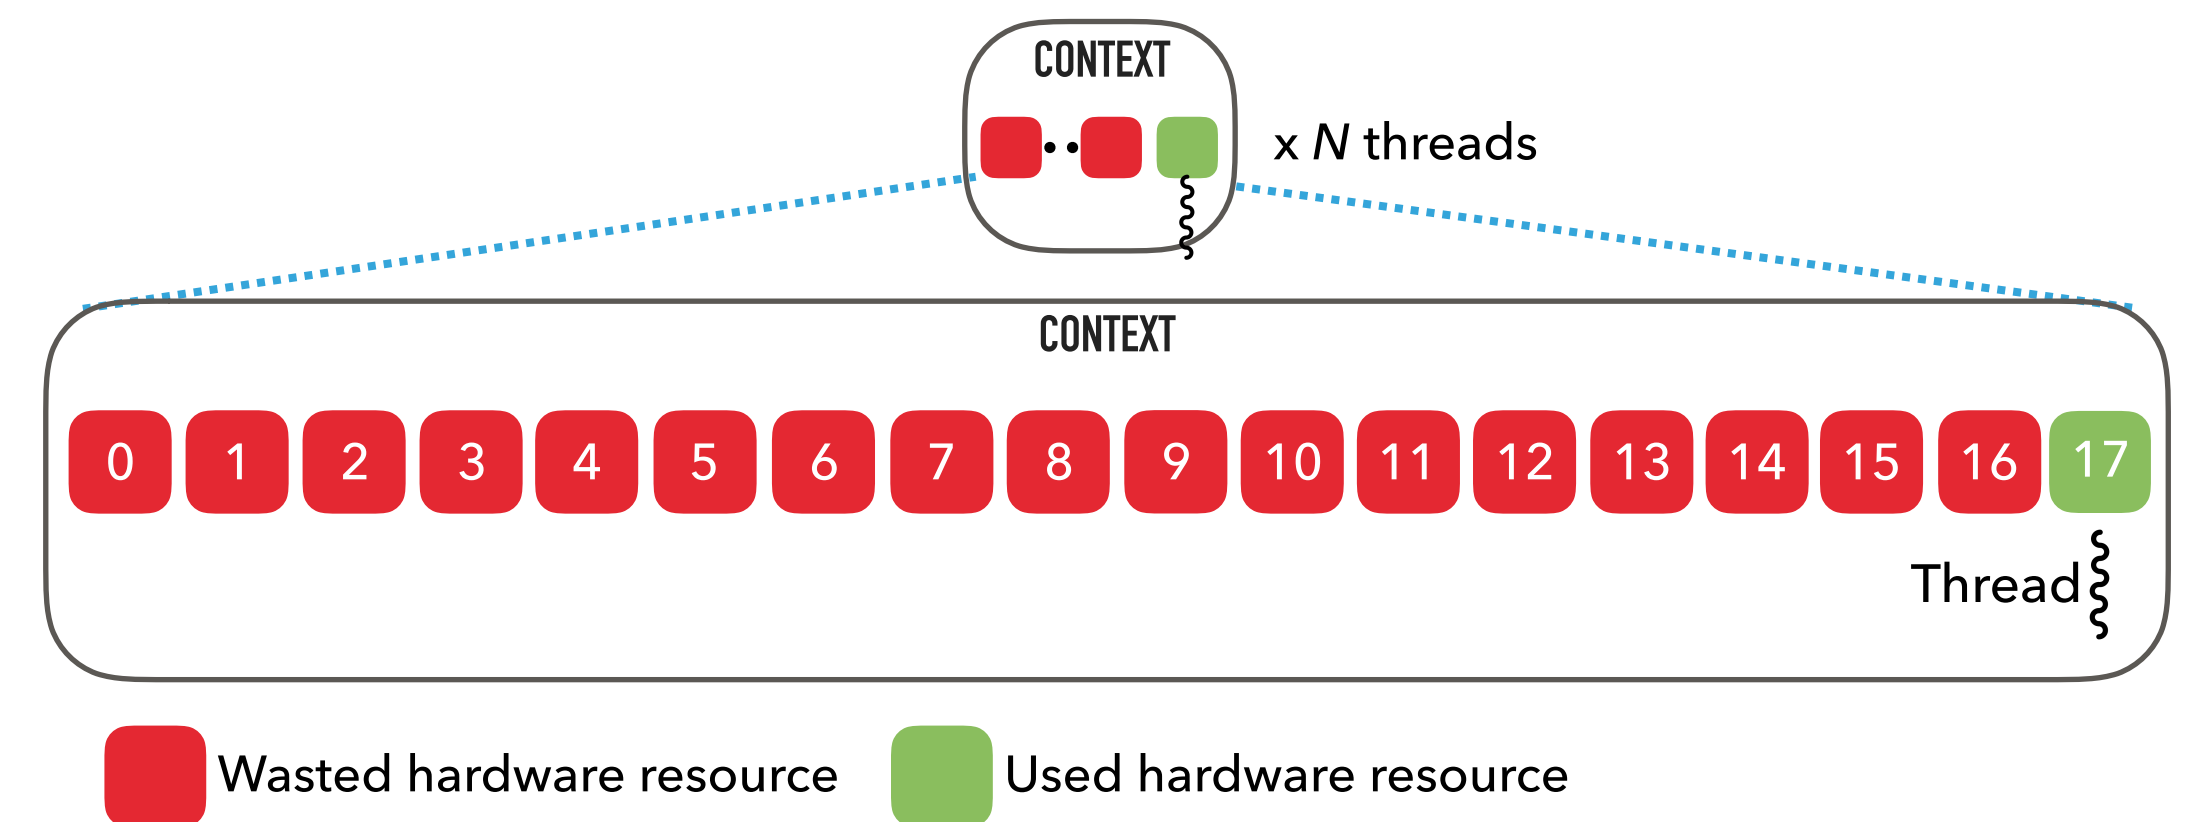
TABLE I  
BYTES USED BY VERBS RESOURCES

CTXs	PDs	MRs	QPs	CQs	Total
256K	144	144	80K	9K	345K



## HARDWARE RESOURCE USAGE

- ▶ Total number of uUARs on the ConnectX-4 is 16,336.
  - ▶ Translates to a max of 907 CTXs since each CTX allocates 18 uUARs.
- ▶ Naive approach translates to only 1 uUAR being utilized every 18 uUARs.
  - ▶ Translates to 94% resource wastage.
  - ▶ Need a 2nd NIC only after using 6% on the 1st.
- ▶ Could've reduced cost and power usage of the device.




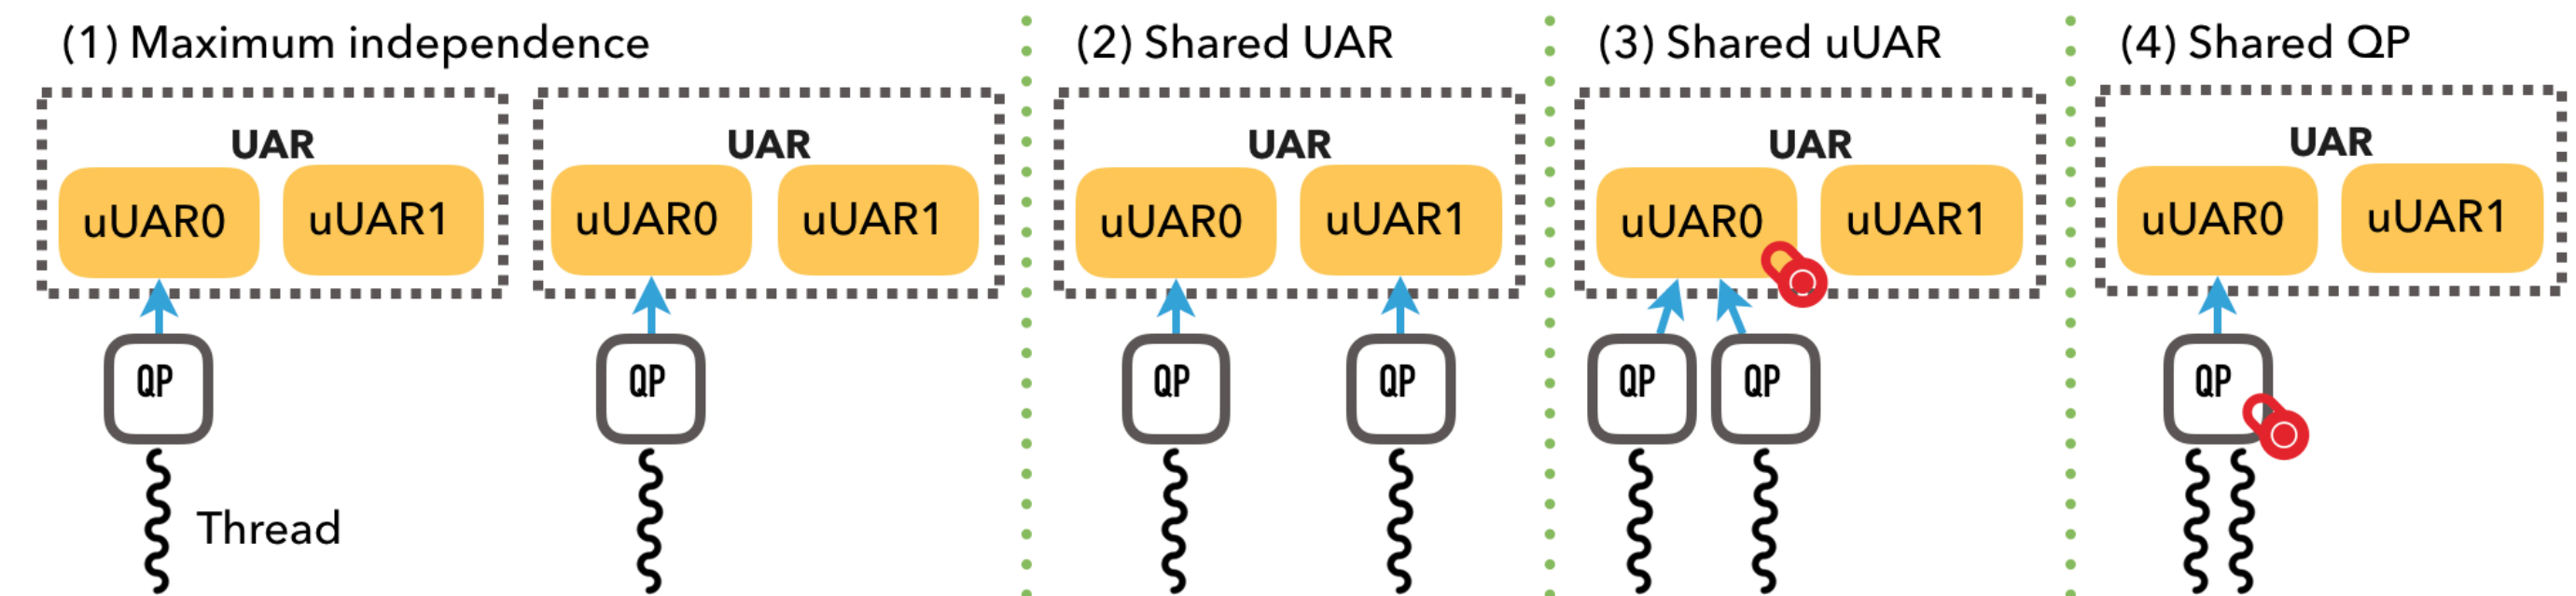


# COMMUNICATION RESOURCE SHARING

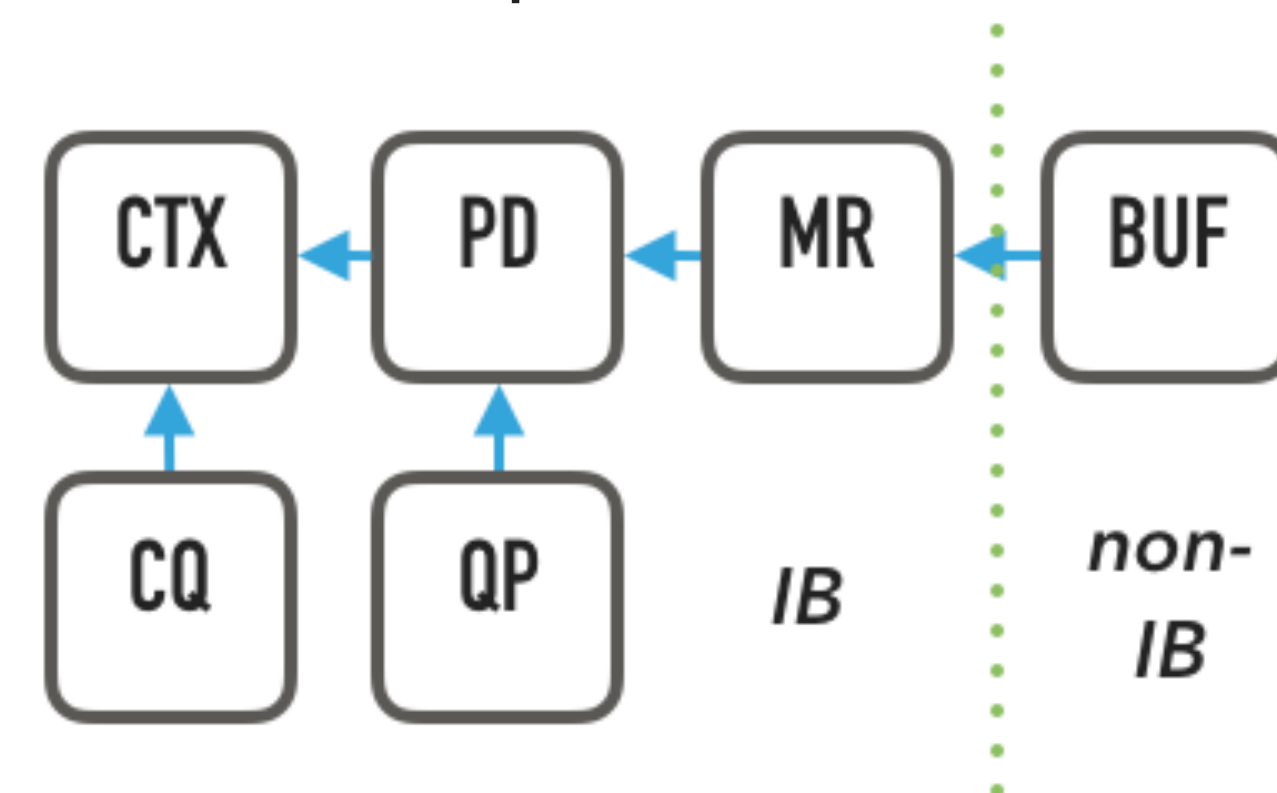
Analytically, 4 ways for initiation  
interface

Arbitrary number of ways for  
completion interface

Sharing analysis based on IB's  
resource hierarchy


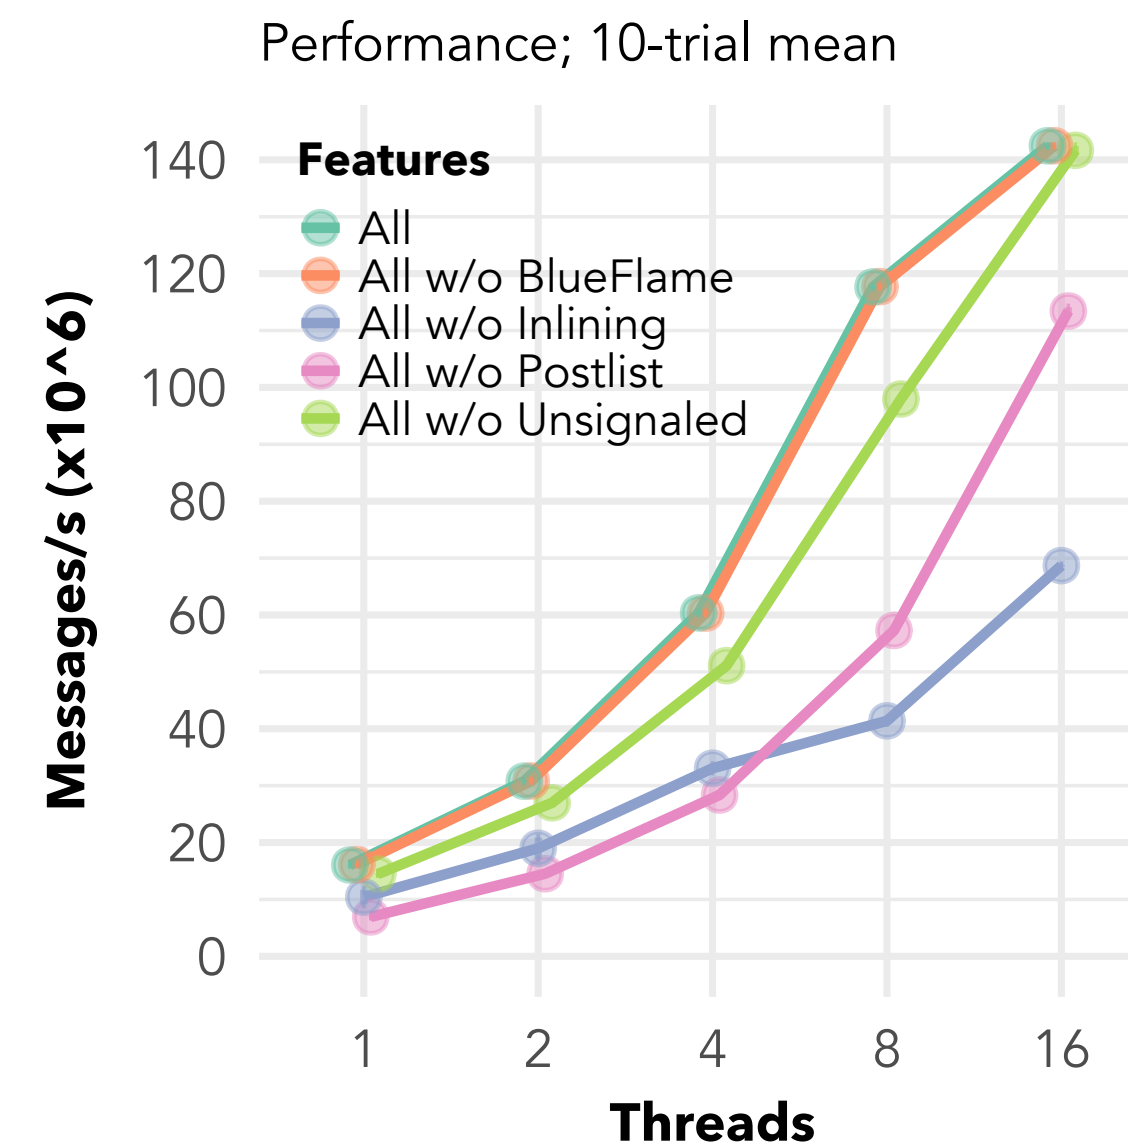


Arrow points to owner



## EVALUATION SETUP

- ▶ Multi-threaded 2-byte RDMA-write sender-receiver message rate benchmark
- ▶ Intel Haswell @ 2.5 GHz + ConnectX-4
- ▶ 16 threads
- ▶ QP depth: 64
- ▶ Postlist: 32; Unsignaled: 64
- ▶ To study effect of feature  $f$ : "All w/o  $f$ "

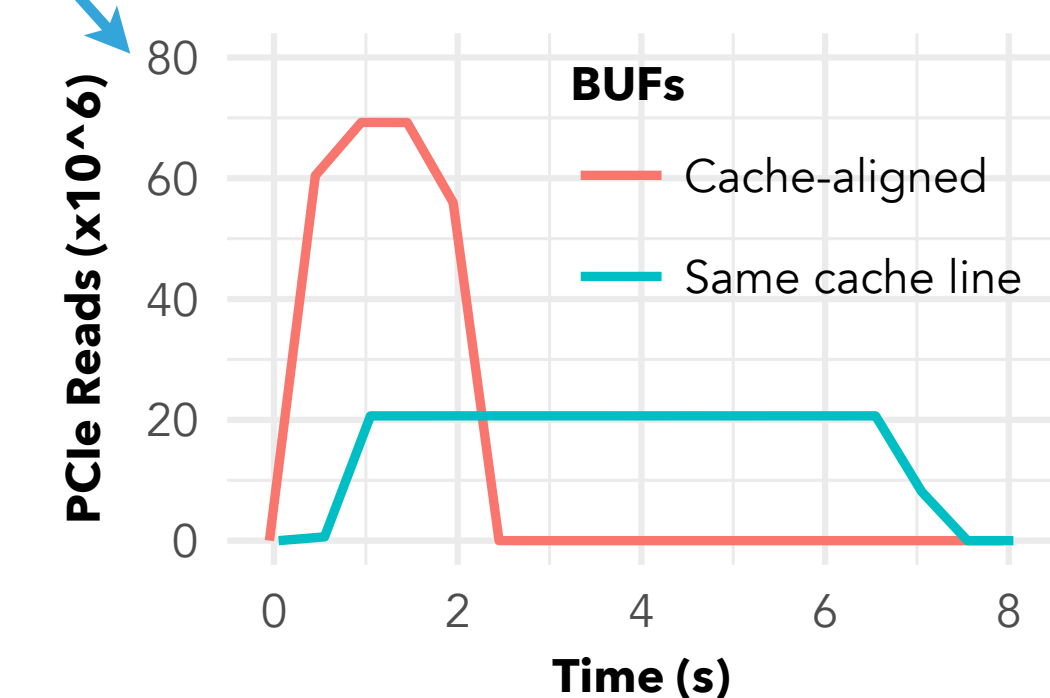
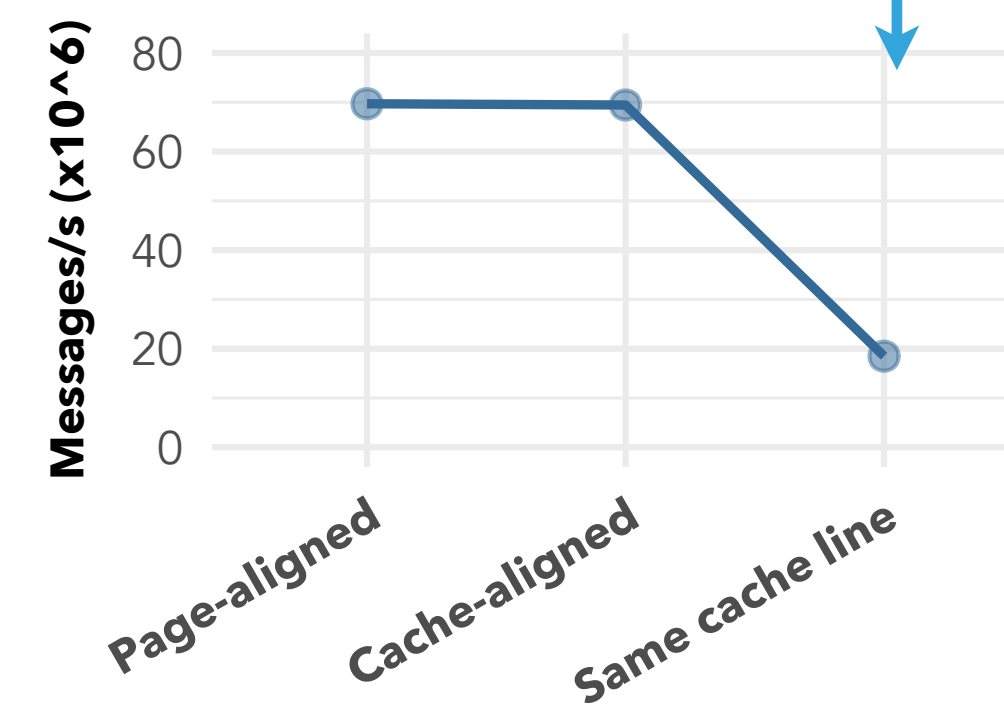
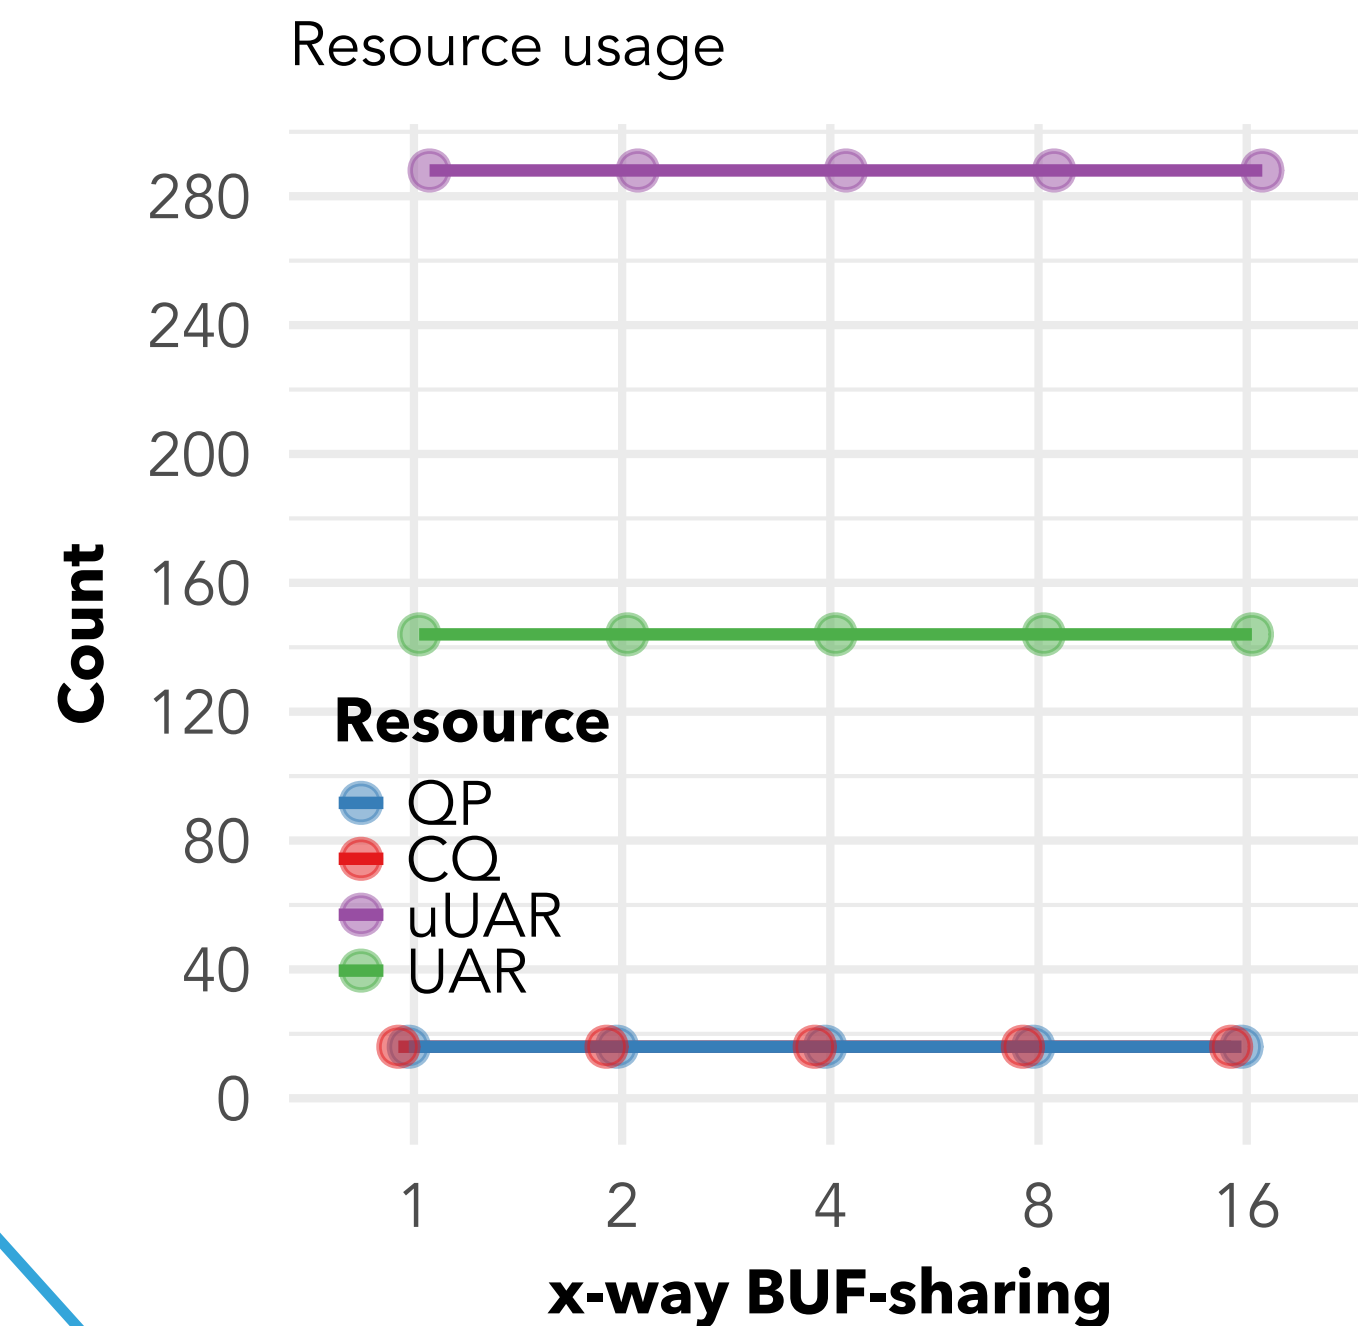



# BUFFER SHARING

Hash function of NIC's parallel TLB  
based on cache line.

Without inlining, parallel DMA reads  
to same address are serialized.



No impact on communication  
resources





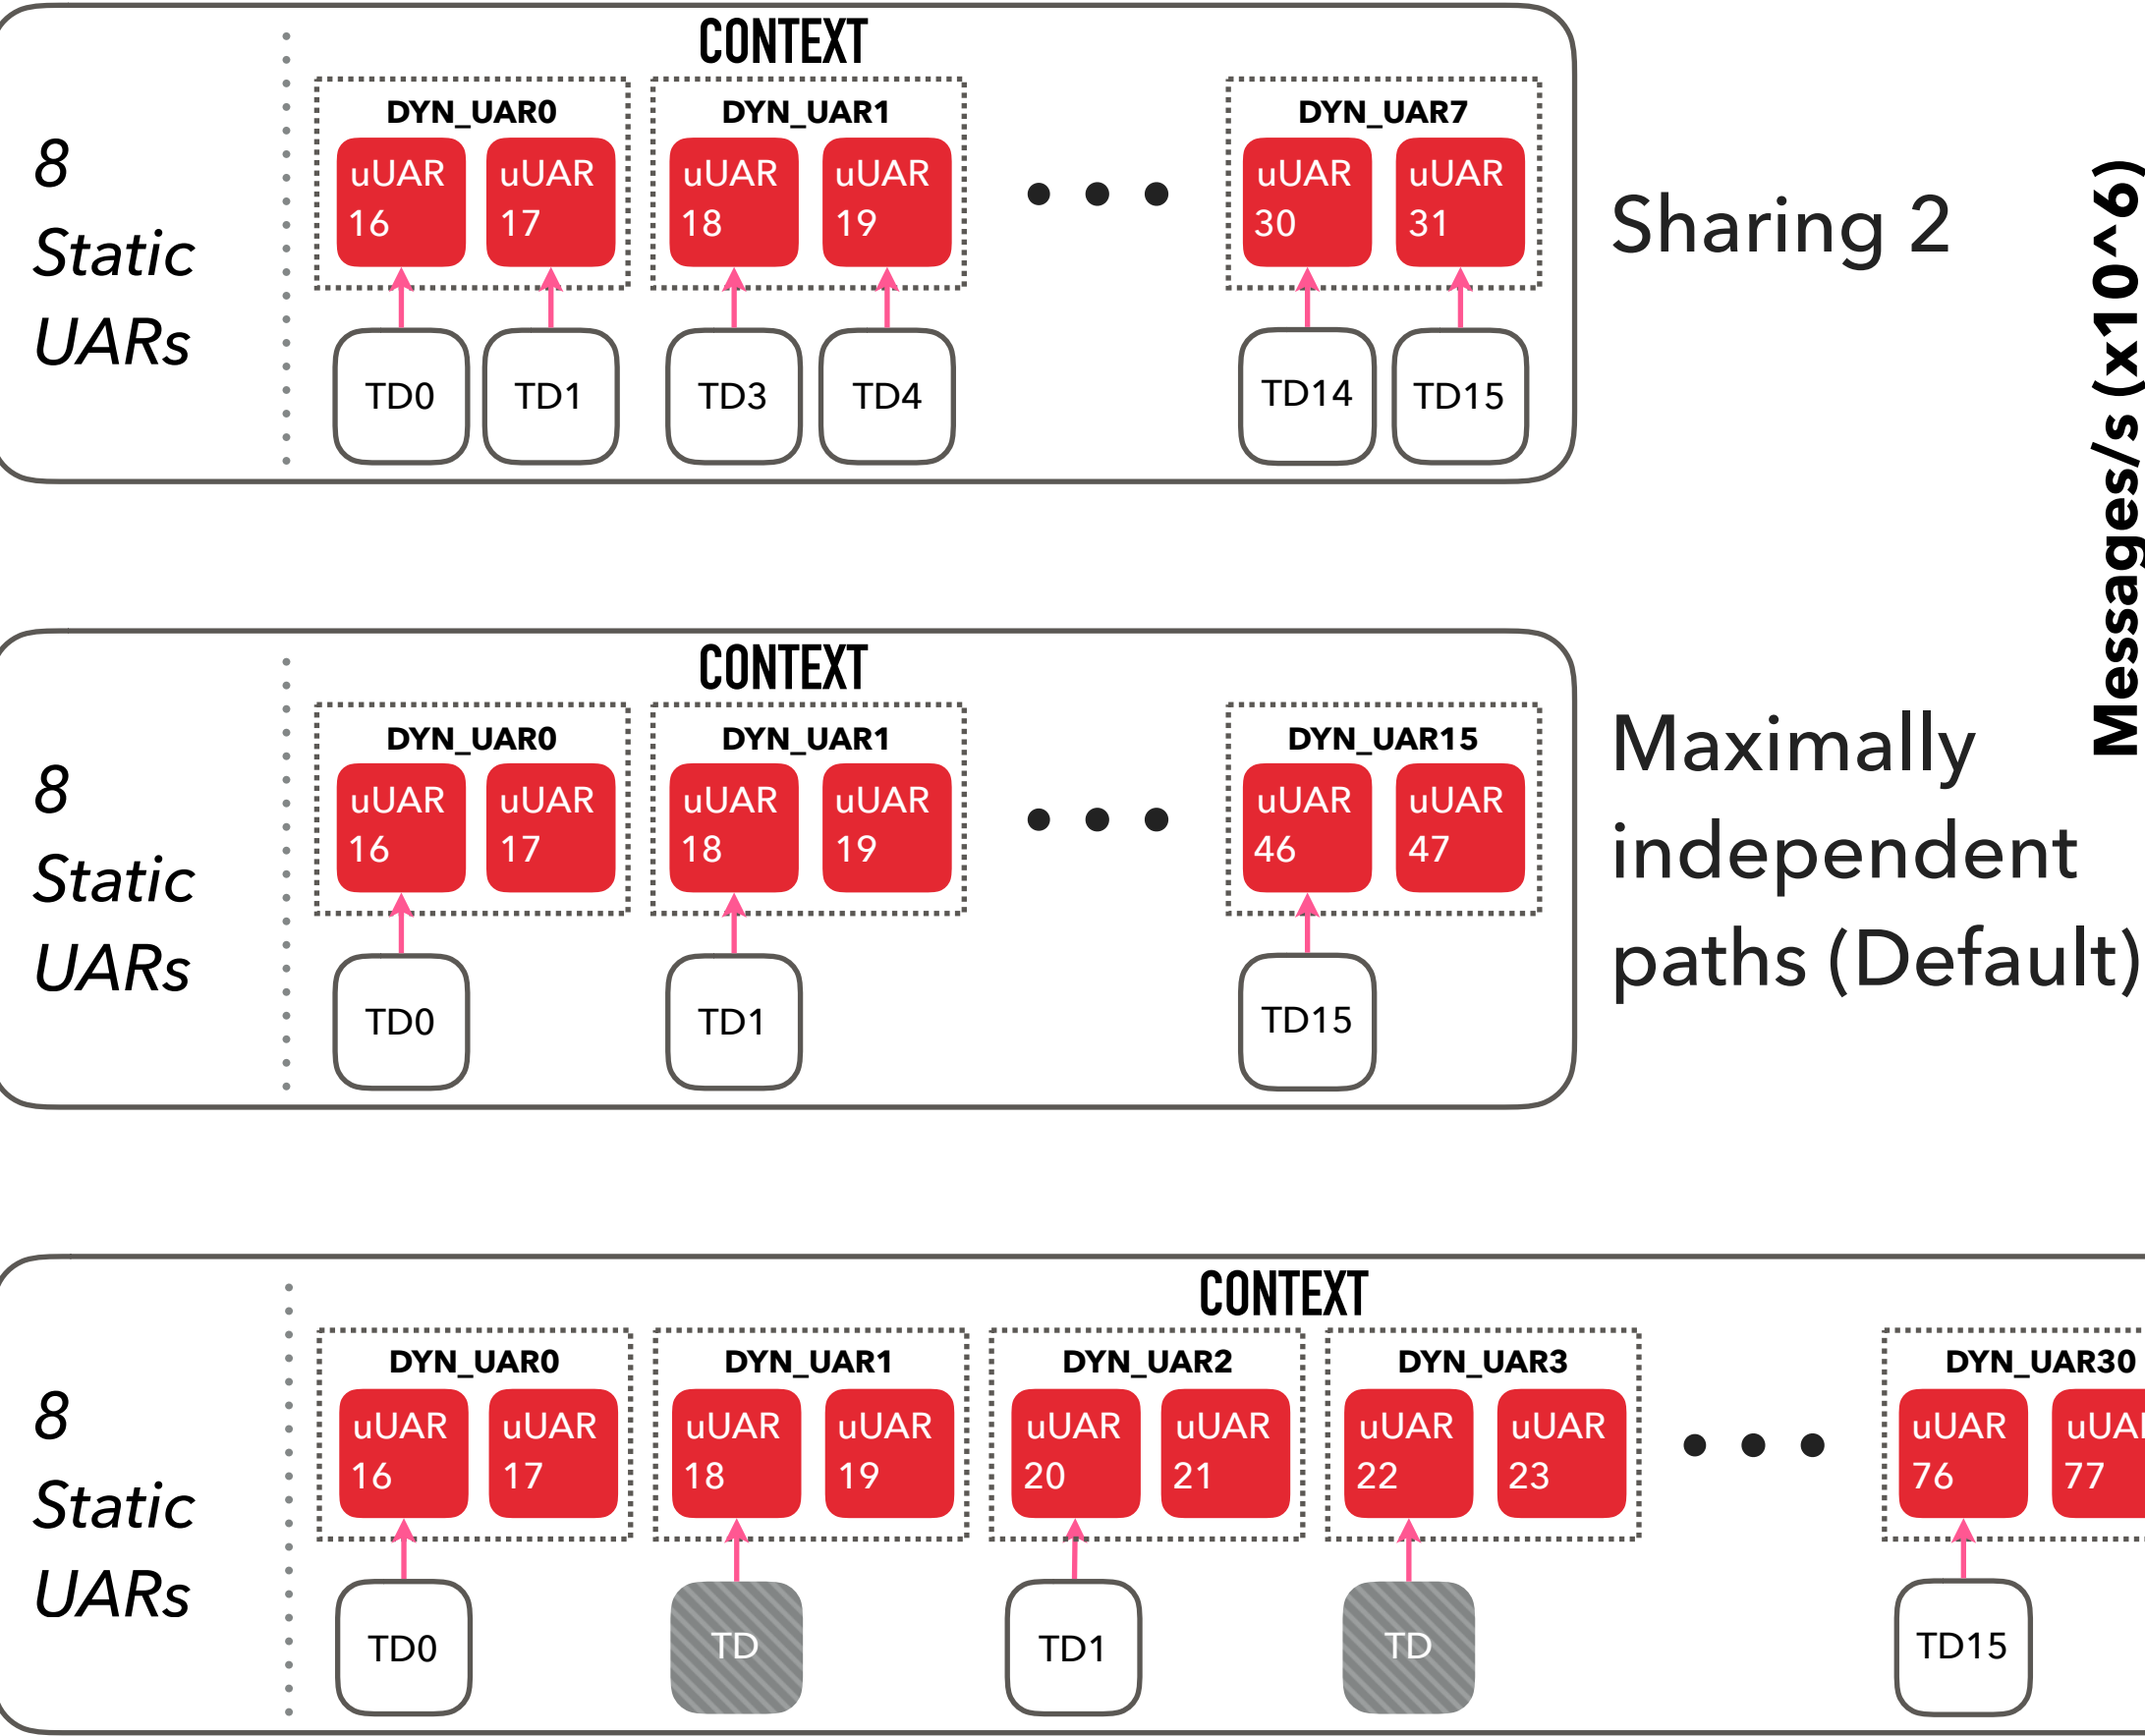
# CONTEXT SHARING

Sharing the UAR hurts performance with BlueFlame

No impact on QPs and CQs. Very impactful on UAR and uUAR usage.





CONTEXT SHARING





# PROTECTION DOMAIN OR MEMORY REGION SHARING

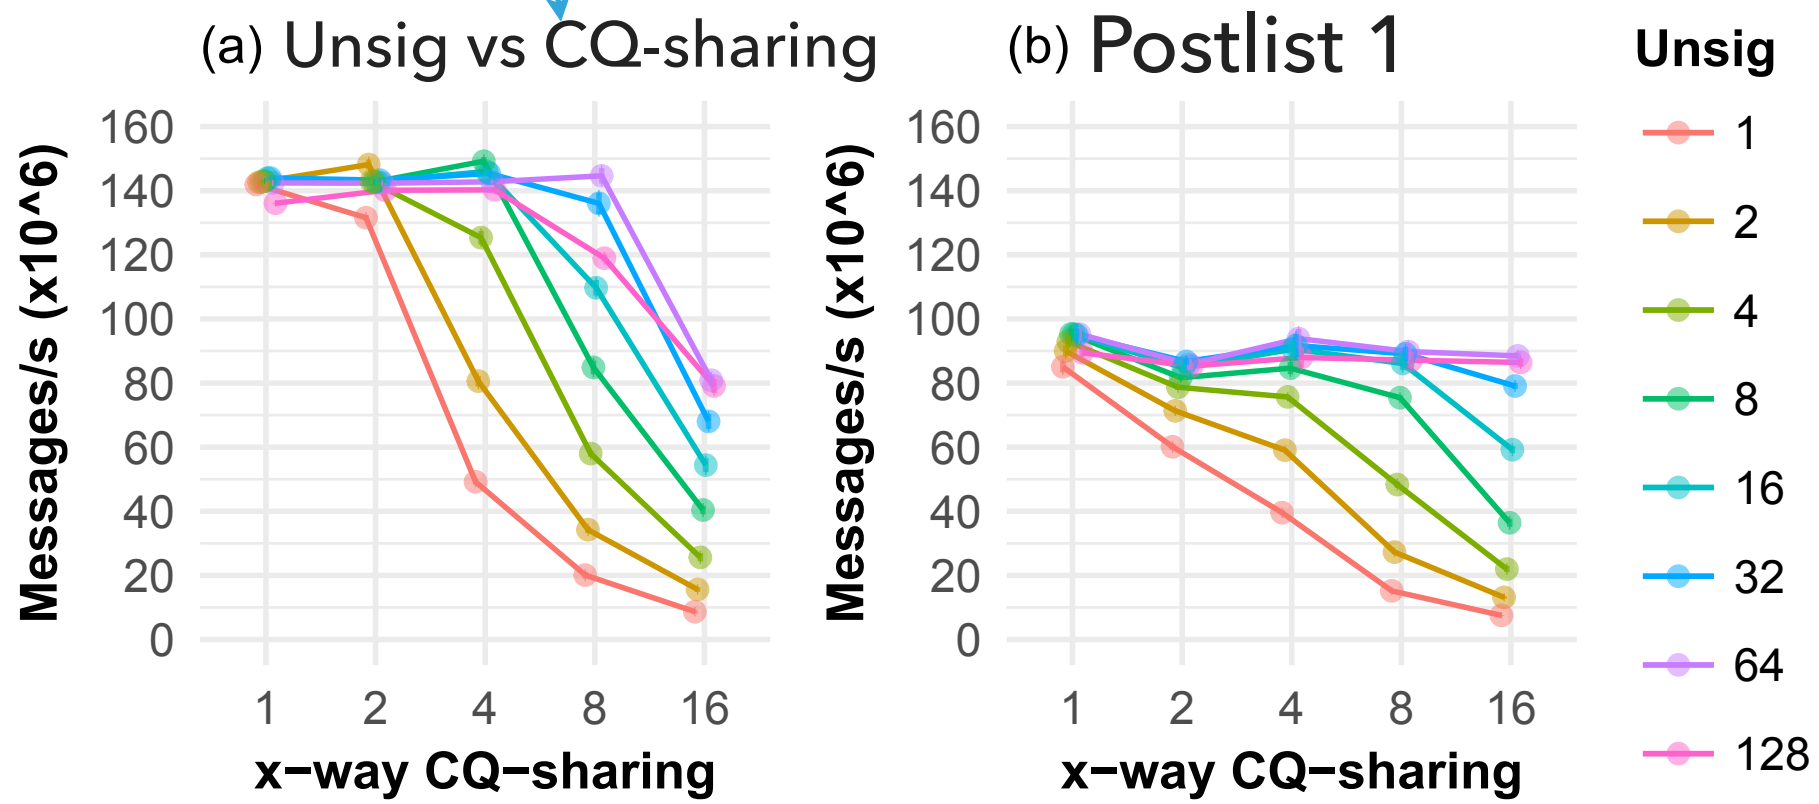


No impact on performance or  
communication resource usage



# COMPLETION QUEUE SHARING

Without unsigaled, lock-contention, and atomics on the CQ degrades performance.



Impacts only CQ usage.



# QUEUE PAIR SHARING

Contention on QP's lock, atomics, and decrease in utilization of underlying network parallelism.

Impacts only QP and CQ usage.





# LESSONS LEARNED FROM ANALYSIS

- ▶ Each thread must have its own cache-aligned buffer
- ▶ Can use Protection Domain and Memory Region at will.
- ▶ CTX-sharing the most critical for hardware resource usage.

16 threads	% Performance		Lesser uUARs
No CTX sharing	100%		288
2xDynamic	100%		3.6x (80)
Default TDs	80%		6x (48)
Sharing 2	50%		9x (32)

### LESSONS LEARNED FROM ANALYSIS

- ▶ Each thread must have its own cache-aligned buffer
- ▶ Can use Protection Domain and Memory Region at will.
- ▶ CTX-sharing the most critical for hardware resource usage.
- ▶ Only QP- and CQ-sharing impact memory usage
  - ▶ However reduction in memory usage not as meaningful (1.1x) when compared to the consequent drop in performance (18x with 16-way CQ-sharing)

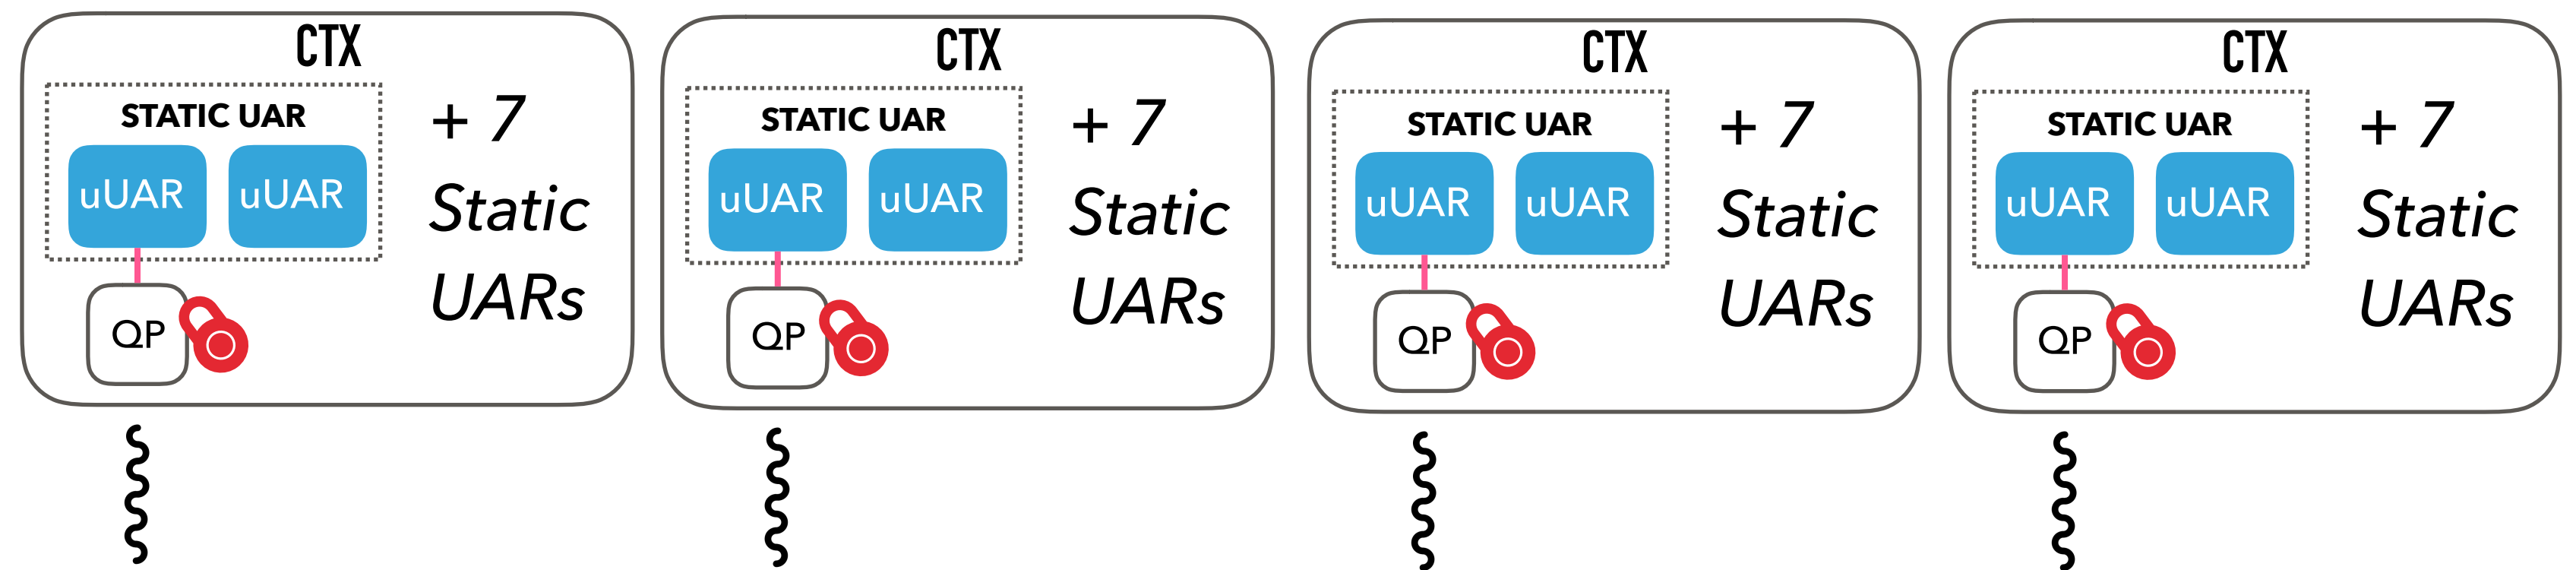
# SCALABLE ENDPOINTS

Max performance but not  
best


15 uUARs wasted per thread

Memory usage identity  
function of threads

MPI-everywhere



(1) Maximum independence

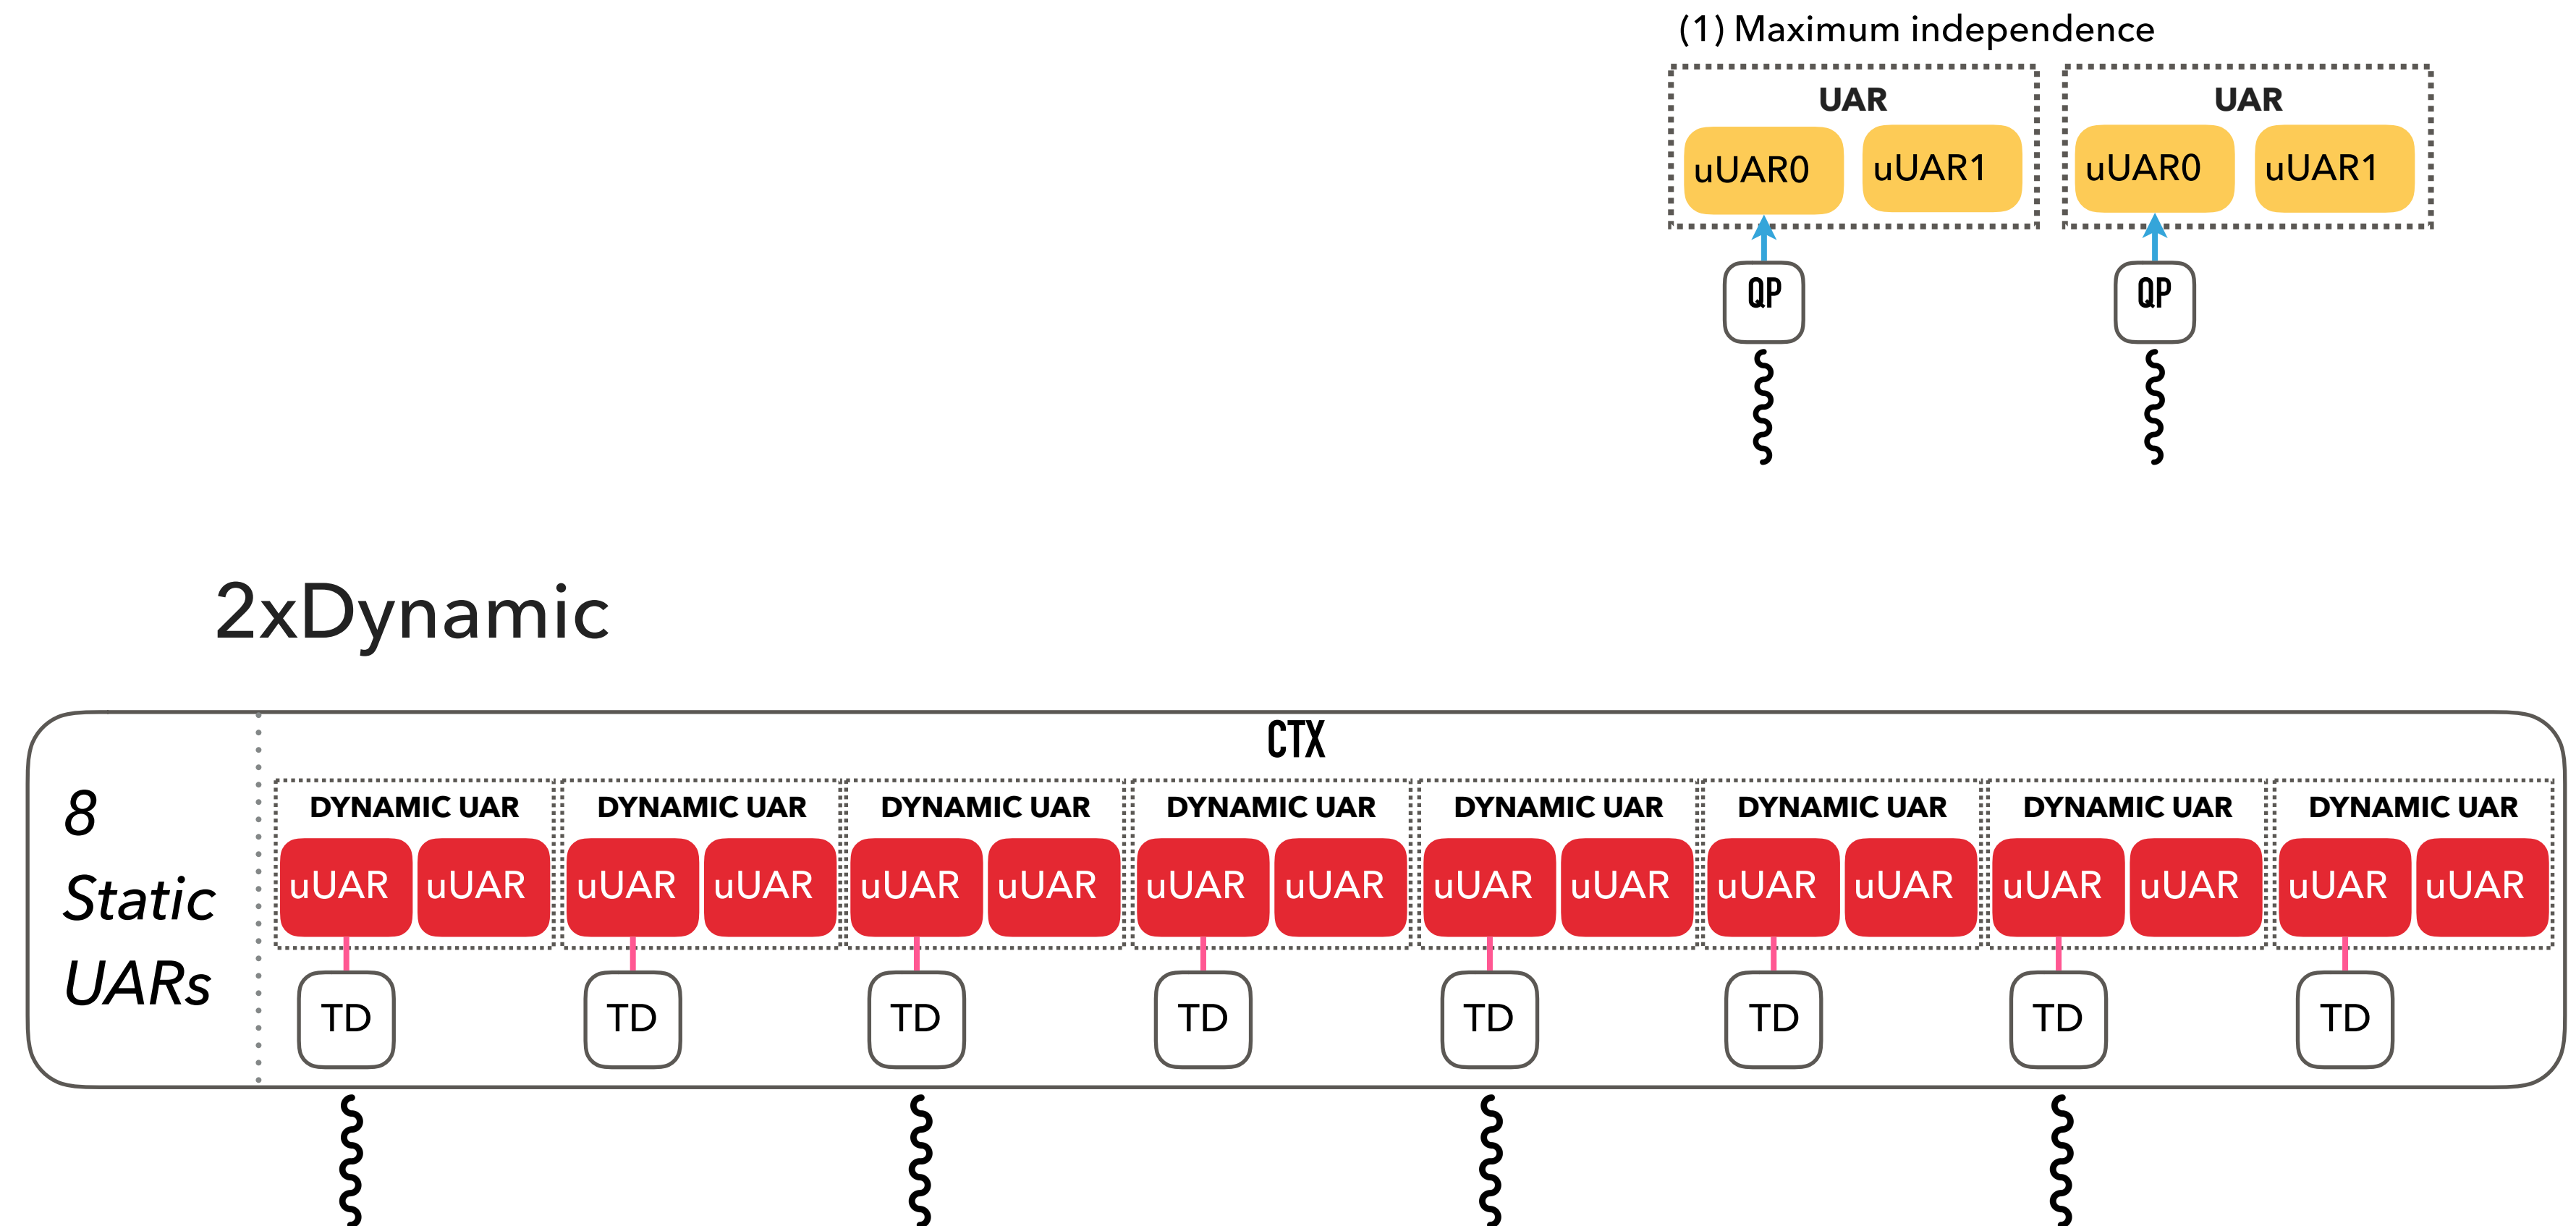


# SCALABLE ENDPOINTS

Best performance

3 uUARs per thread + 16  
uUARs wasted

Memory usage twice as that  
of MPI-everywhere




# SCALABLE ENDPOINTS

Will hurt performance

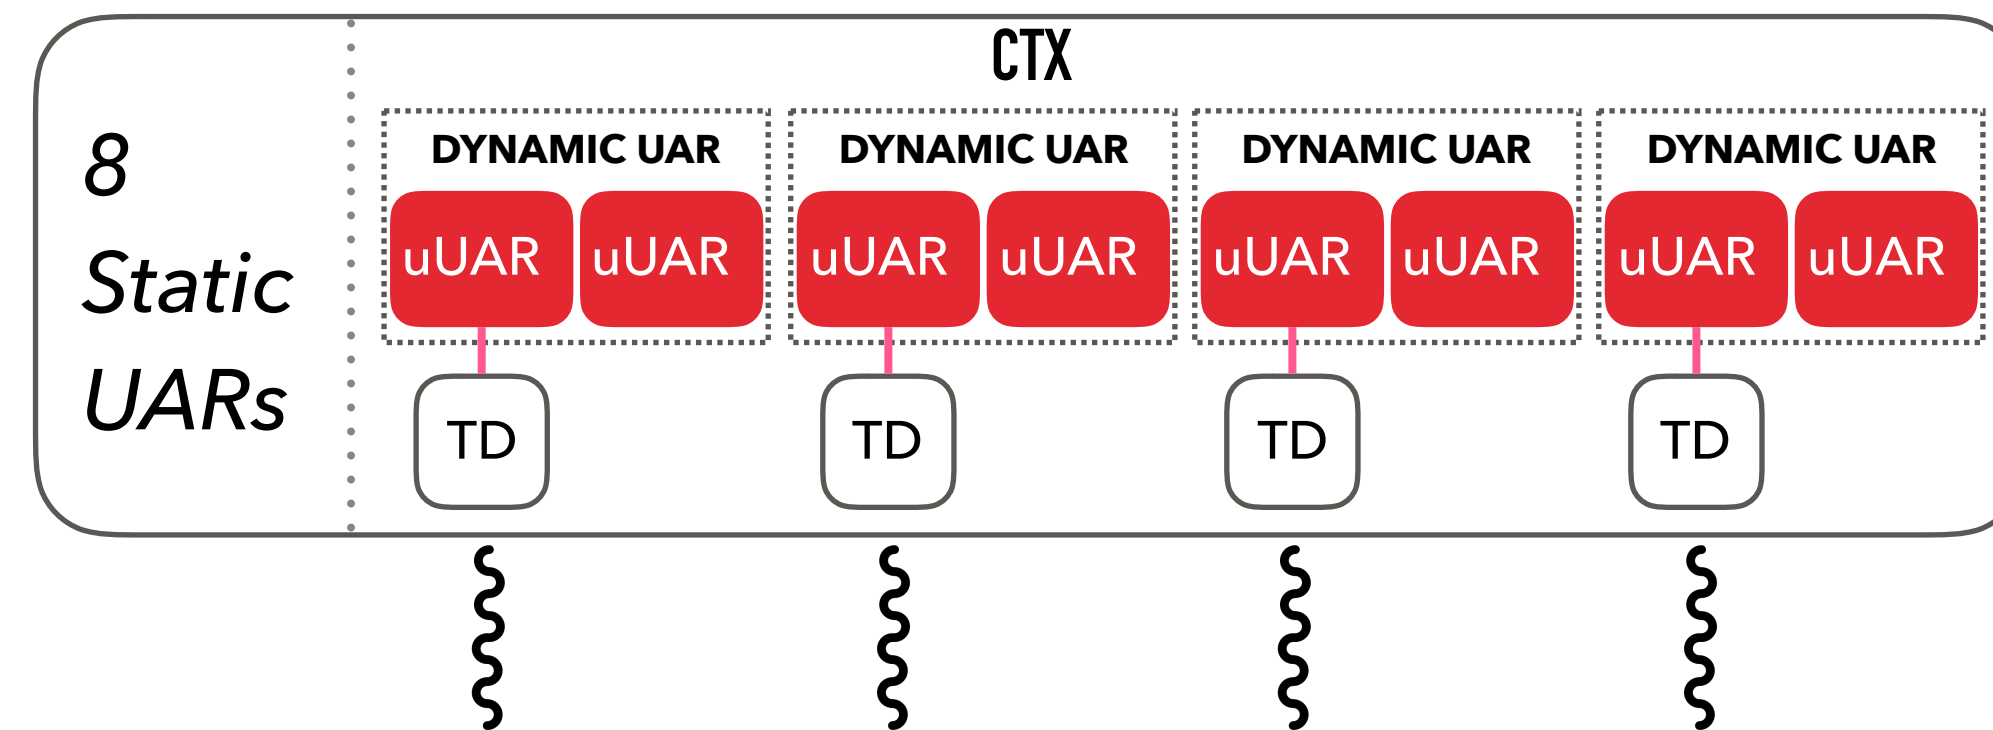
1 uUAR per thread + 16  
uUARs wasted

Memory usage same as MPI-  
everywhere

(1) Maximum independence



Dynamic




# SCALABLE ENDPOINTS

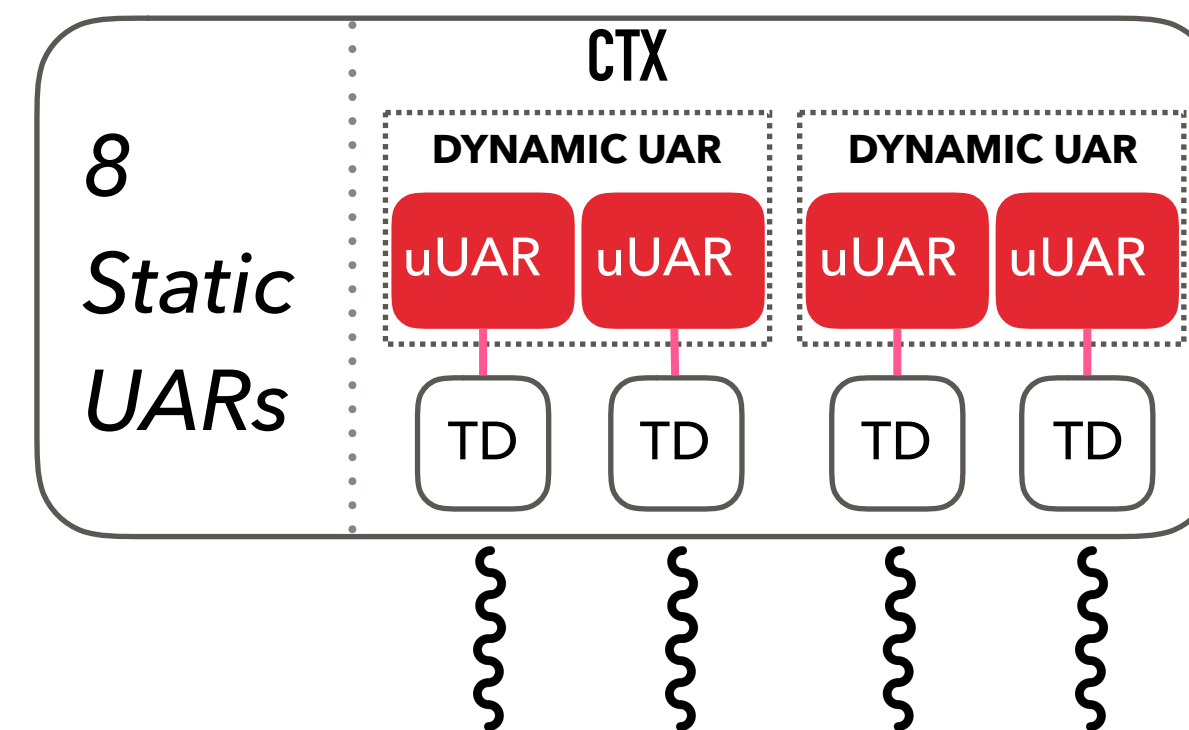
Will hurt performance

16 uUARs wasted

Memory usage same as  
MPI-everywhere



## Shared Dynamic




# SCALABLE ENDPOINTS

Locks + UAR sharing will  
hurt performance


Memory usage same as  
MPI-everywhere

(2) Shared UAR




+

(3) Shared uUAR



Static

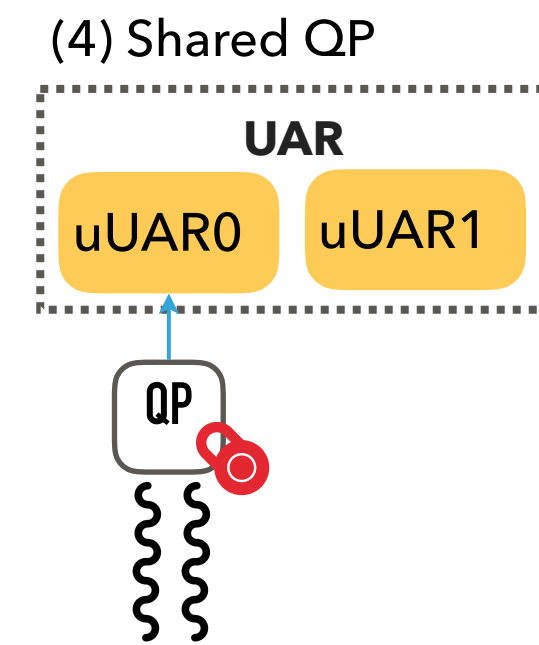



# SCALABLE ENDPOINTS

Worst performance

Least resource usage

MPI+threads






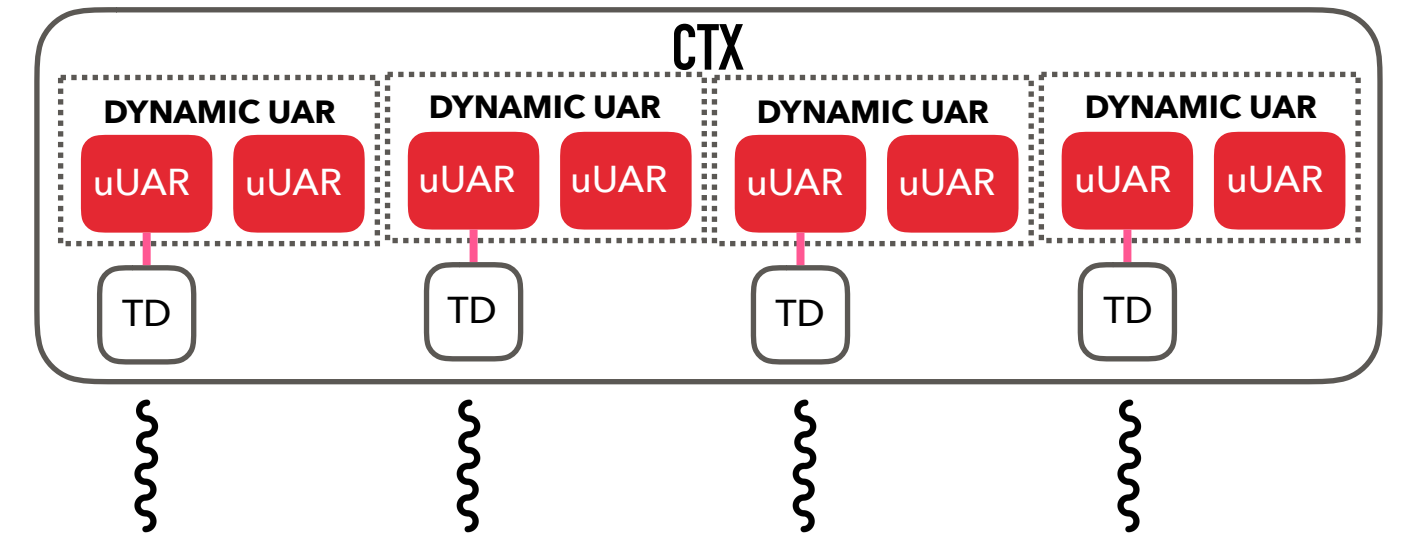
# SCALABLE ENDPOINTS

CQ can be shared in any manner for any of these categories.

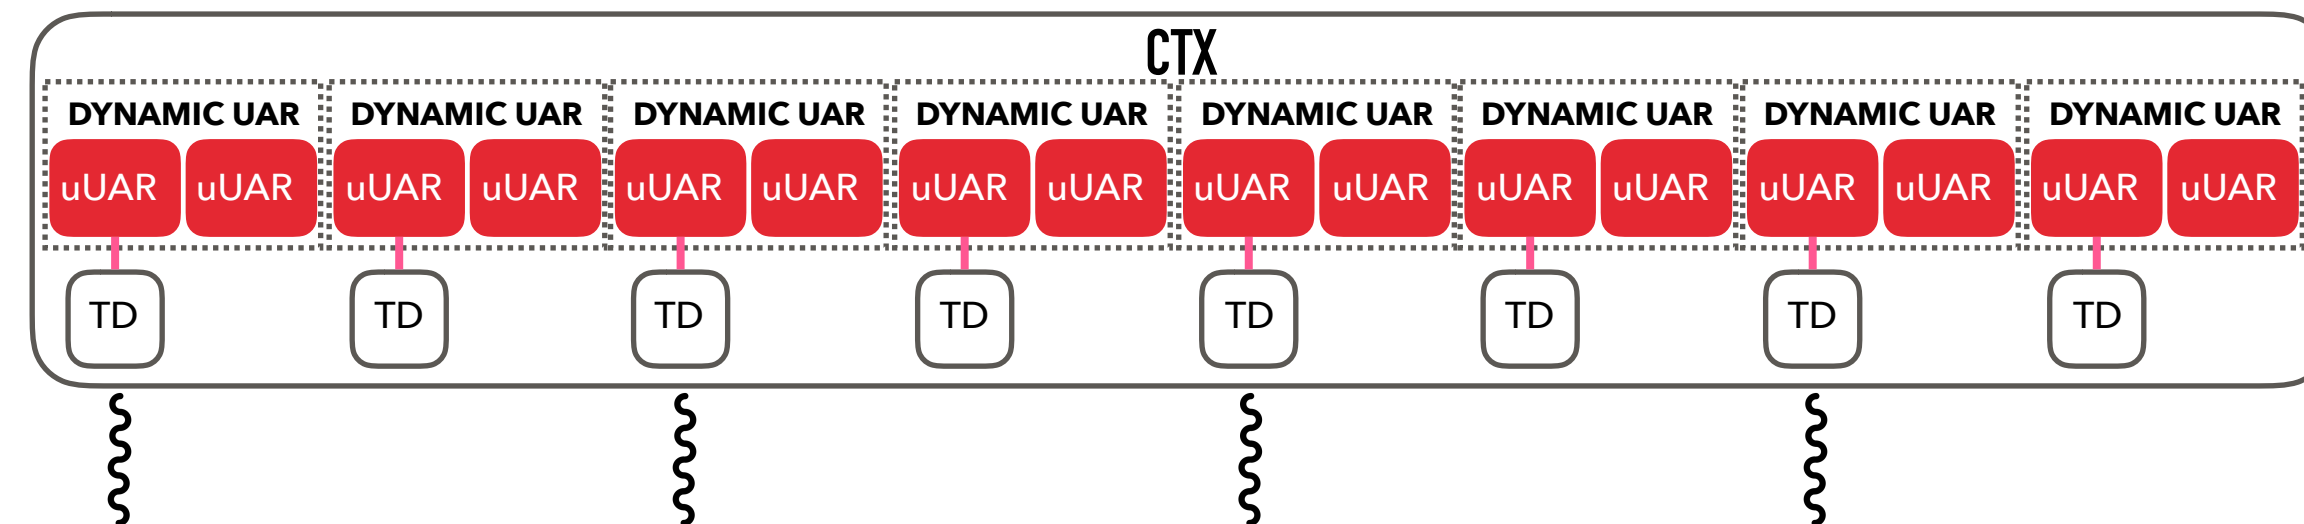
MPI-everywhere




Dynamic



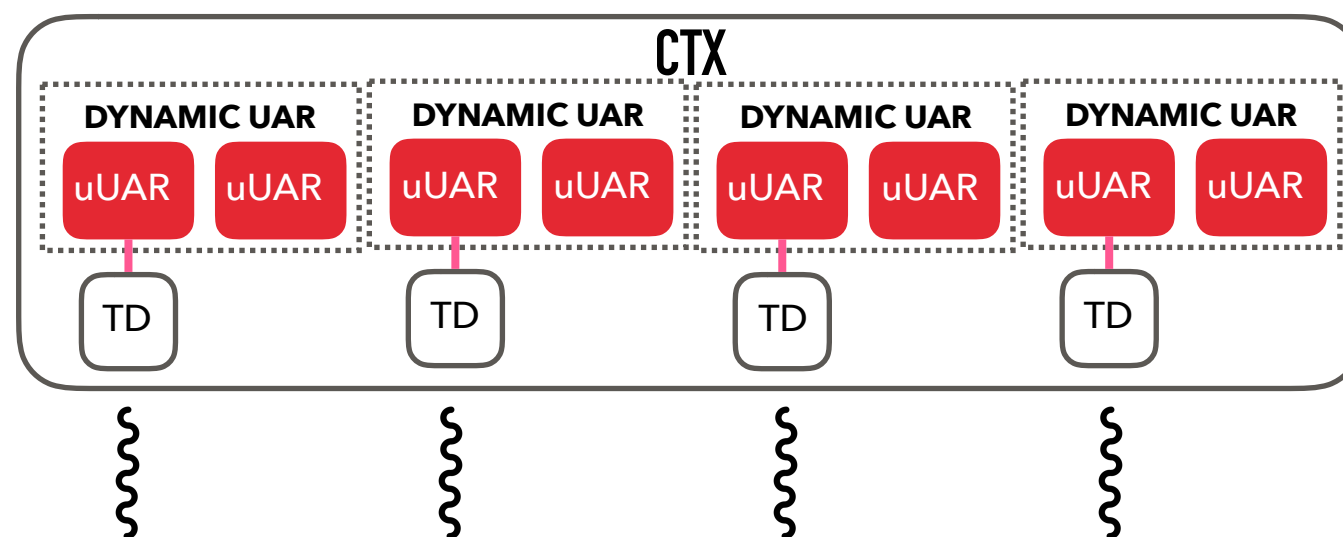
2xDynamic




Static



Dynamic





MPI+threads




## GLOBAL ARRAY EVALUATION

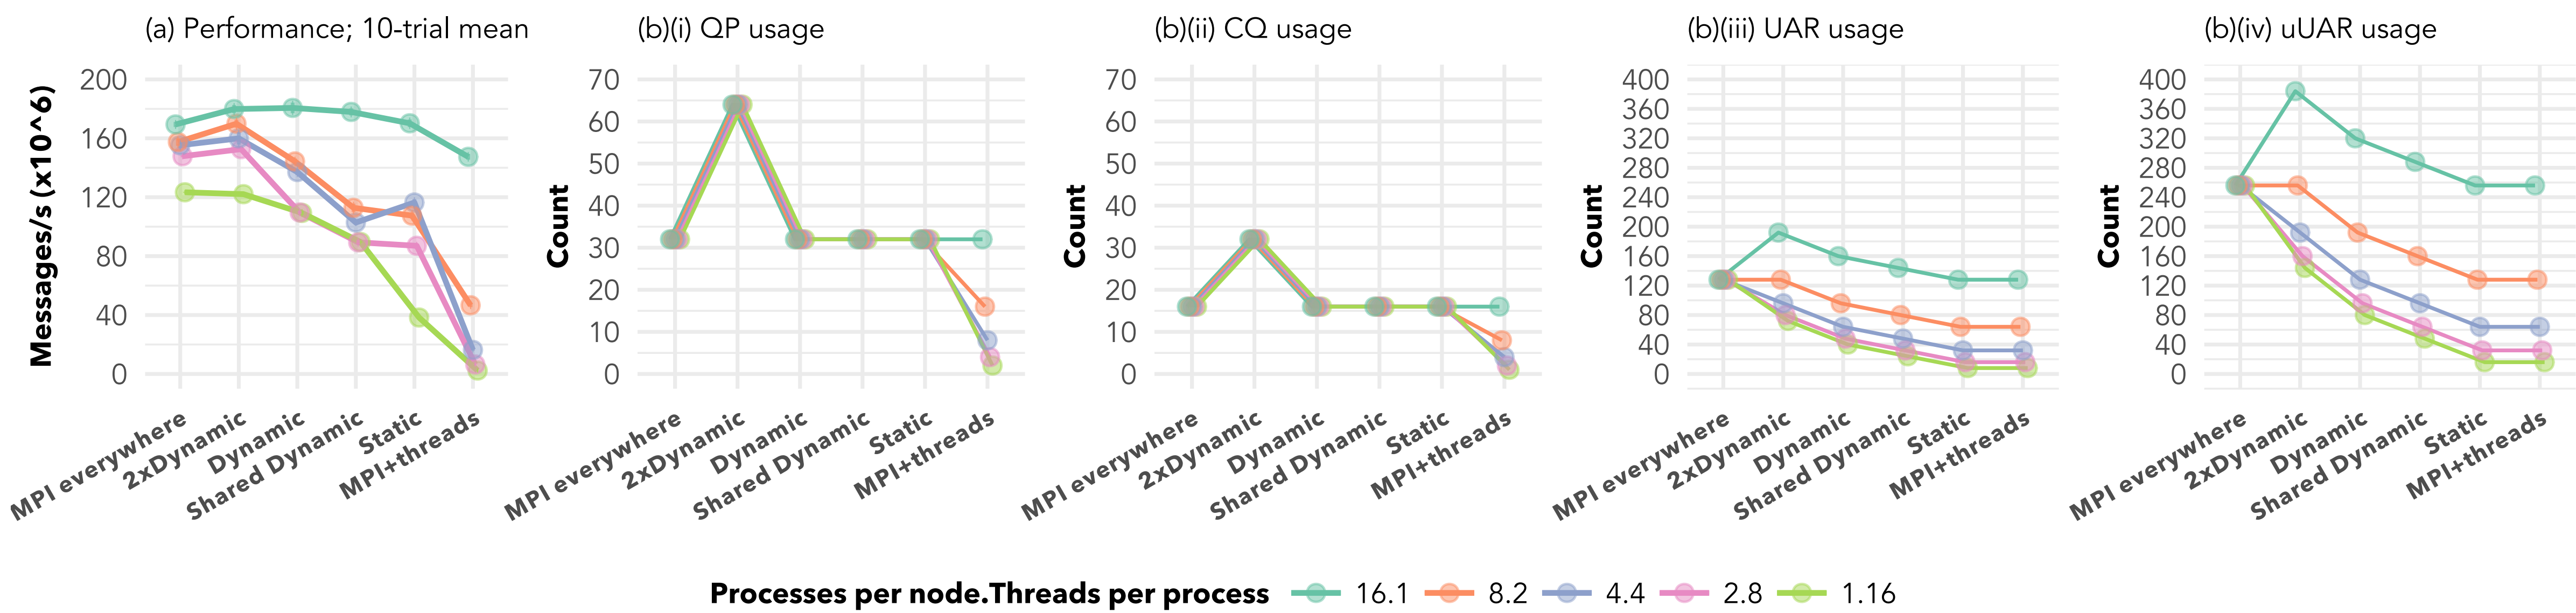
- ▶ Fetching and writing tiles from a global array
- ▶ Double precision global matrix multiply (DGEMM)  $A \times B = C$ 
  - ▶ All three sit on a server node
  - ▶ Client node RDMA-gets tiles, multiplies them, RDMA-puts them back
- ▶ Trend: Performance decreases with increasing resource sharing



# STENCIL EVALUATION



- ▶ 5-pt stencil, 1D partitioning with varying hybrid environments for 16 threads.
- ▶ Similar trend of decreasing performance with increasing resource sharing



### RELATED WORK AND FUTURE DIRECTION

- ▶ MPI Endpoints proposal (<https://github.com/mpi-forum/mpi-issues/issues/56>)
  - ▶ Top-down approach: create extra MPI processes to serve as “endpoints” for threads.
  - ▶ James Dinan et al.
- ▶ Ours is a bottom-up approach.
- ▶ Figure out how to expose this to MPI through OFI/UCX middleware.
- ▶ MPI library can handle endpoints OR expose to MPI user?

---

**BACKUP SLIDES**

## MAPPING UAR TO QP

- ▶ During QP creation, mlx5 maps the QP to a uUAR
- ▶ The 16 uUARs considered to be “statically” allocated. mlx5 categorizes 4 of these into low latency (blue) and 11 into regular (yellow) and the zeroth into high latency (grey).
  - ▶ BlueFlame-writes possible only on low latency and regular uUARs. High latency uUARs allow only Doorbells.
  - ▶ Only 1 QP mapped to low latency uUAR. No lock needed for BlueFlame-writes. Hence, “low latency”.
  - ▶ Regular uUARs may be mapped to multiple QPs. Lock necessary for concurrent BlueFlame-writes.
  - ▶ High latency uUAR needs no locks for concurrent Doorbells since Doorbells are atomic. Used only when no regular uUARs are present.
- ▶ TDs assigned to new, “dynamically” allocated uUARs. Even TDs allocate new UAR page. Odd TDs use uUAR of already allocated UAR.

